

EECS151 : Introduction to Digital Design and ICs

Lecture 8 – RISC-V Datapath and Control

Bora Nikolić and Sophia Shao



Dual-Core RISC-V 64Bit K210 AI Board – Kendryte KD233

The Kendryte K210 is a system-on-chip (SoC) that integrates machine vision and machine hearing. Using TSMC's ultra-low-power 28-nm advanced process with dual-core 64-bit processors for better power efficiency, stability and reliability.



<https://www.analoglamb.com/product/dual-core-risc-v-64bit-k210-ai-board-kendryte-kd233/>



Review

- State machines:
 - Specify circuit function
 - Draw state transition diagram
 - Write down symbolic state-transition table
 - Assign encodings (bit patterns) to symbolic states
 - Code as Verilog behavioral description
- RISC-V processor
 - A large state machine
 - Datapath + control
 - Reviewed R-format instructions

Implementing other R-Format instructions

00000000	rs2	rs1	000	rd	0110011	add
01000000	rs2	rs1	000	rd	0110011	sub
00000000	rs2	rs1	001	rd	0110011	sll
00000000	rs2	rs1	010	rd	0110011	slt
00000000	rs2	rs1	011	rd	0110011	sltu
00000000	rs2	rs1	100	rd	0110011	xor
00000000	rs2	rs1	101	rd	0110011	srl
01000000	rs2	rs1	101	rd	0110011	sra
00000000	rs2	rs1	110	rd	0110011	or
00000000	rs2	rs1	111	rd	0110011	and

- All implemented by decoding funct3 and funct7 fields and selecting appropriate ALU function

Instruction Encoding

- Instructions are encoded to simplify logic
 - **sub** and **sra** differ in Inst[30] from **add** and **srl**
- **RV64I** widens registers (XLEN=64)
- Additional instructions manipulate 32-bit values, identified by a suffix **W**
 - **ADDW, SUBW**
 - RV64I opcode field for 'W' instructions is **0111011** (**0110011** for RV32I)

00000000	rs2	rs1	000	rd	0111011
00000000	rs2	rs1	000	rd	0110011

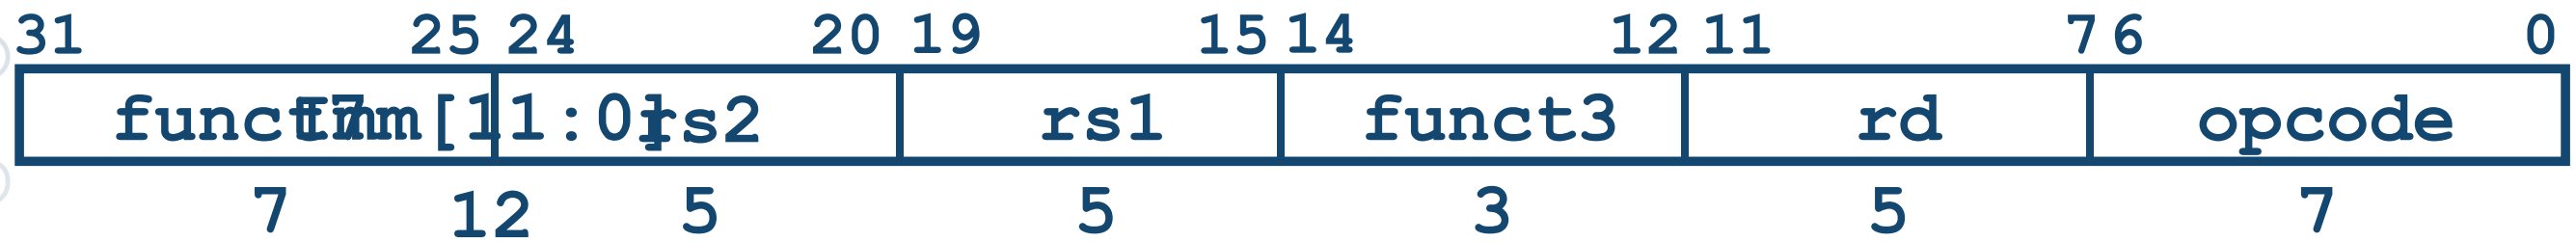
32b

addw

add

64b

I-Format Instruction Layout



- Only one field is different from R-format, rs2 and funct7 replaced by 12-bit signed immediate, **imm[11:0]**
- Remaining fields (rs1, funct3, rd, opcode) same as before
- imm[11:0] can hold values in range $[-2048_{\text{ten}}, +2047_{\text{ten}}]$
- Immediate is always sign-extended to 32-bits before use in an arithmetic operation
- Other instructions handle immediates > 12 bits

All RV32 I-format Arithmetic Instructions

imm[11:0]		rs1	000	rd	0010011	addi
imm[11:0]		rs1	010	rd	0010011	slti
imm[11:0]		rs1	011	rd	0010011	sltiu
imm[11:0]		rs1	100	rd	0010011	xori
imm[11:0]		rs1	110	rd	0010011	ori
imm[11:0]		rs1	111	rd	0010011	andi
0000000	shamt	rs1	001	rd	0010011	slli
0000000	shamt	rs1	101	rd	0010011	srli
0100000	shamt	rs1	101	rd	0010011	srai

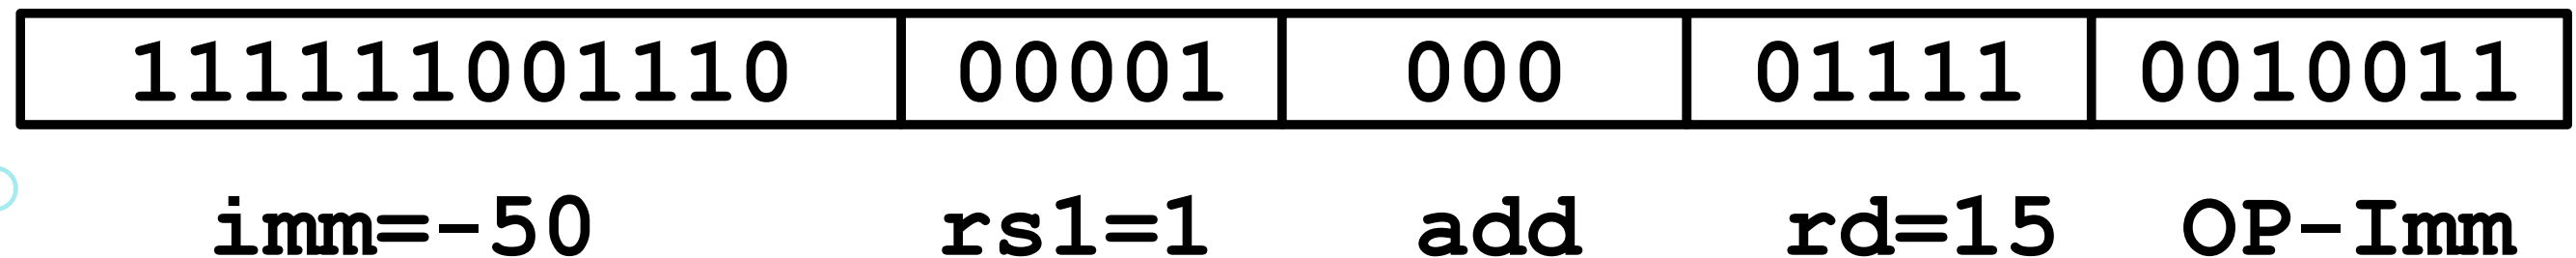
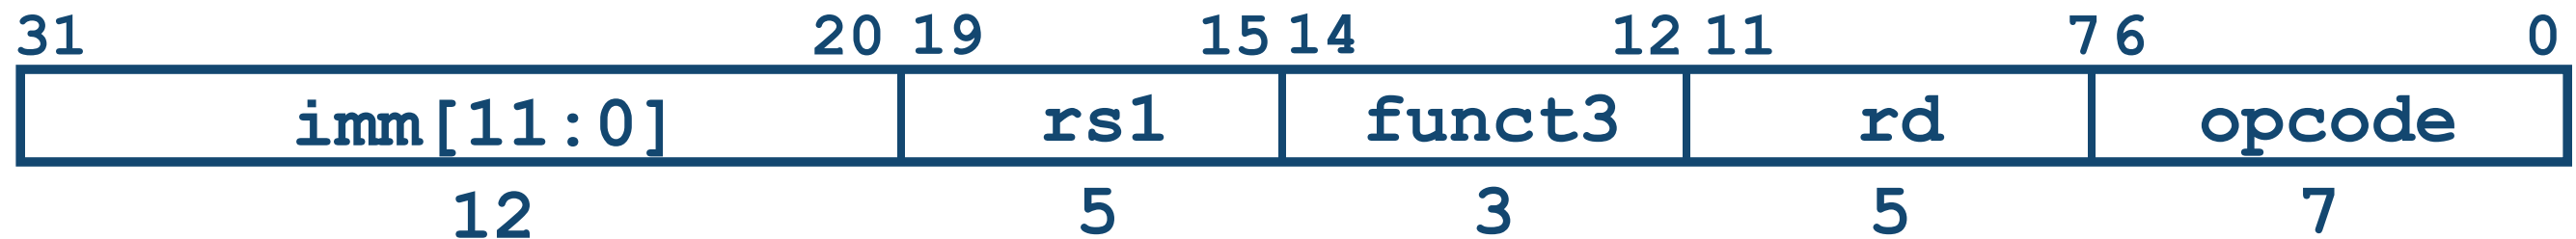
The same Inst[30] immediate bit is used to distinguish “shift right logical” (SRLI) from “shift right arithmetic” (SRAI)

“Shift-by-immediate” instructions only use lower 5 bits of the immediate value for shift amount (can only shift by 0-31 bit positions)

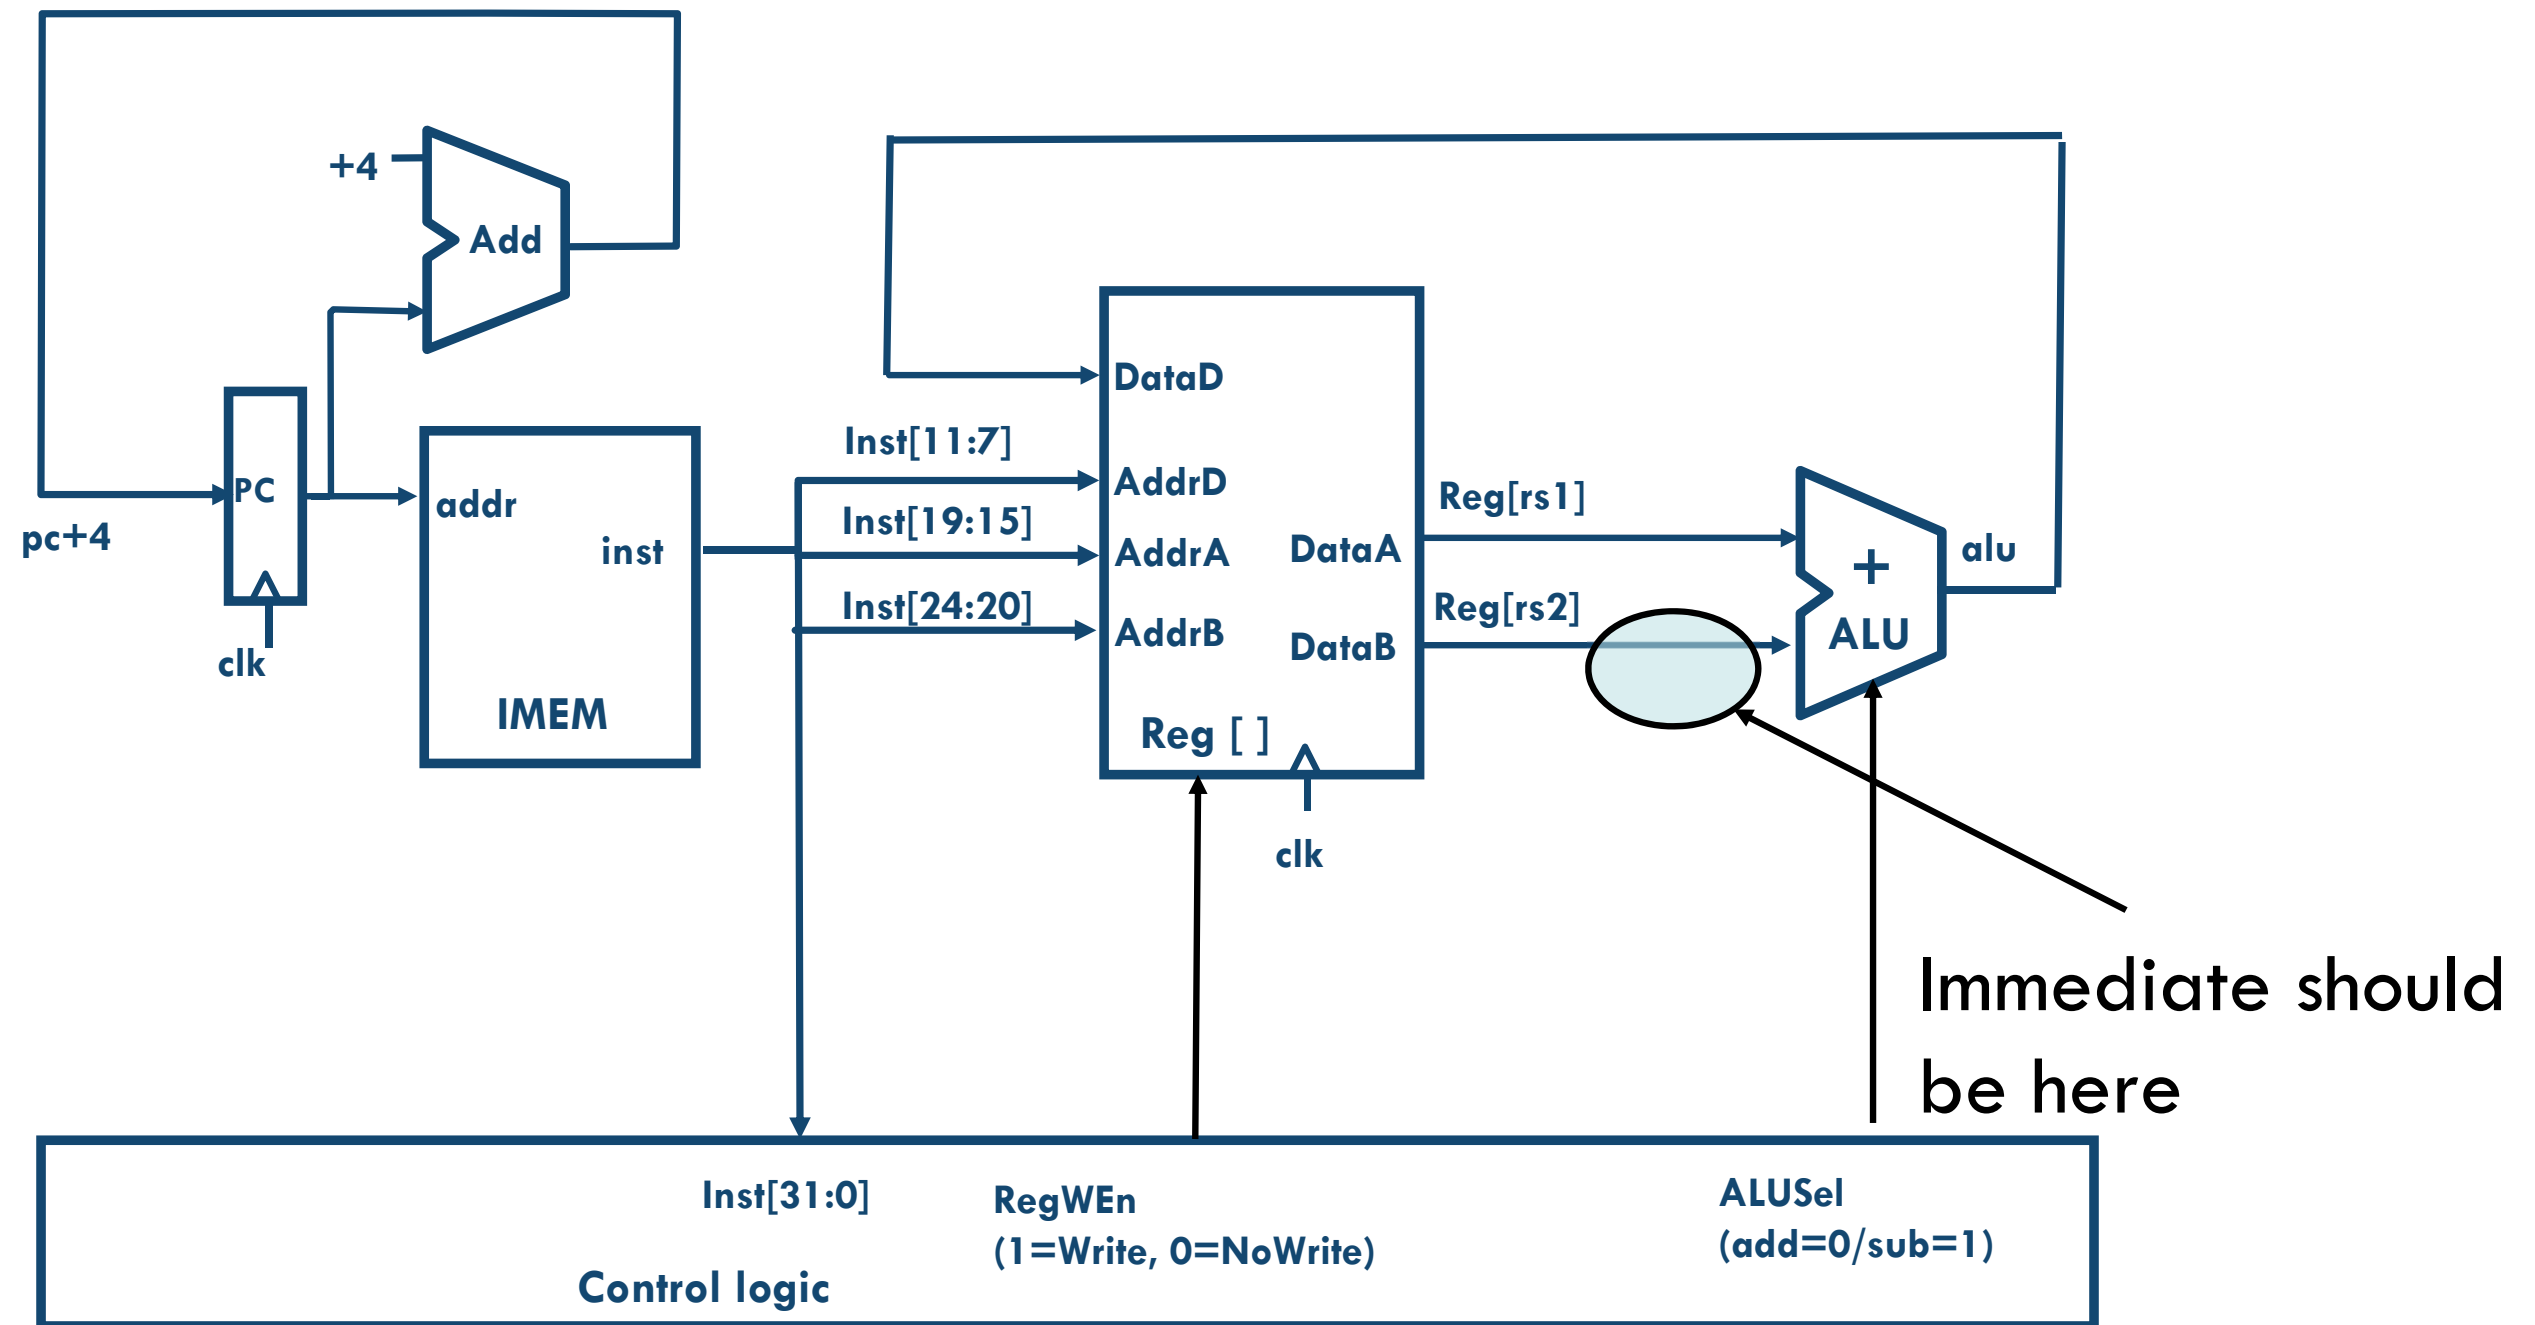
Implementing I-Format - `addi` instruction

- **RISC-V Assembly Instruction – add immediate:**

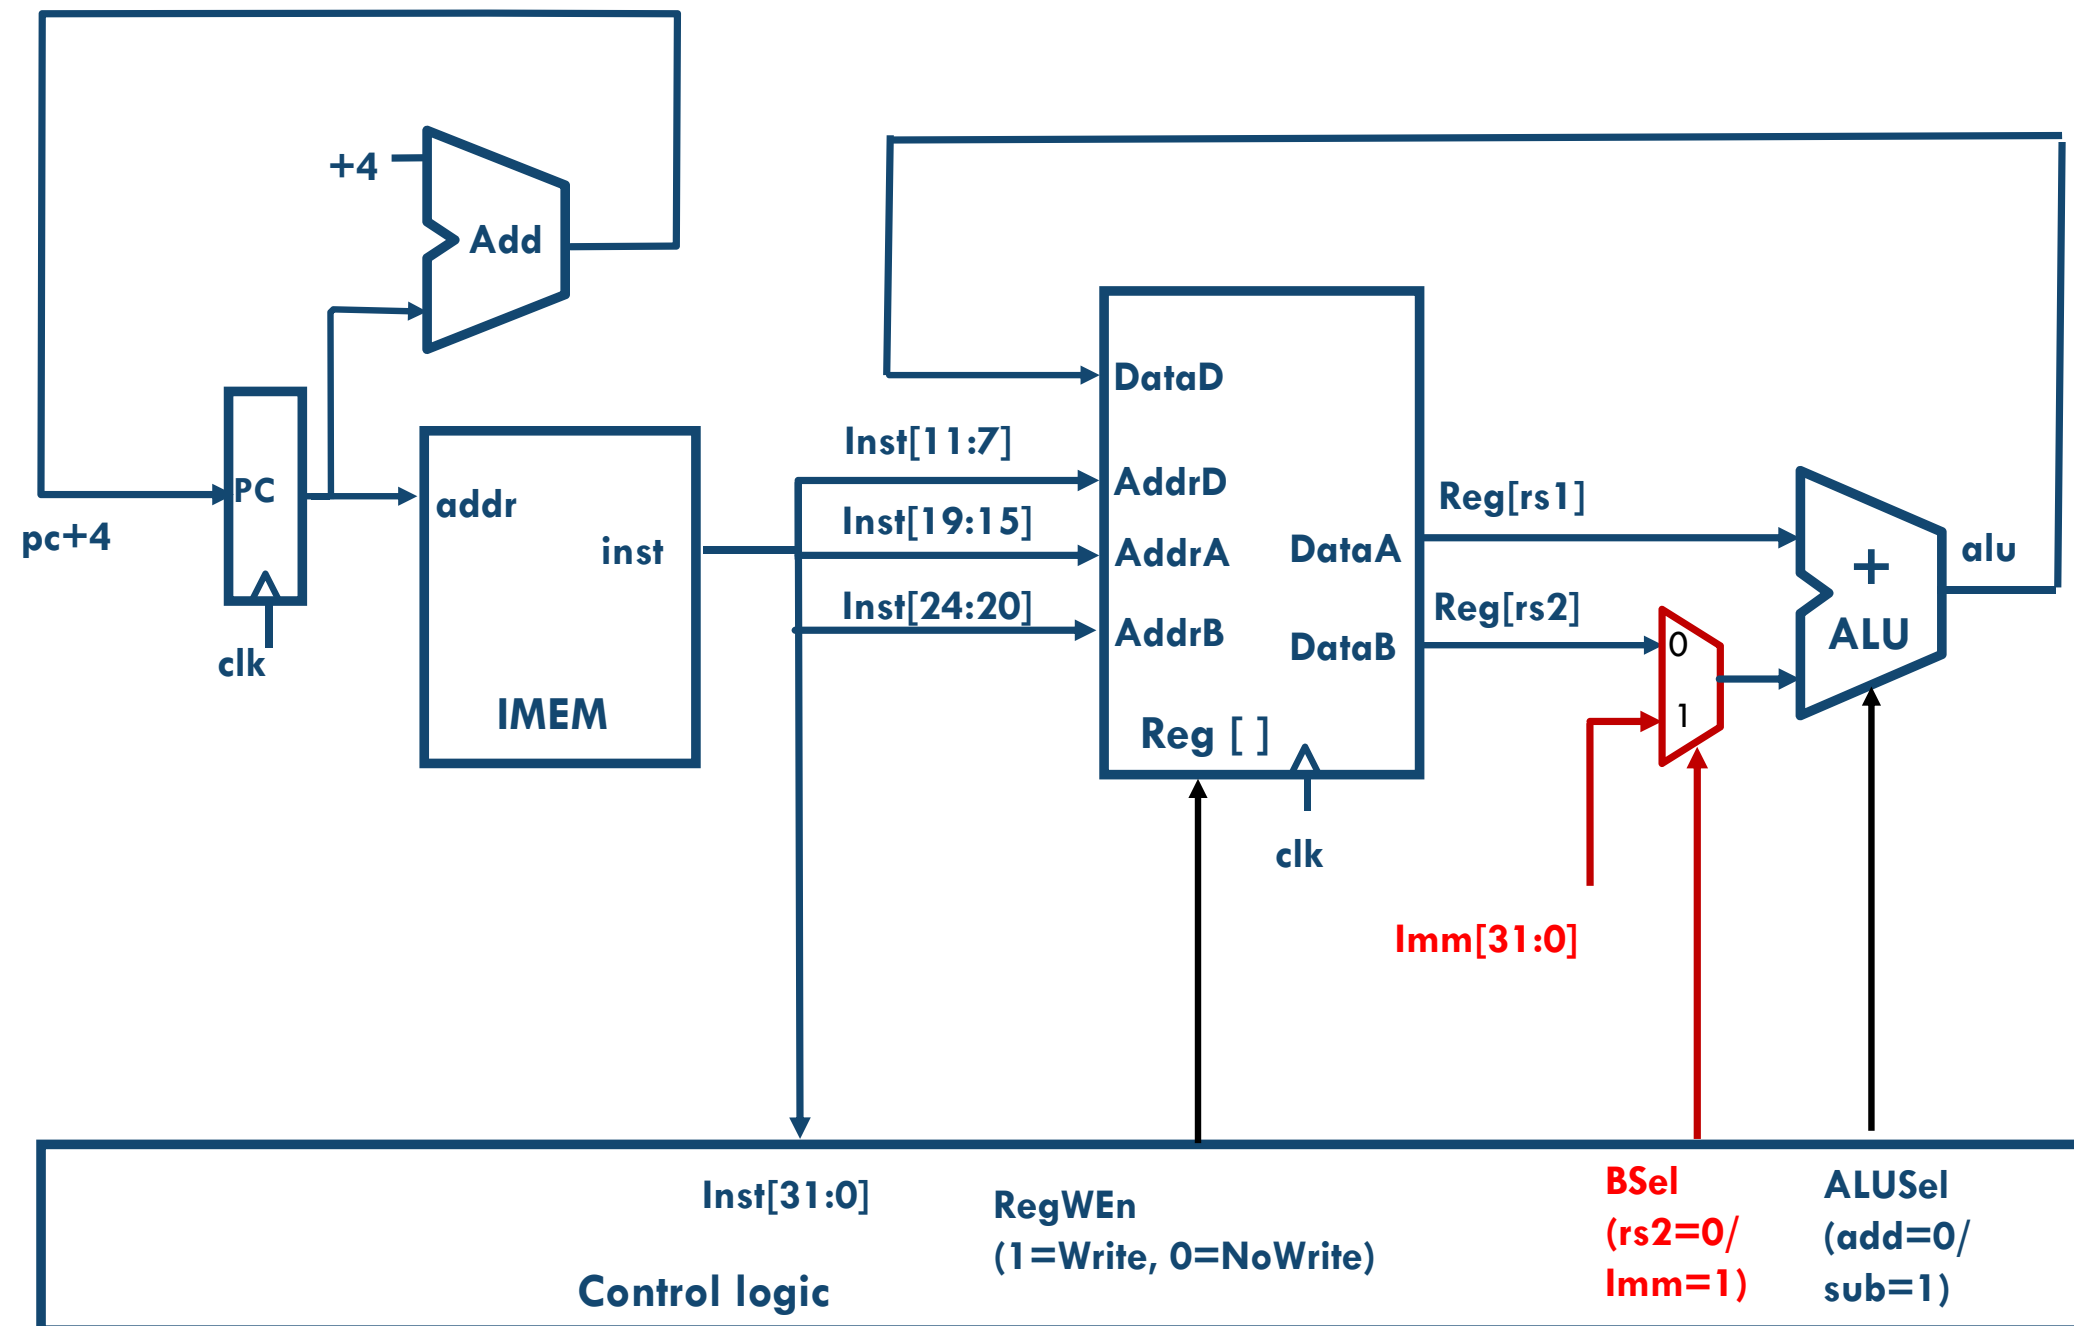
`addi x15,x1,-50`



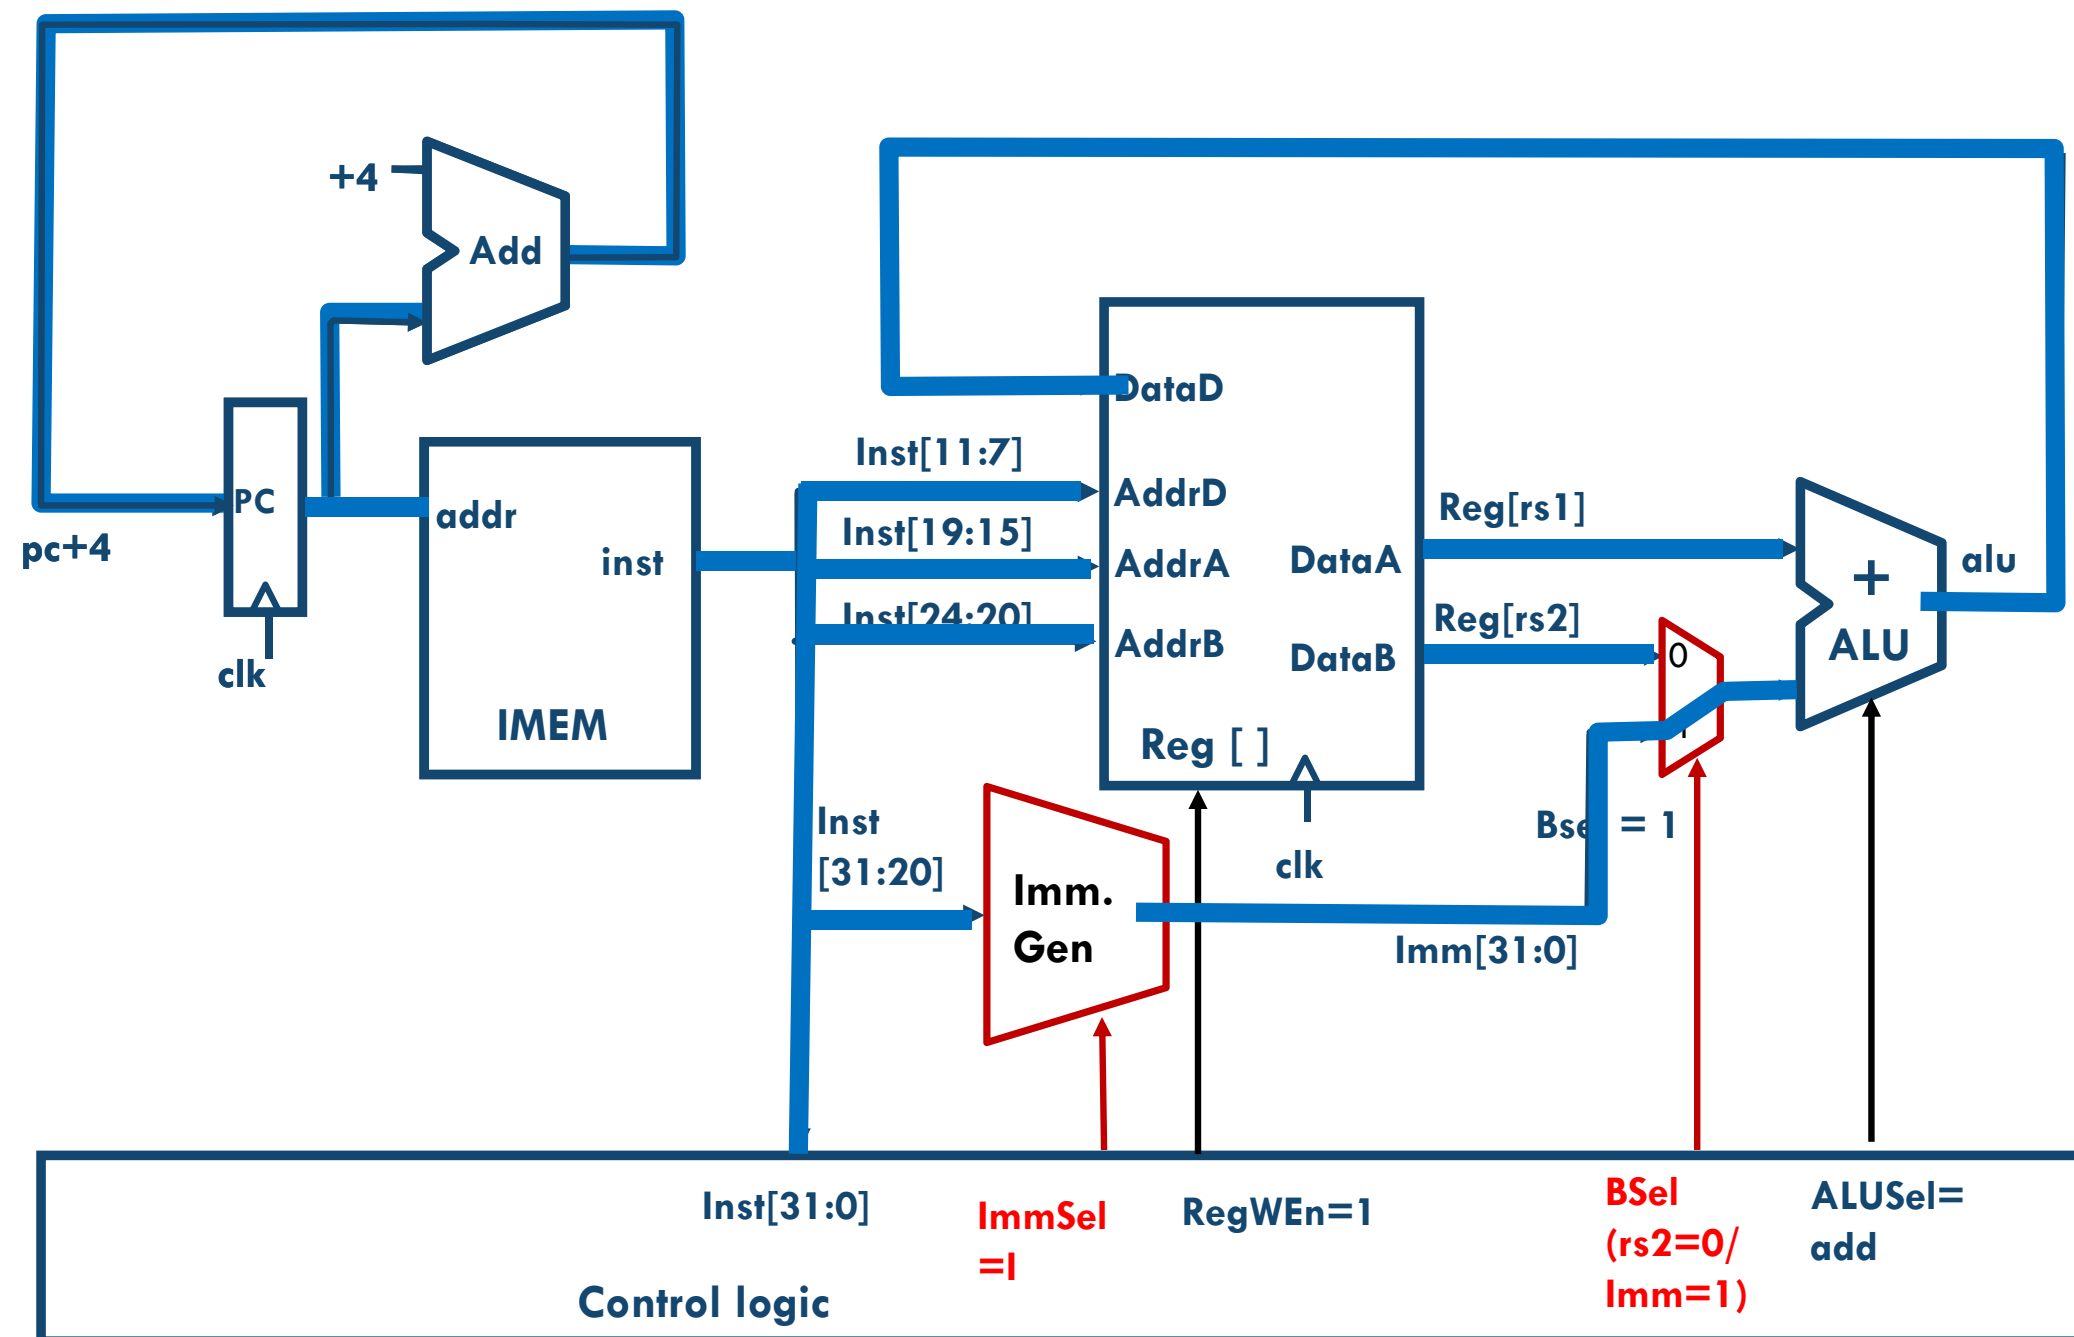
Datapath for add/sub



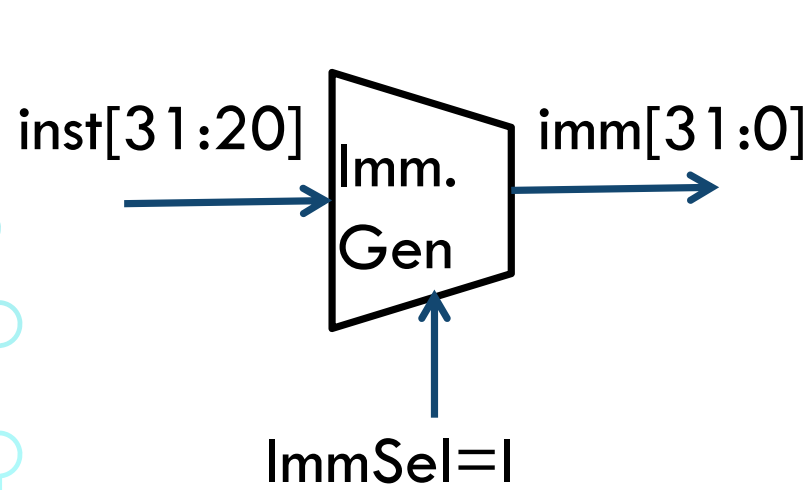
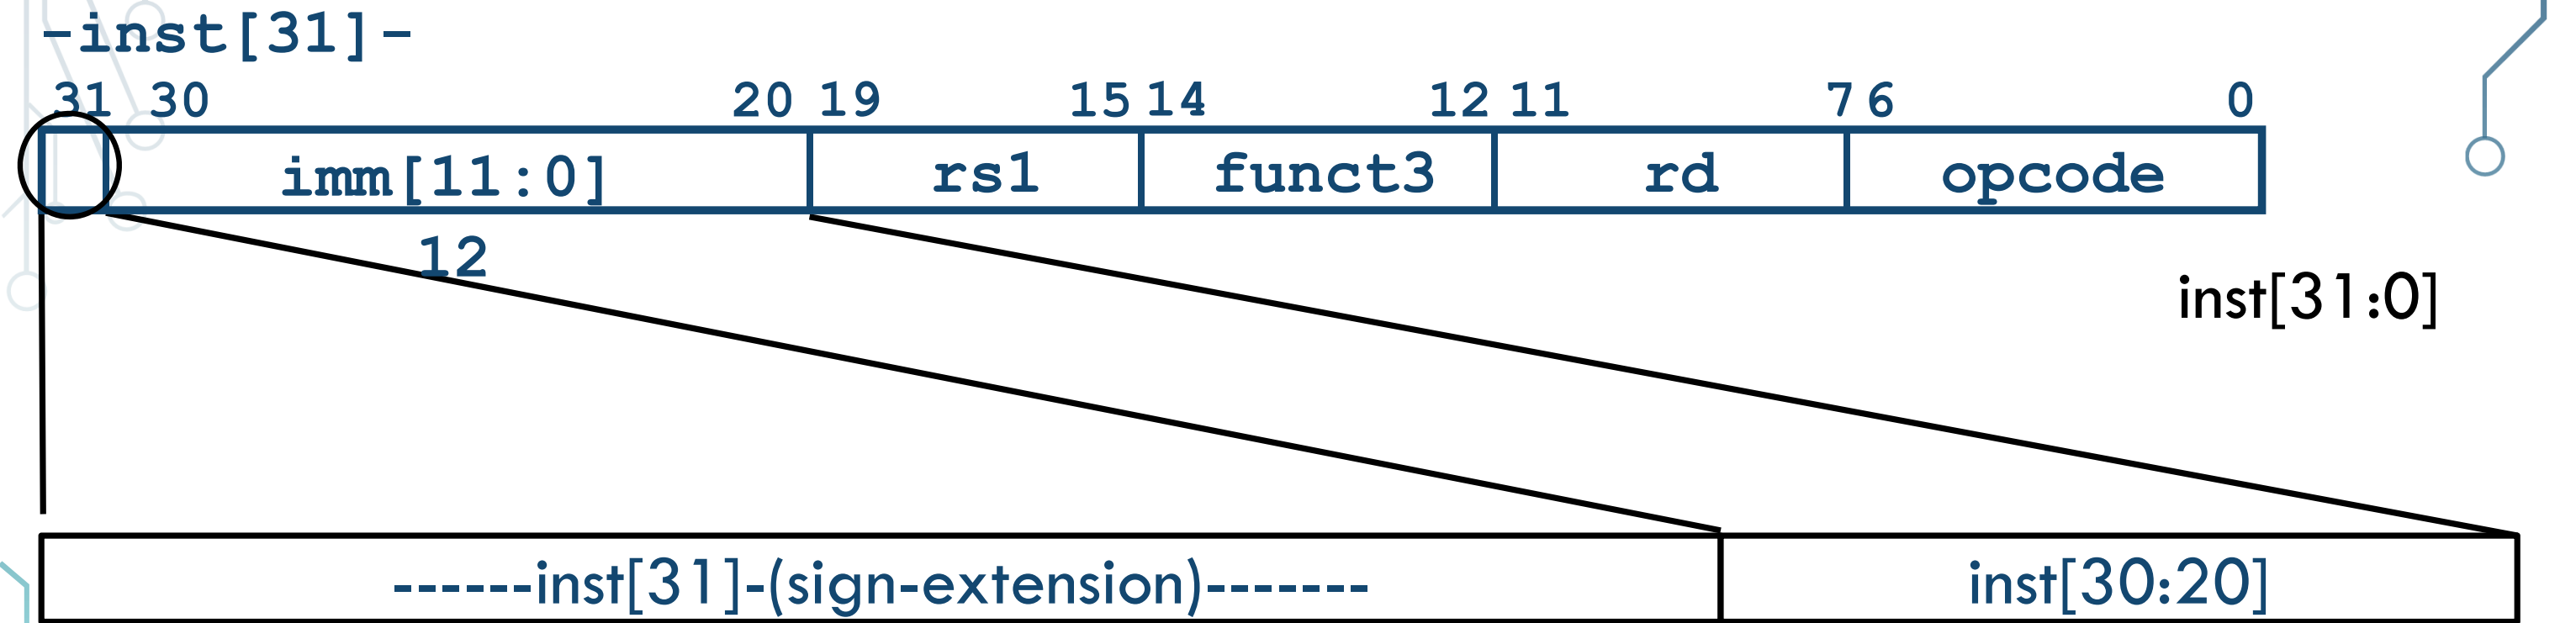
Adding addi to Datapath



Adding addi to Datapath

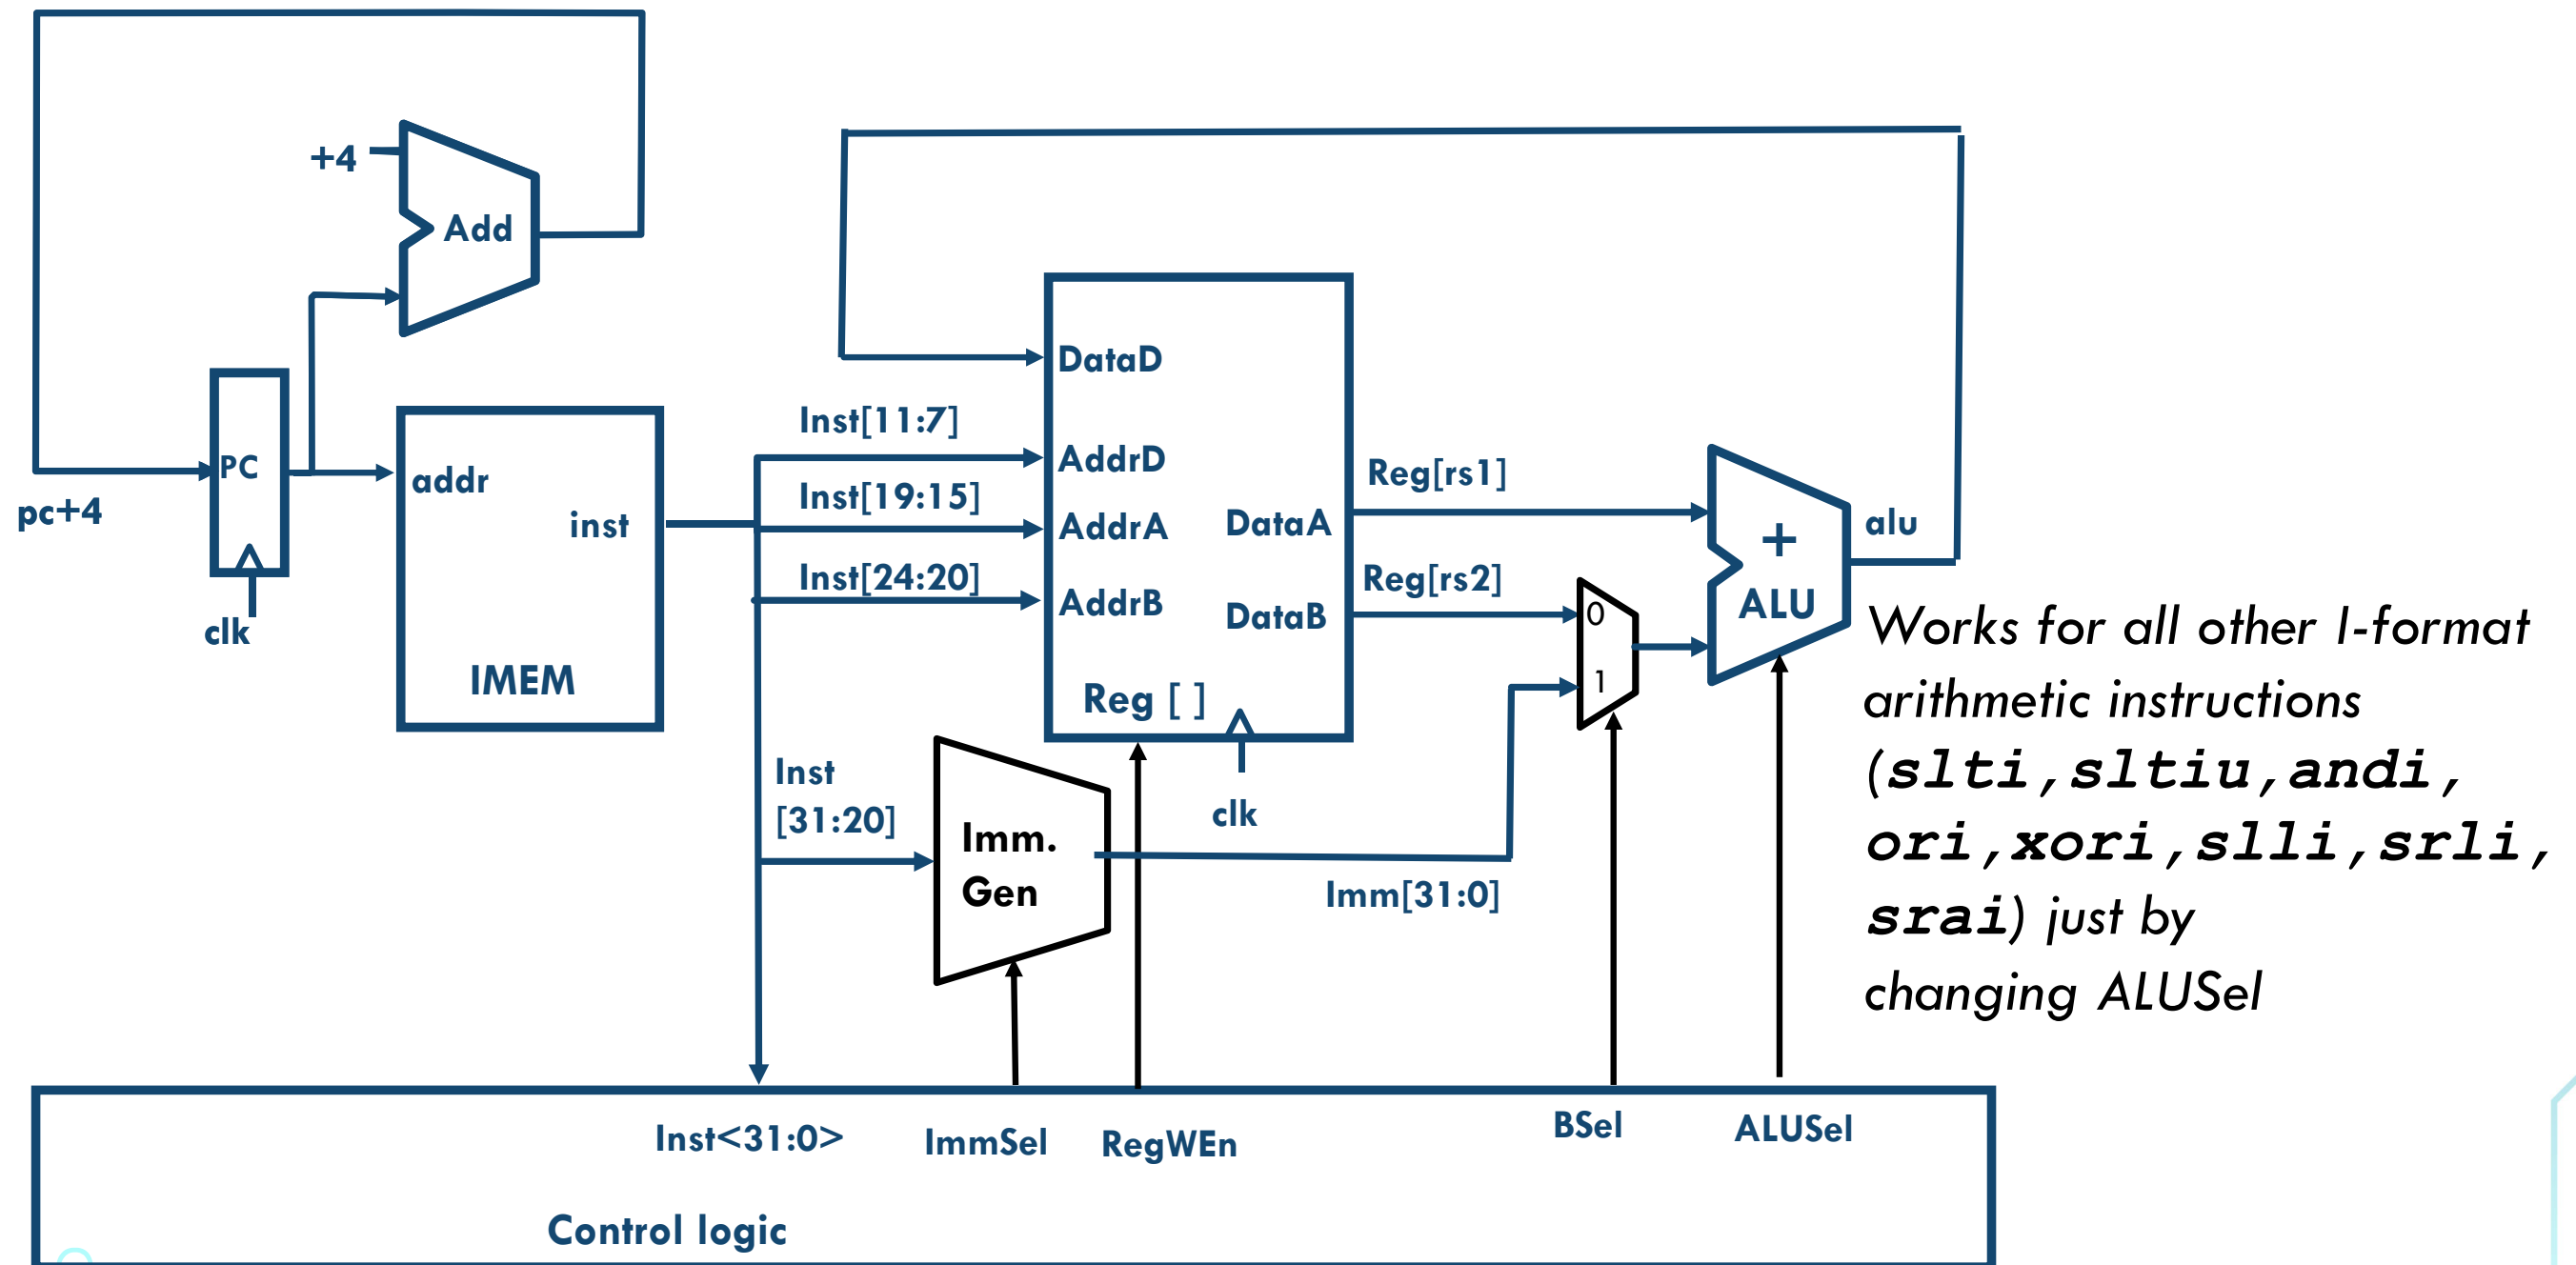


I-Format immediates

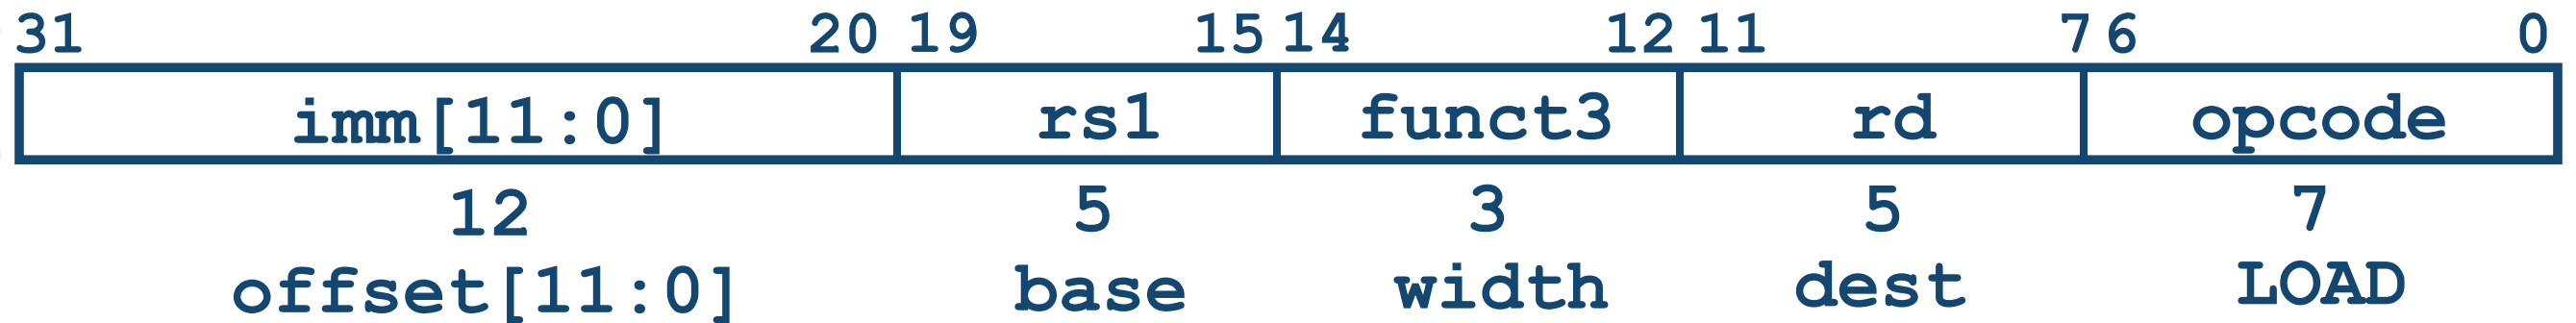


- **High 12 bits of instruction (inst[31:20]) copied to low 12 bits of immediate (imm[11:0])**
- **Immediate is sign-extended by copying value of inst[31] to fill the upper 20 bits of the immediate value (imm[31:12])**
- **Sign extension often in critical path**

R+I Datapath



Load Instructions are also I-Type

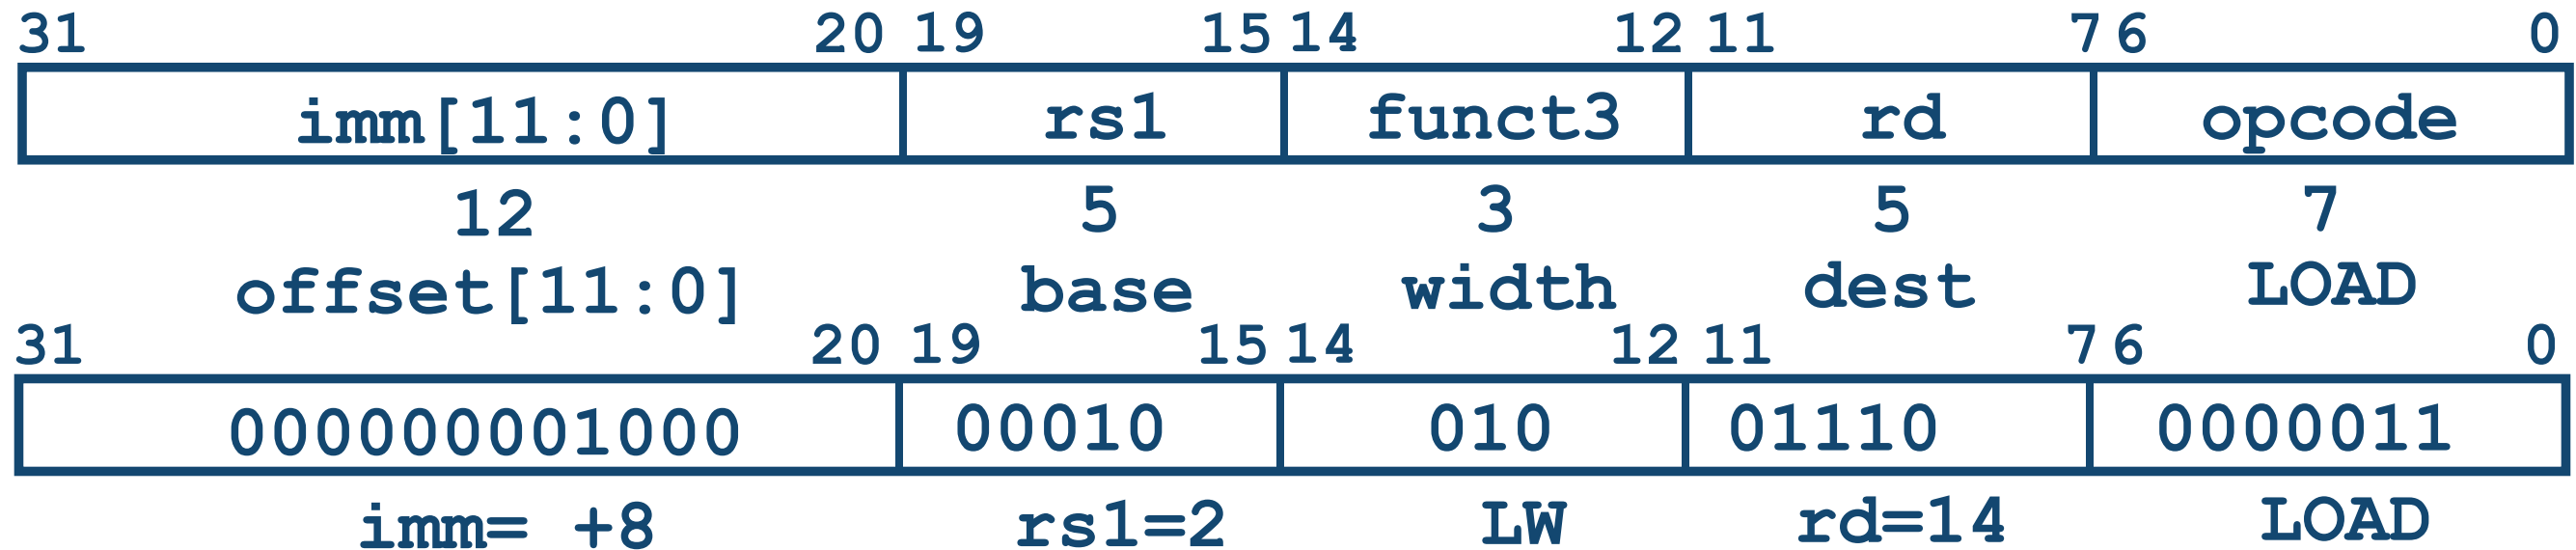


- The 12-bit signed immediate is added to the base address in register **rs1** to form the memory address
 - This is very similar to the add-immediate operation but used to create address not to create final result
- The value loaded from memory is stored in register rd

Add lw to Datapath

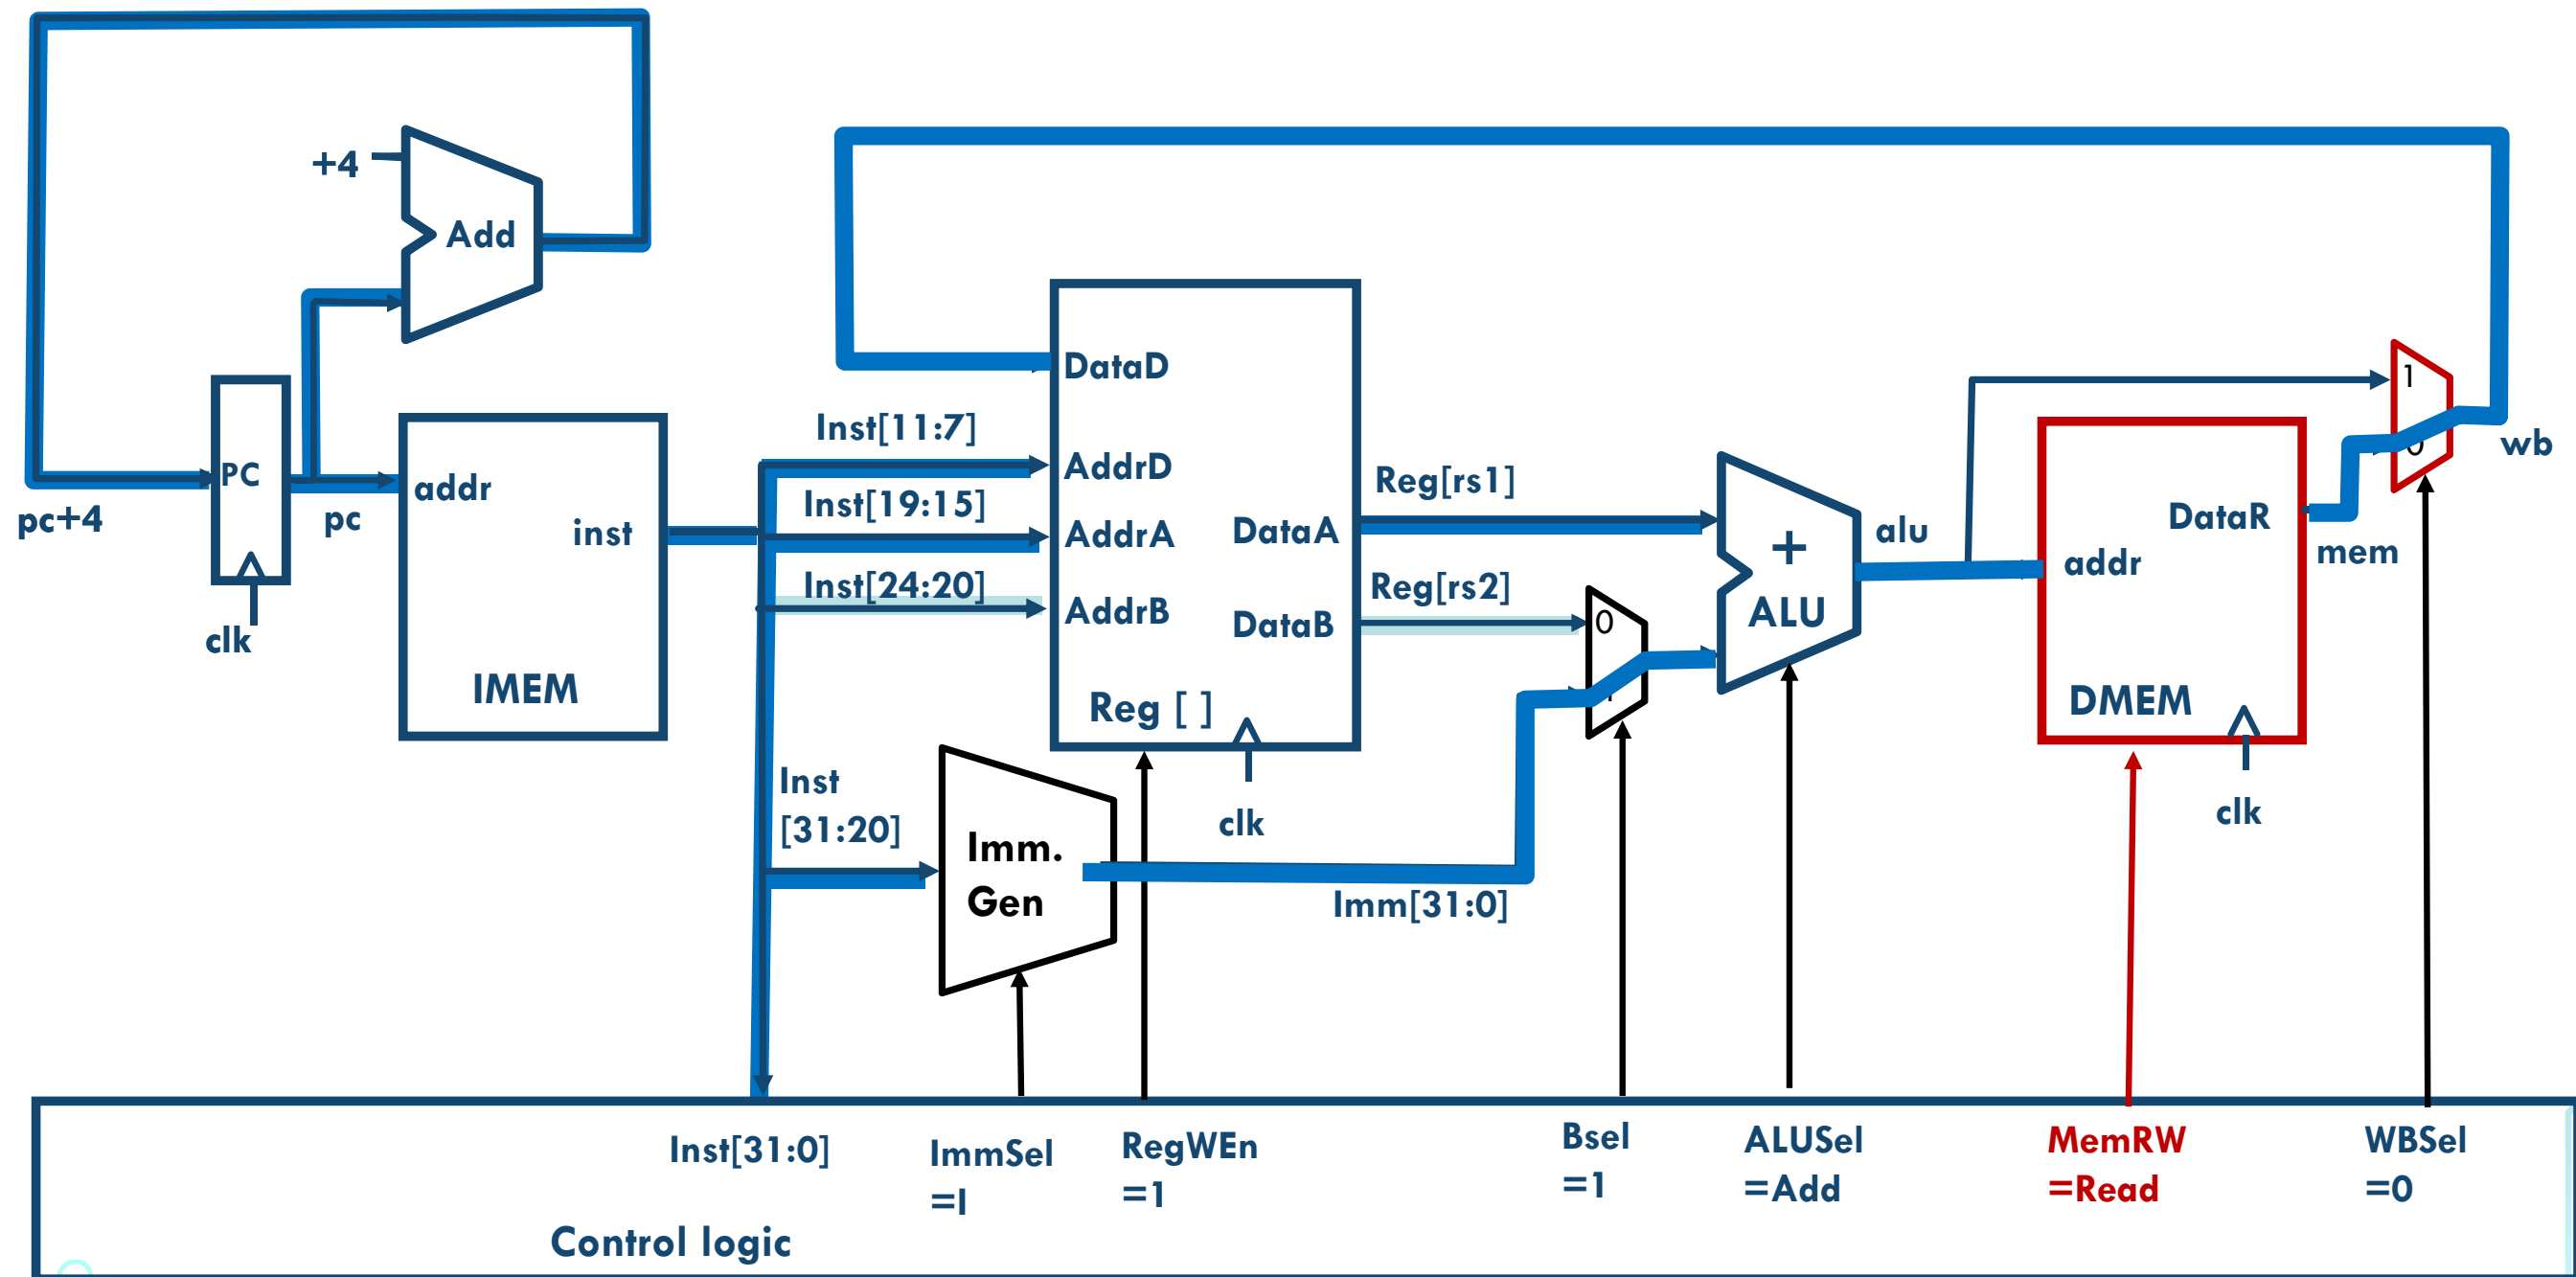
- RISC-V Assembly Instruction (I-type):

lw x14, 8(x2)



- The 12-bit signed immediate is added to the base address in register **rs1** to form the **memory** address
 - This is very similar to the add-immediate operation but used to create address not to create final result
- The value loaded from **memory** is stored in register **rd**

Adding lw to Datapath



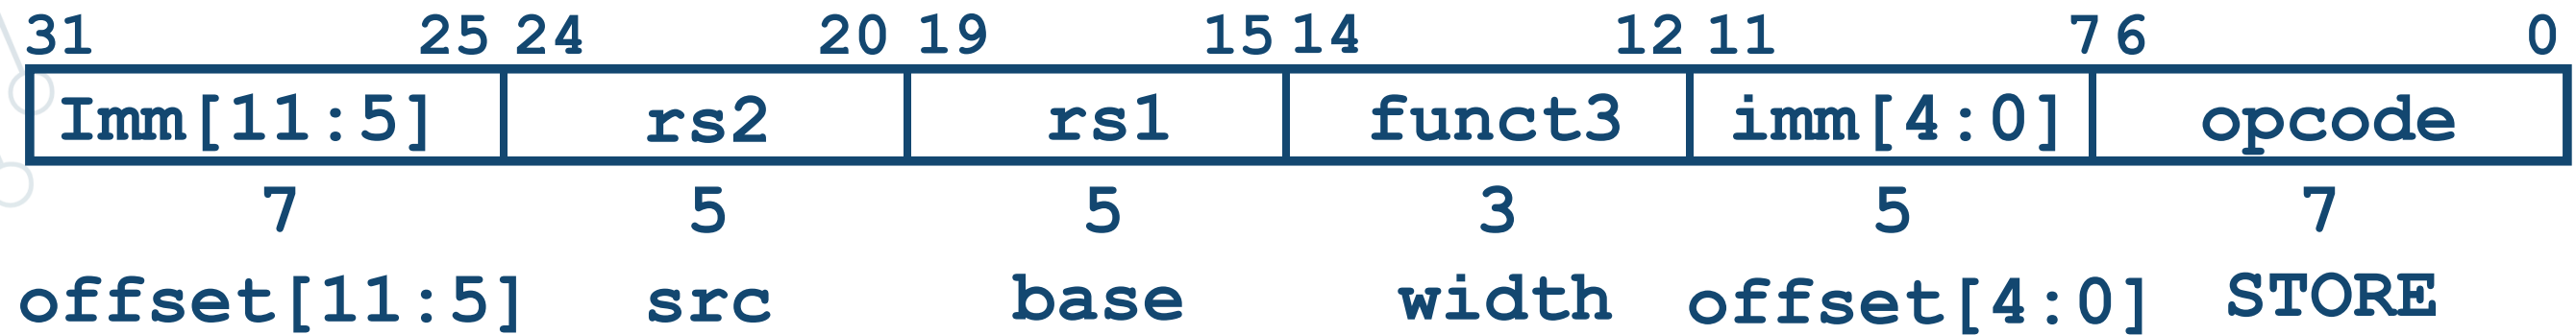
All RV32 Load Instructions

<code>imm[11:0]</code>	<code>rs1</code>	<code>000</code>	<code>rd</code>	<code>0000011</code>	<code>lb</code>
<code>imm[11:0]</code>	<code>rs1</code>	<code>001</code>	<code>rd</code>	<code>0000011</code>	<code>lh</code>
<code>imm[11:0]</code>	<code>rs1</code>	<code>010</code>	<code>rd</code>	<code>0000011</code>	<code>lw</code>
<code>imm[11:0]</code>	<code>rs1</code>	<code>100</code>	<code>rd</code>	<code>0000011</code>	<code>lbu</code>
<code>imm[11:0]</code>	<code>rs1</code>	<code>101</code>	<code>rd</code>	<code>0000011</code>	<code>lhu</code>

funct3 field encodes size and
'signedness' of load data

- Supporting the narrower loads requires additional logic to extract the correct byte/halfword from the value loaded from memory, and sign- or zero-extend the result to 32 bits before writing back to register file.
 - It is just a mux for load extend, similar to sign extension for immediates

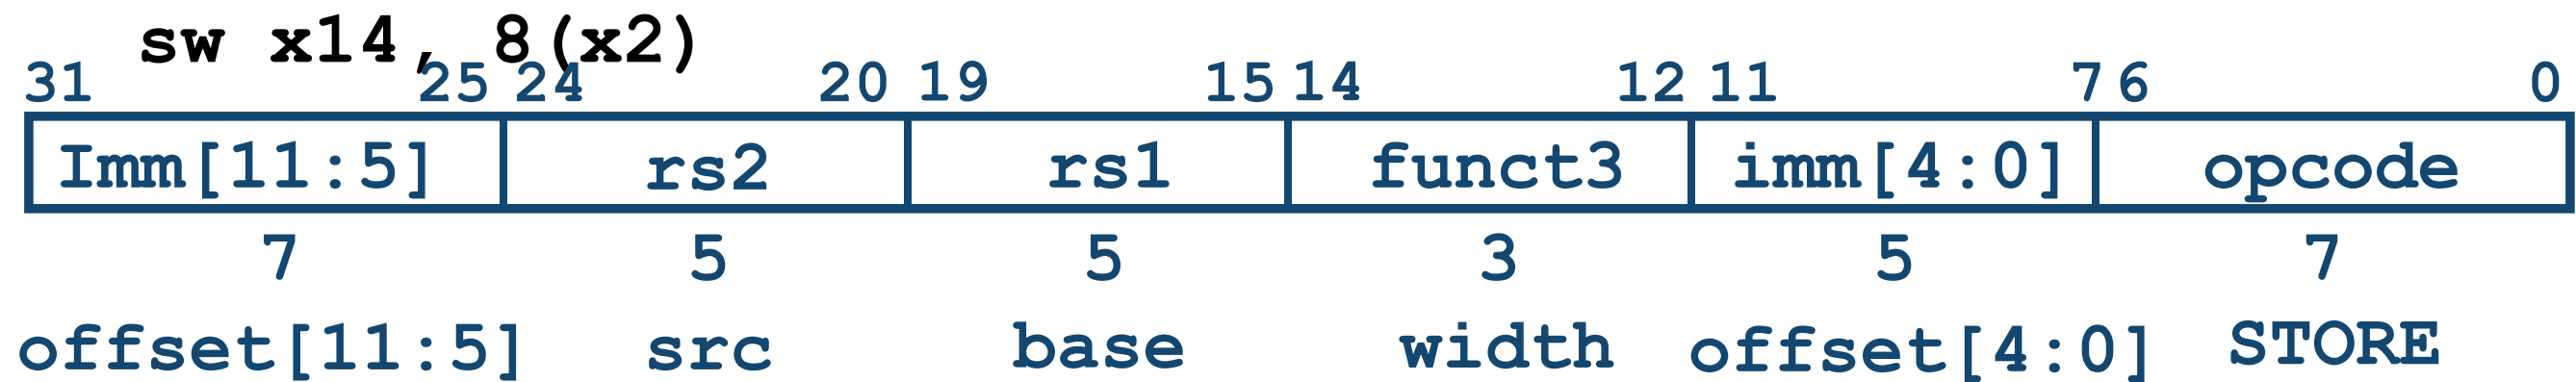
S-Format Used for Stores



- Store needs to read two registers, rs1 for base memory address, and rs2 for data to be stored, as well immediate offset!
- Can't have both rs2 and immediate in same place as other instructions!
- Note that stores don't write a value to the register file, **no rd!**
- RISC-V design decision is move low 5 bits of immediate to where rd field was in other instructions – keep rs1/rs2 fields in same place
 - register names more critical than immediate bits in hardware design

Adding SW Instruction

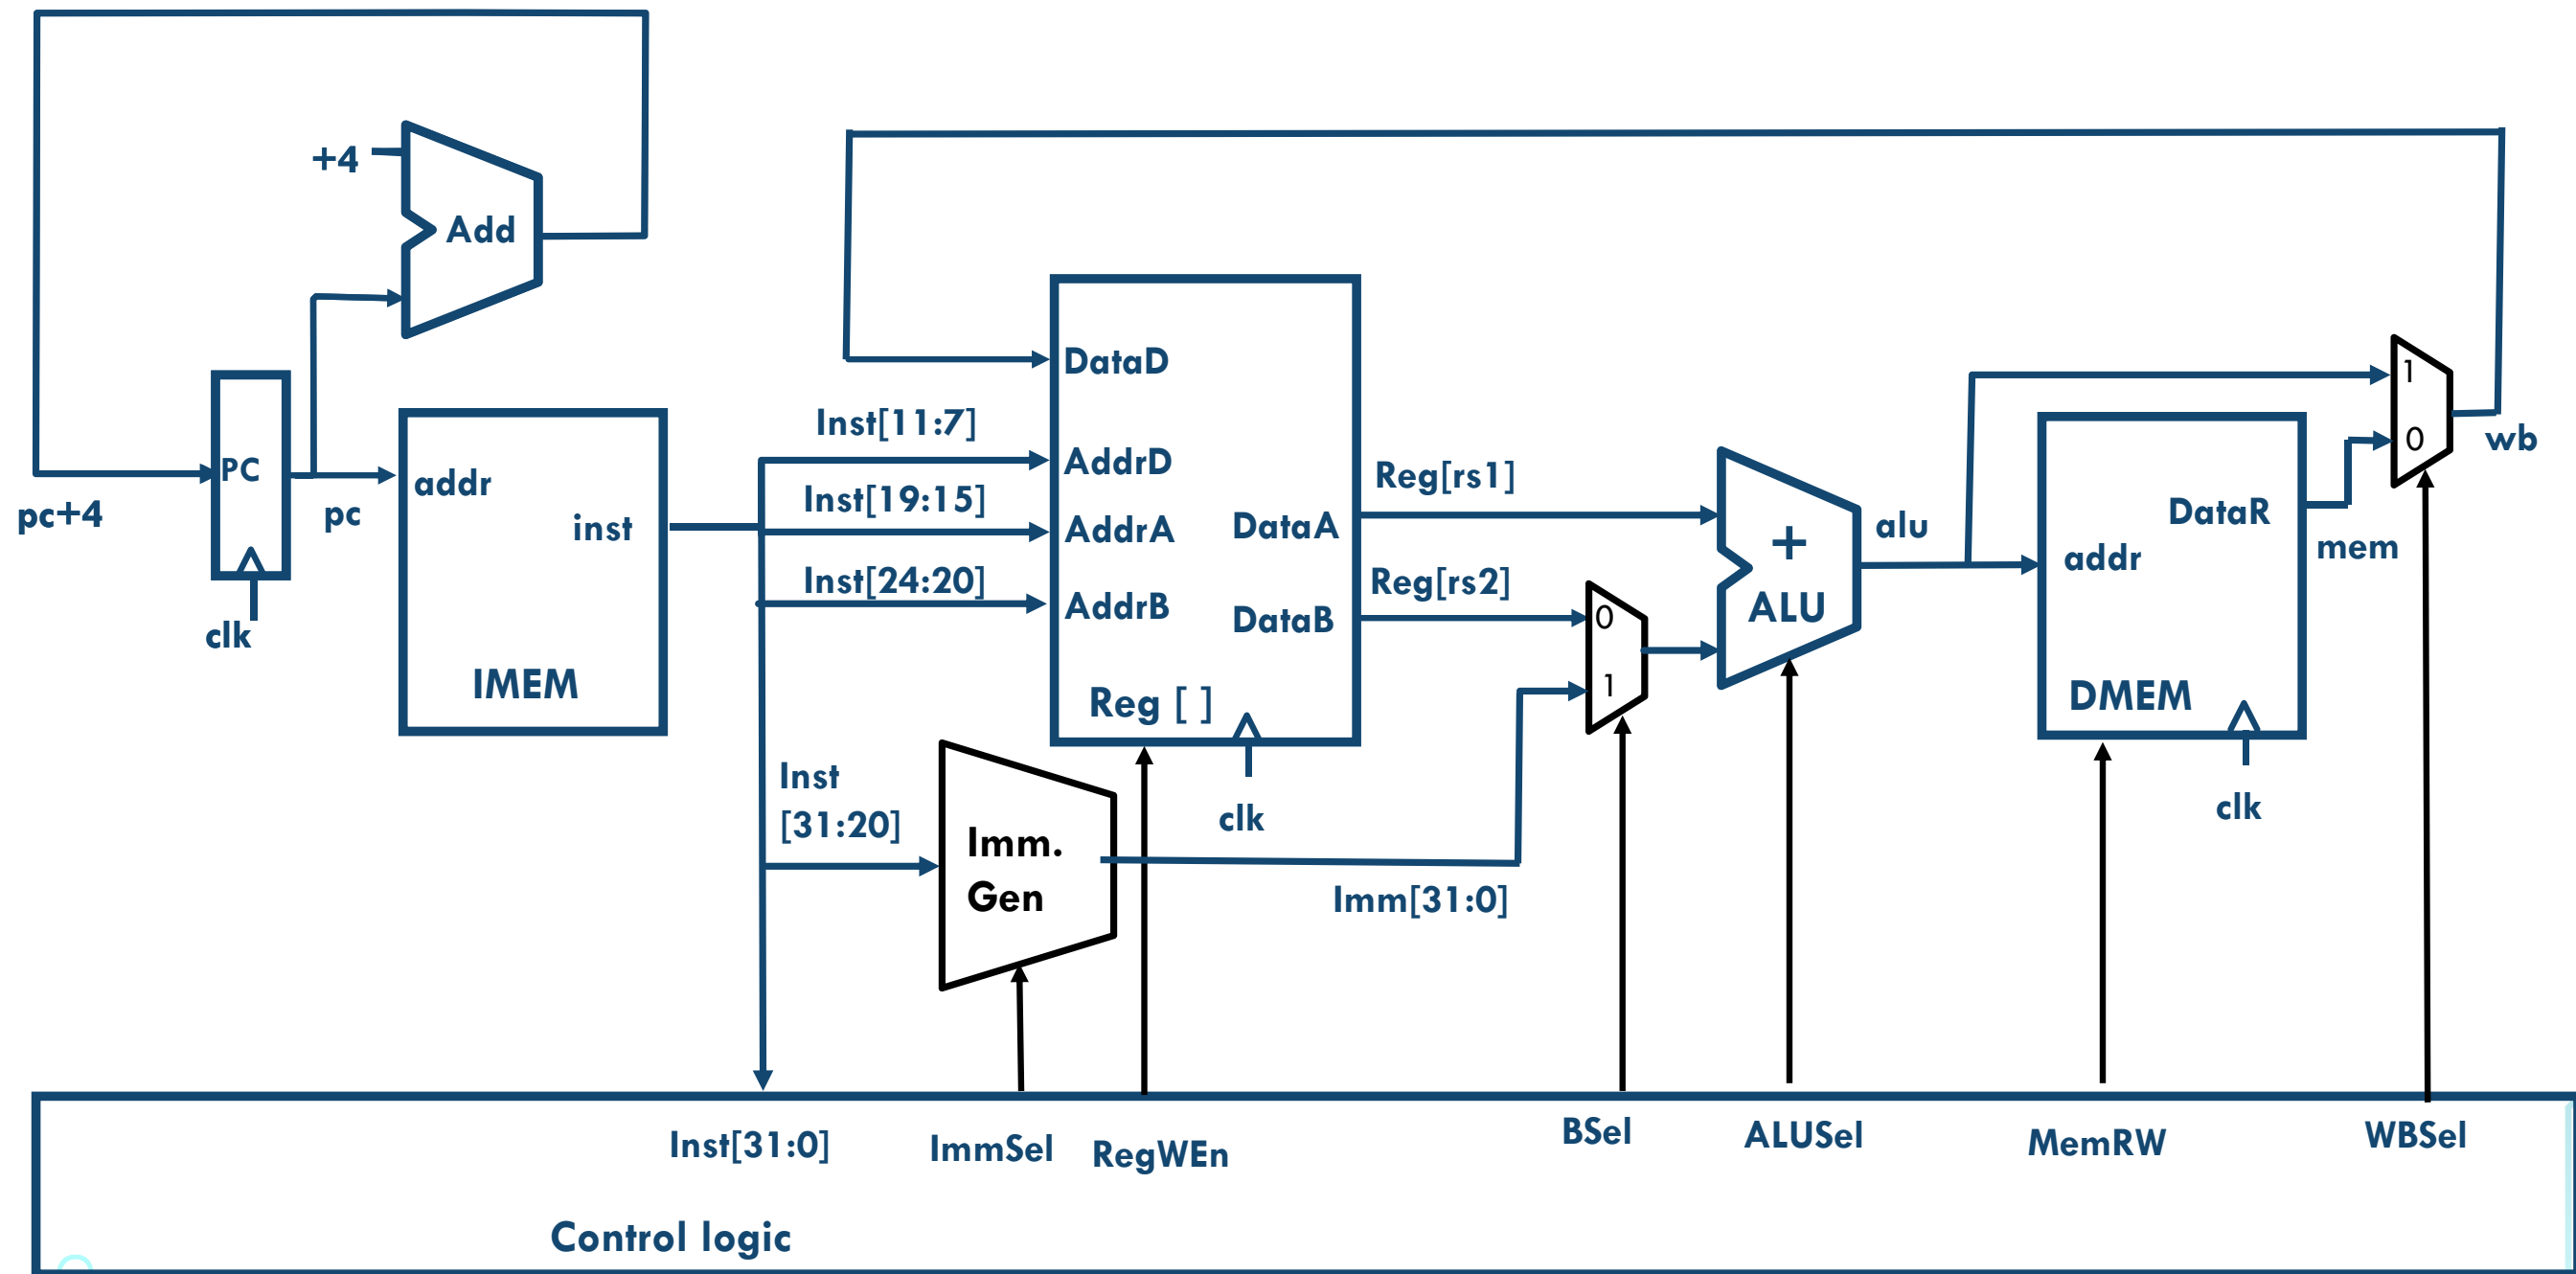
- sw: Reads two registers, rs1 for base memory address, and rs2 for data to be stored, as well immediate offset!



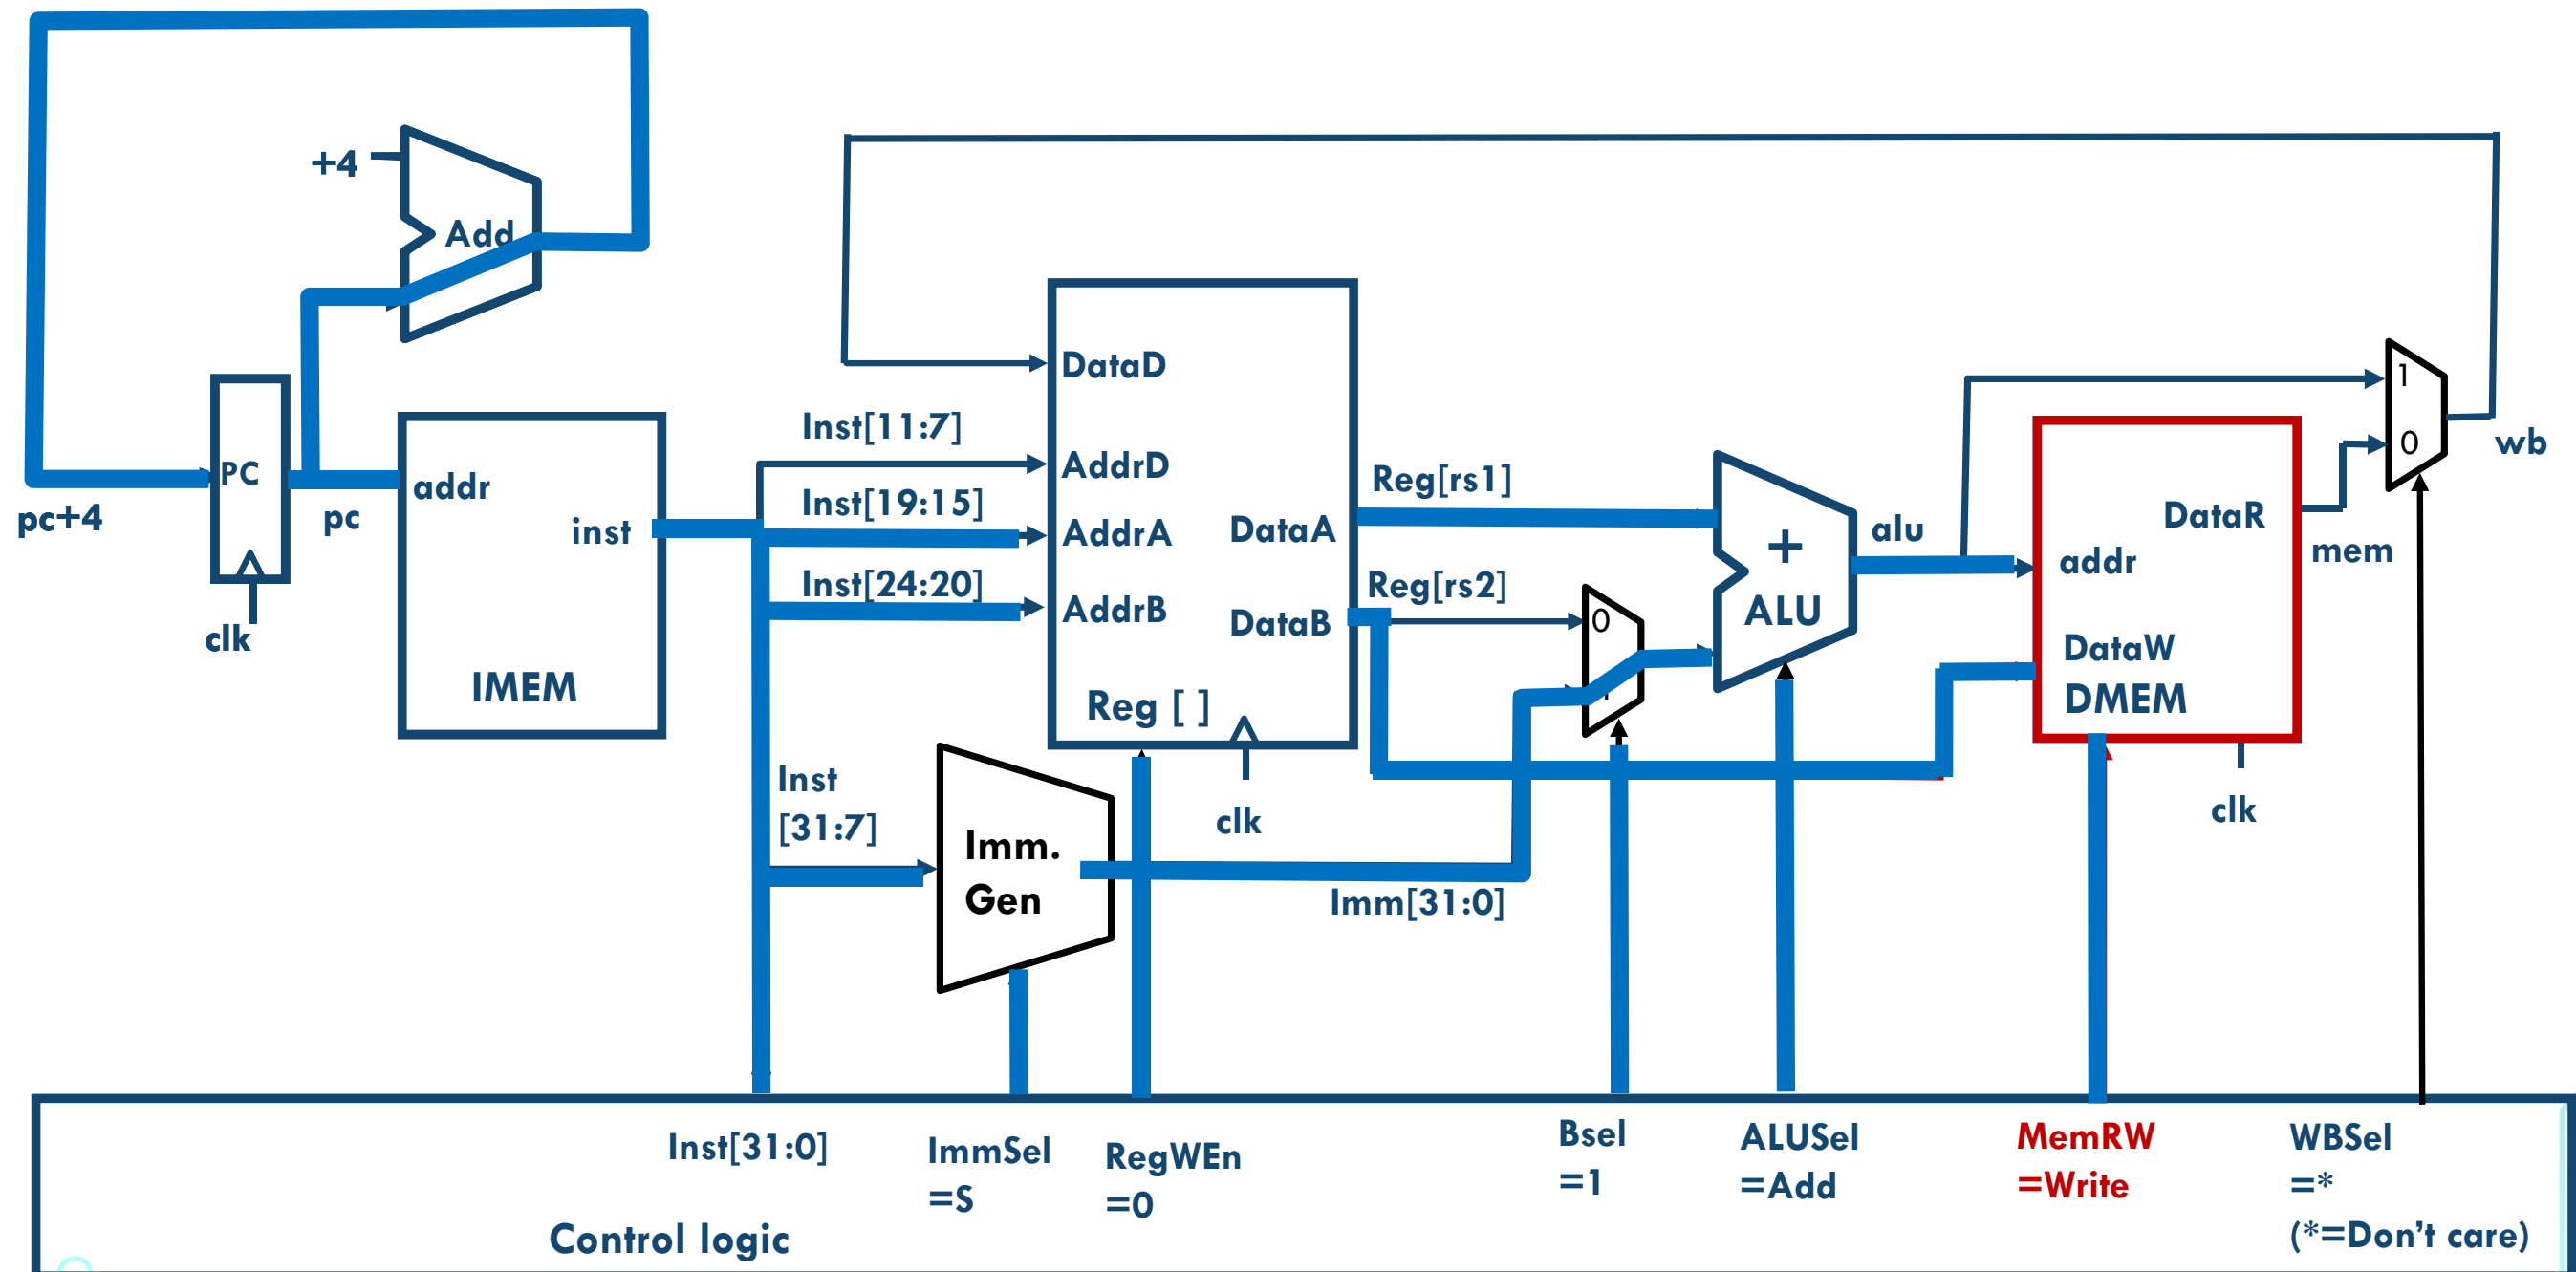
offset[11:5] rs2=14 rs1=2 SW offset[4:0] STORE
=0 =8



Datapath with 1w



Adding SW to Datapath



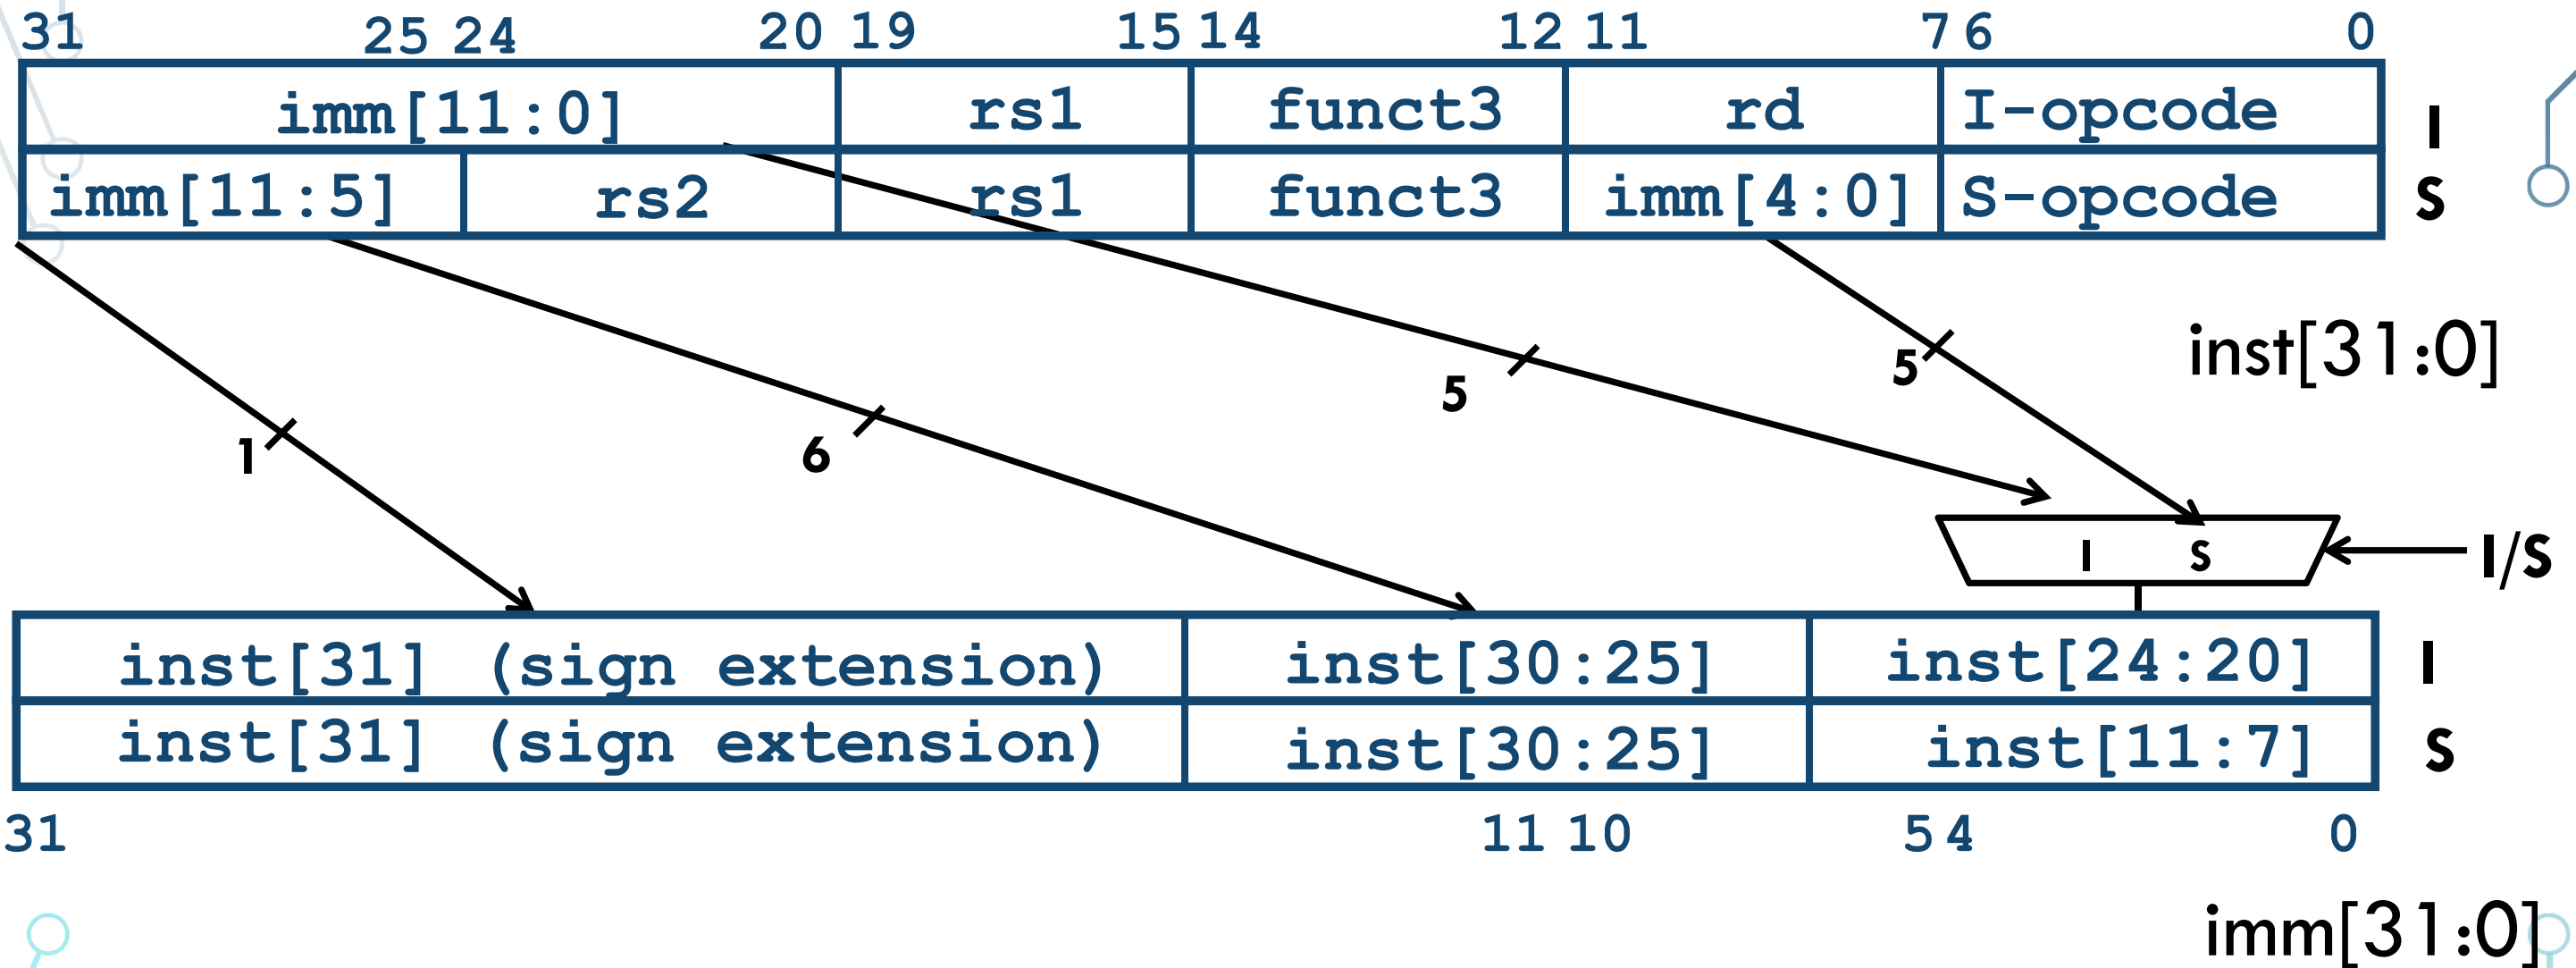
All RV32 Store Instructions

Imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	sb
Imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	sh
Imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	sw

width

- Store byte, halfword, word

I+S Immediate Generation



- Just need a 5-bit mux to select between two positions where low five bits of immediate can reside in instruction
- Other bits in immediate are wired to fixed positions in instruction

Administrivia

- Midterm 1: October 2, in class
 - Covers material up to today
 - One double-sided page of hand-written notes
 - Review session on Monday
- Midterm 2: November 6

B-Format - RISC-V Conditional Branches

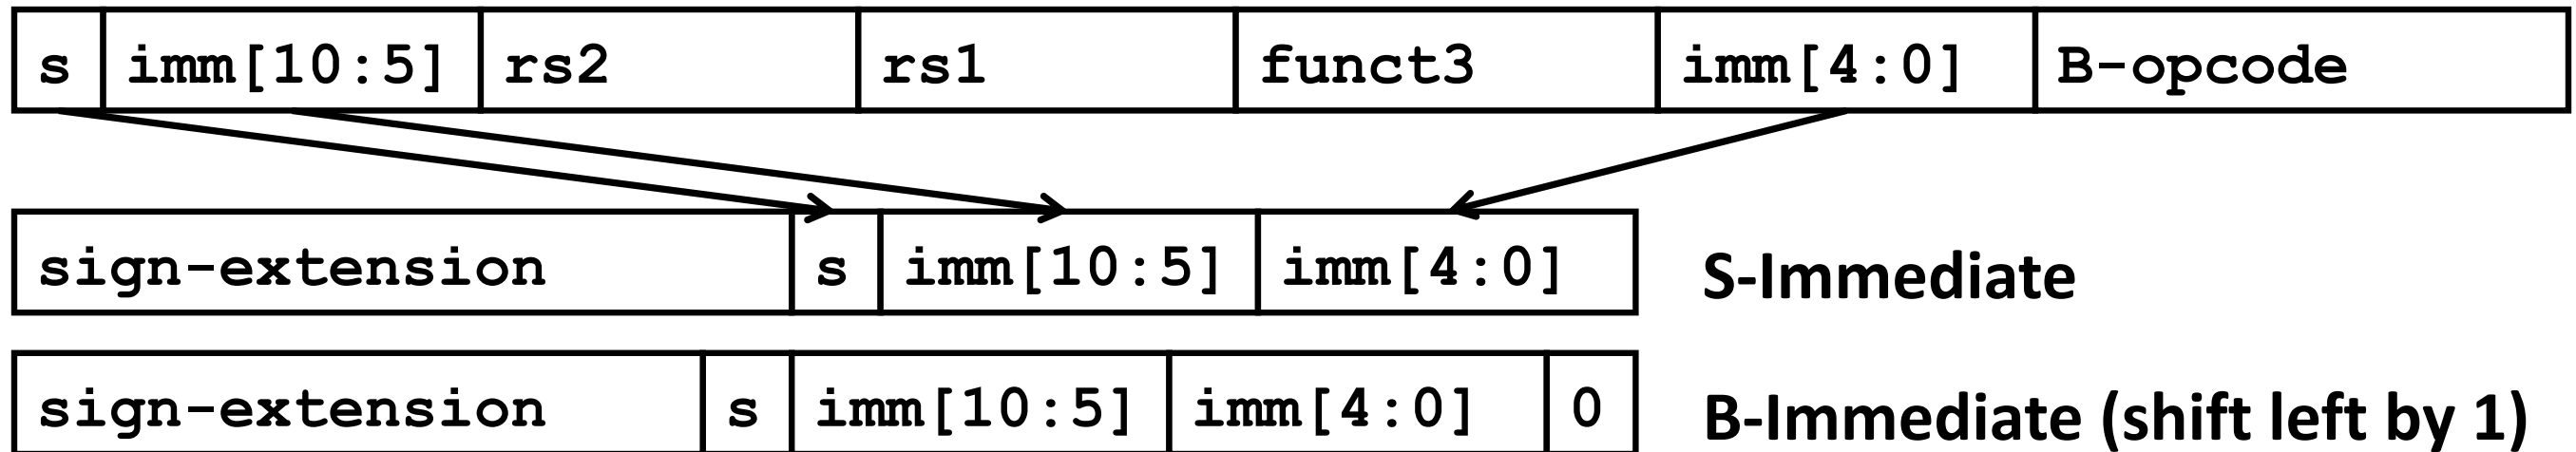
- E.g., `BEQ x1, x2, Label`
- Branches read two registers but don't write a register (similar to stores)
- How to encode label, i.e., where to branch to?

RISC-V Feature, $n \times 16$ -bit instructions

- Extensions to RISC-V base ISA support 16-bit compressed instructions and also variable-length instructions that are multiples of 16-bits in length
- To enable this, RISC-V scales the branch offset by 2 bytes even when there are no 16-bit instructions
- Reduces branch reach by half and means that $\frac{1}{2}$ of possible targets will be errors on RISC-V processors that only support 32-bit instructions (as used in this class)
- RISC-V conditional branches can only reach $\pm 2^{10} \times 32$ -bit instructions on either side of PC

RISC-V Branch Immediates

- 12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes
- RISC-V approach: keep 11 immediate bits in fixed position in output value, and rotate LSB of S-format to be bit 12 of B-format



Only one bit changes position between S and B, so only need a single-bit 2-way mux

RISC-V Immediate Encoding

Instruction encodings, inst[31:0]

31	30	25	24	20	19	15	14	12	11	8	7	6	0																	
funct7					rs2					rs1					funct3					rd					opcode					R-type
imm[11:0]										rs1					funct3					rd					opcode					I-type
imm[11:5]					rs2					rs1					funct3					imm[4:0]					opcode					S-type
imm[12 10:5]					rs2					rs1					funct3					imm[4:1 11]					opcode					B-type

32-bit immediates produced, imm[31:0]

31	25	24	12	11	10	5	4	1	0			
-inst[31]-					inst[30:25]		inst[24:21]		inst[20]		I-imm.	
-inst[31]-					inst[30:25]		inst[11:8]		inst[7]		S-imm.	
-inst[31]-				inst[7]		inst[30:25]		inst[11:8]		0		B-imm.

Upper bits sign-extended from inst[31]
always

Only bit 7 of instruction changes role in
immediate between S and B

Branch Example, complete encoding

beq **x19,x10,** **offset = 16 bytes**

13-bit immediate, imm[12:0], with value 16

0000000010000

imm[0] discarded,
→ always zero

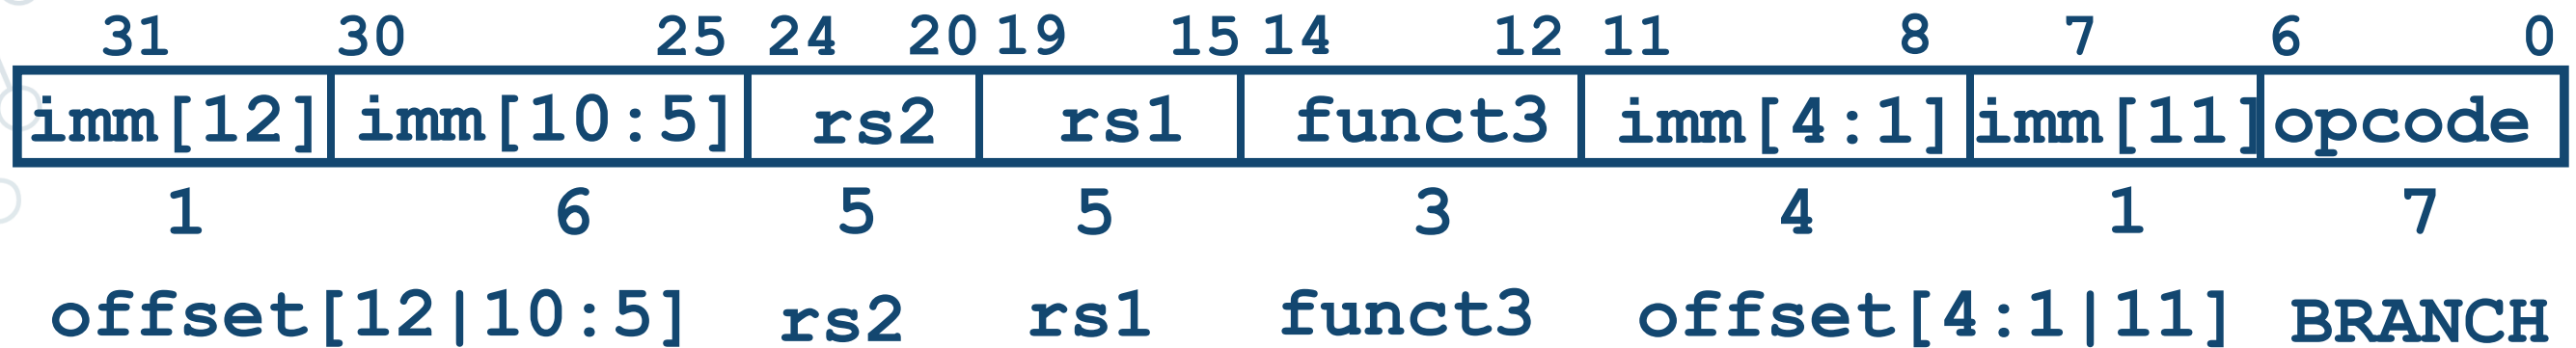
imm[12]

imm[11]

0	000000	01010	10011	000	1000	0	1100011
---	--------	-------	-------	-----	------	---	---------

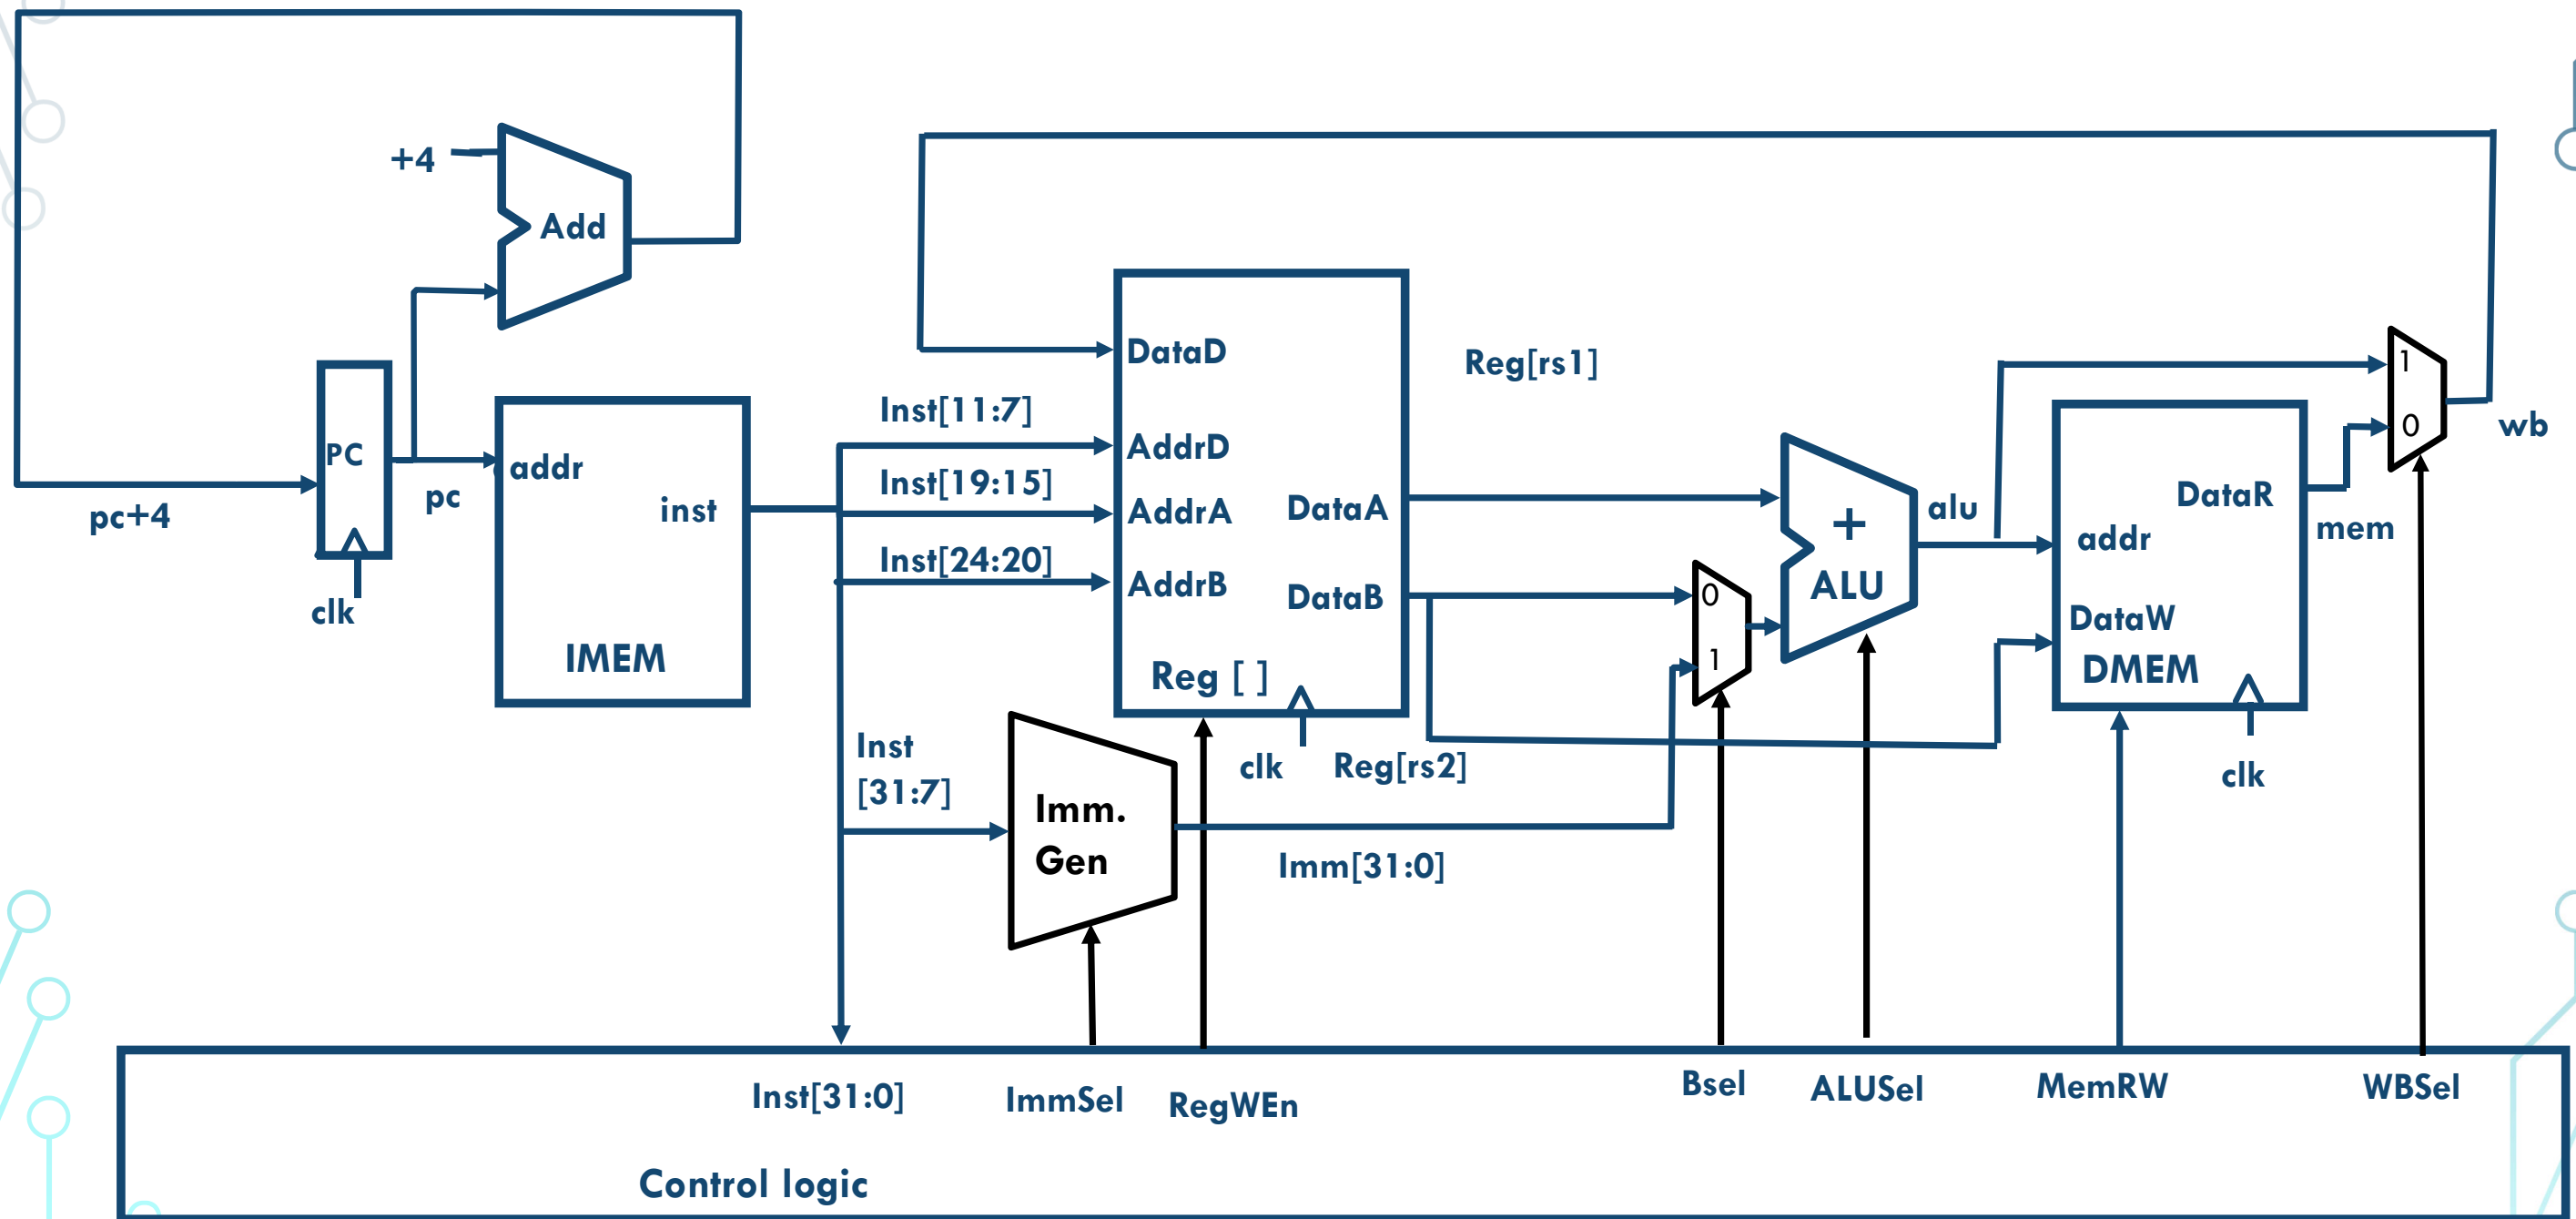
imm[10:5] rs2=10 rs1=19 BEQ imm[4:1] BRANCH

Implementing Branches



- B-format is mostly same as S-format, with two register sources (rs1/rs2) and a 12-bit immediate
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode even 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)

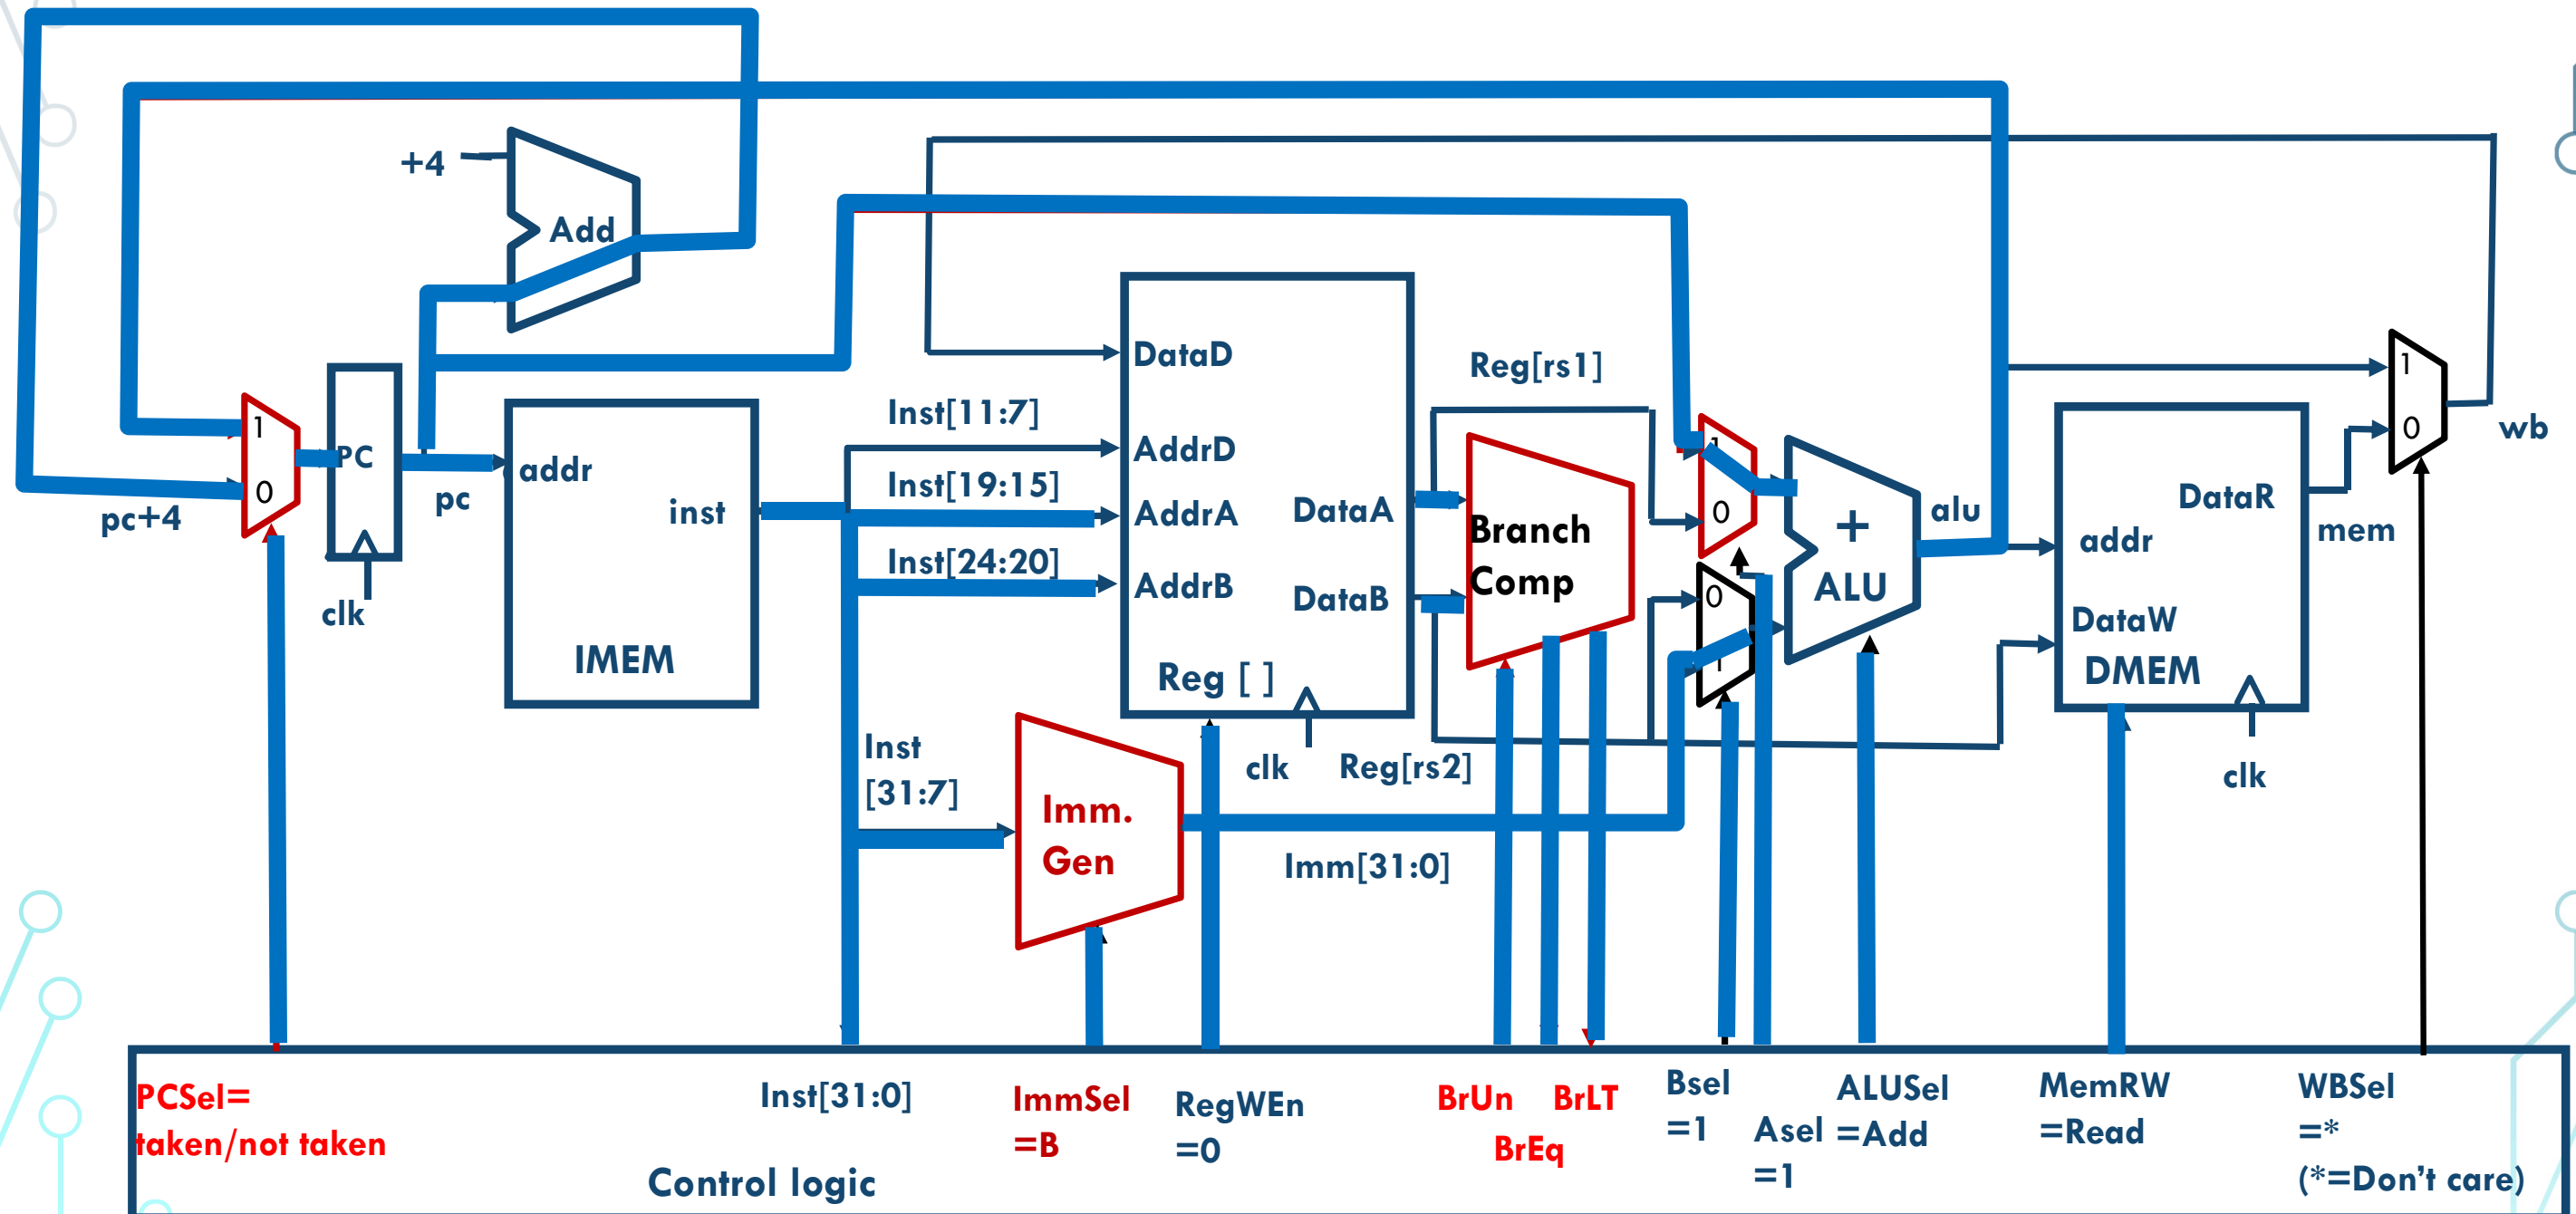
Datapath So Far



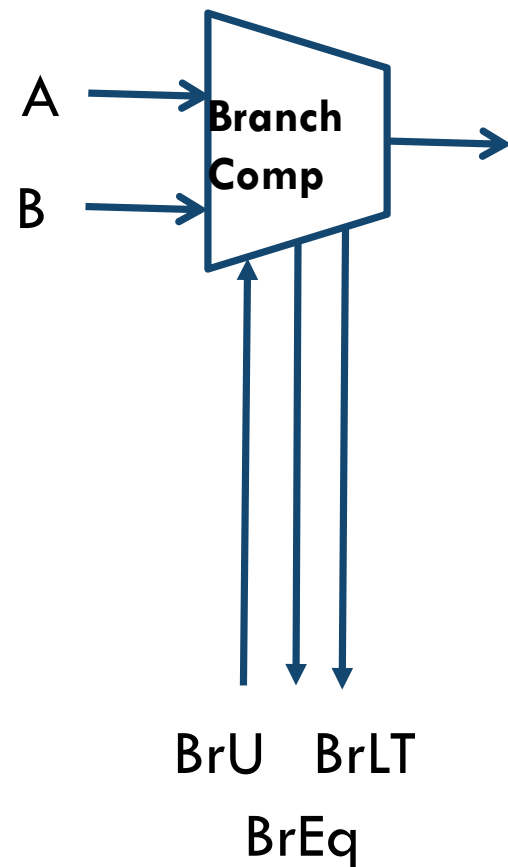
To Add Branches

- Different change to the state:
 - $$PC = \begin{cases} PC + 4, & \text{branch not taken} \\ PC + \text{immediate}, & \text{branch taken} \end{cases}$$
- Six branch instructions: **BEQ**, **BNE**, **BLT**, **BGE**, **BLTU**, **BGEU**
- Need to compute **PC + immediate** and to compare values of **rs1** and **rs2**
 - Need another add/sub unit

Adding Branches



Branch Comparator



- $\text{BrEq} = 1$, if $A = B$
- $\text{BrLT} = 1$, if $A < B$
- $\text{BrUn} = 1$ selects unsigned comparison for BrLT , 0=signed
- BGE branch: $A \geq B$, if $\overline{A < B}$

All RISC-V Branch Instructions

<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>000</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>001</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>100</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>101</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>110</code>	<code>imm[4:1 11]</code>	<code>1100011</code>
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>111</code>	<code>imm[4:1 11]</code>	<code>1100011</code>

BEQ

BNE

BLT

BGE

BLTU

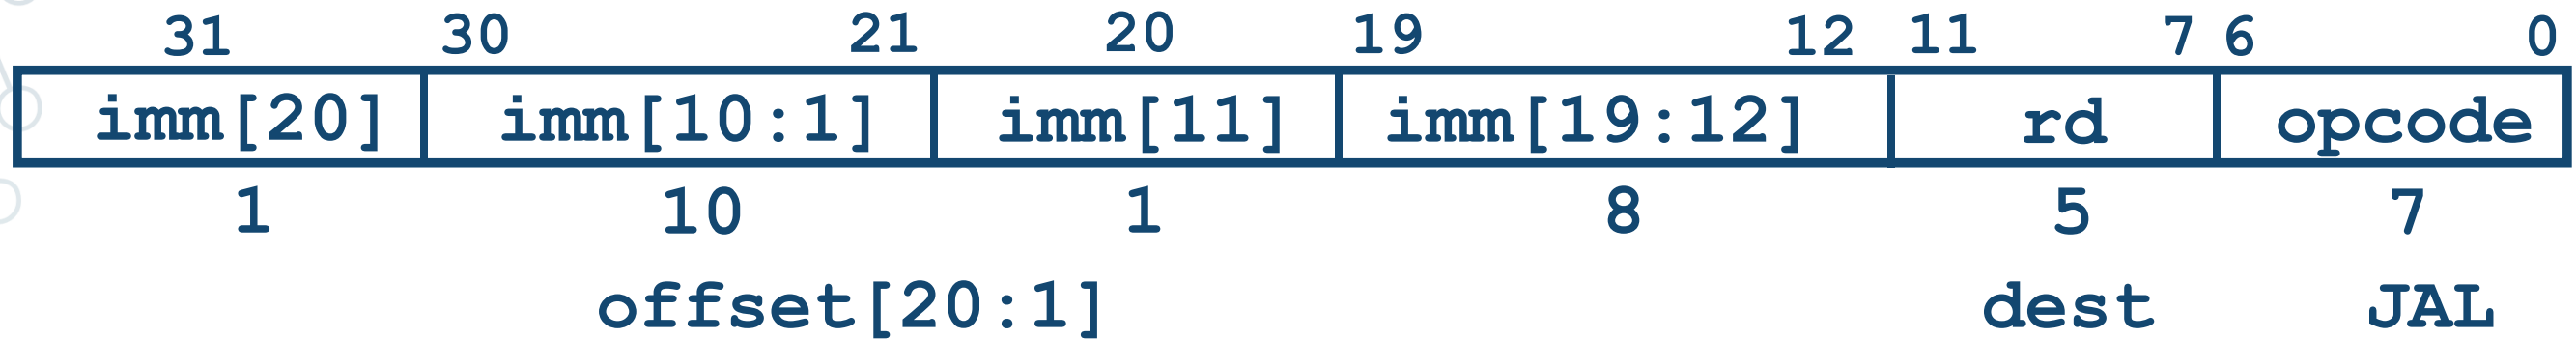
BGEU

Peer Instruction

- 1) We **should use the main ALU** to evaluate branches in order to save some gates in small cores
- 2) Program counter is a register
- 3) The **ALU** is a sequential element

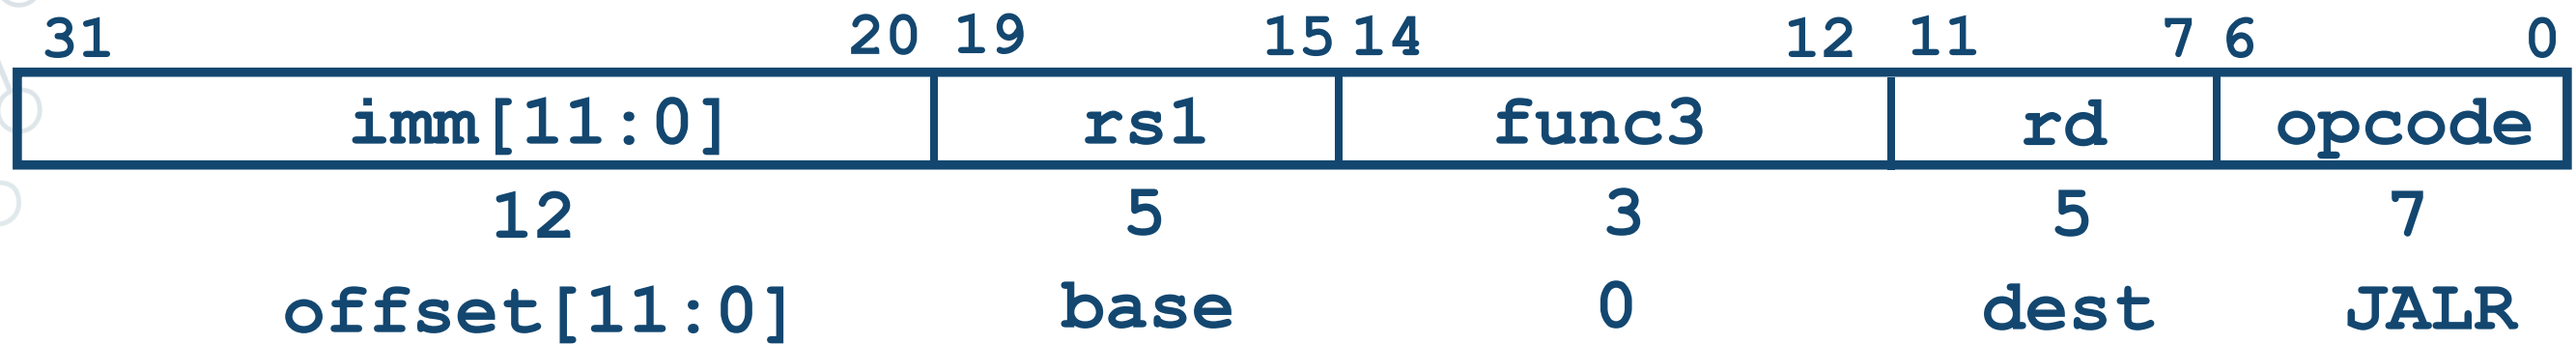
1	2	3
F	F	F
F	F	T
F	T	F
F	T	T
T	F	F
T	F	T
T	T	F
T	T	T

J-Format for Jump Instructions



- JAL saves PC+4 in register `rd` (the return address)
 - Assembler “`j`” jump is pseudo-instruction, uses JAL but sets `rd=x0` to discard return address
- Set $PC = PC + \text{offset}$ (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

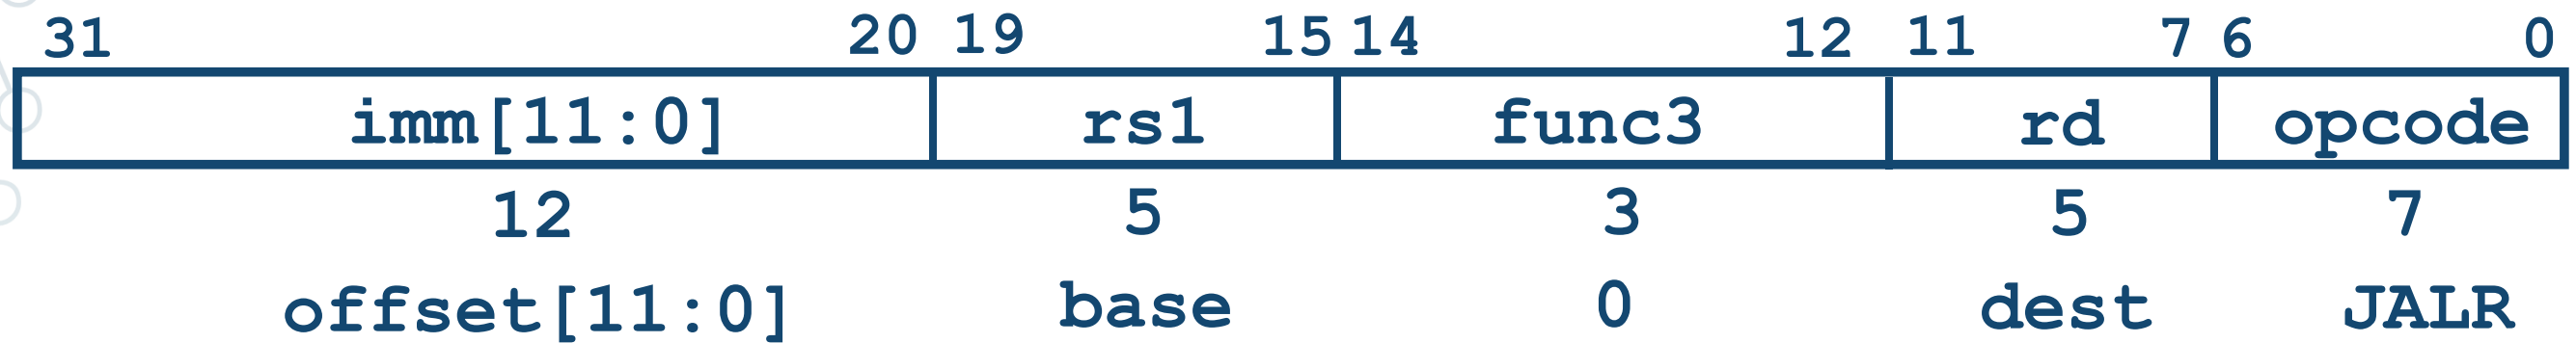
JALR Instruction (I-Format)



- JALR rd, rs, immediate
 - Writes PC+4 to rd (return address)
 - Sets $PC = rs1 + \text{immediate}$ (and sets the LSB to 0)
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes
 - In contrast to branches and JAL

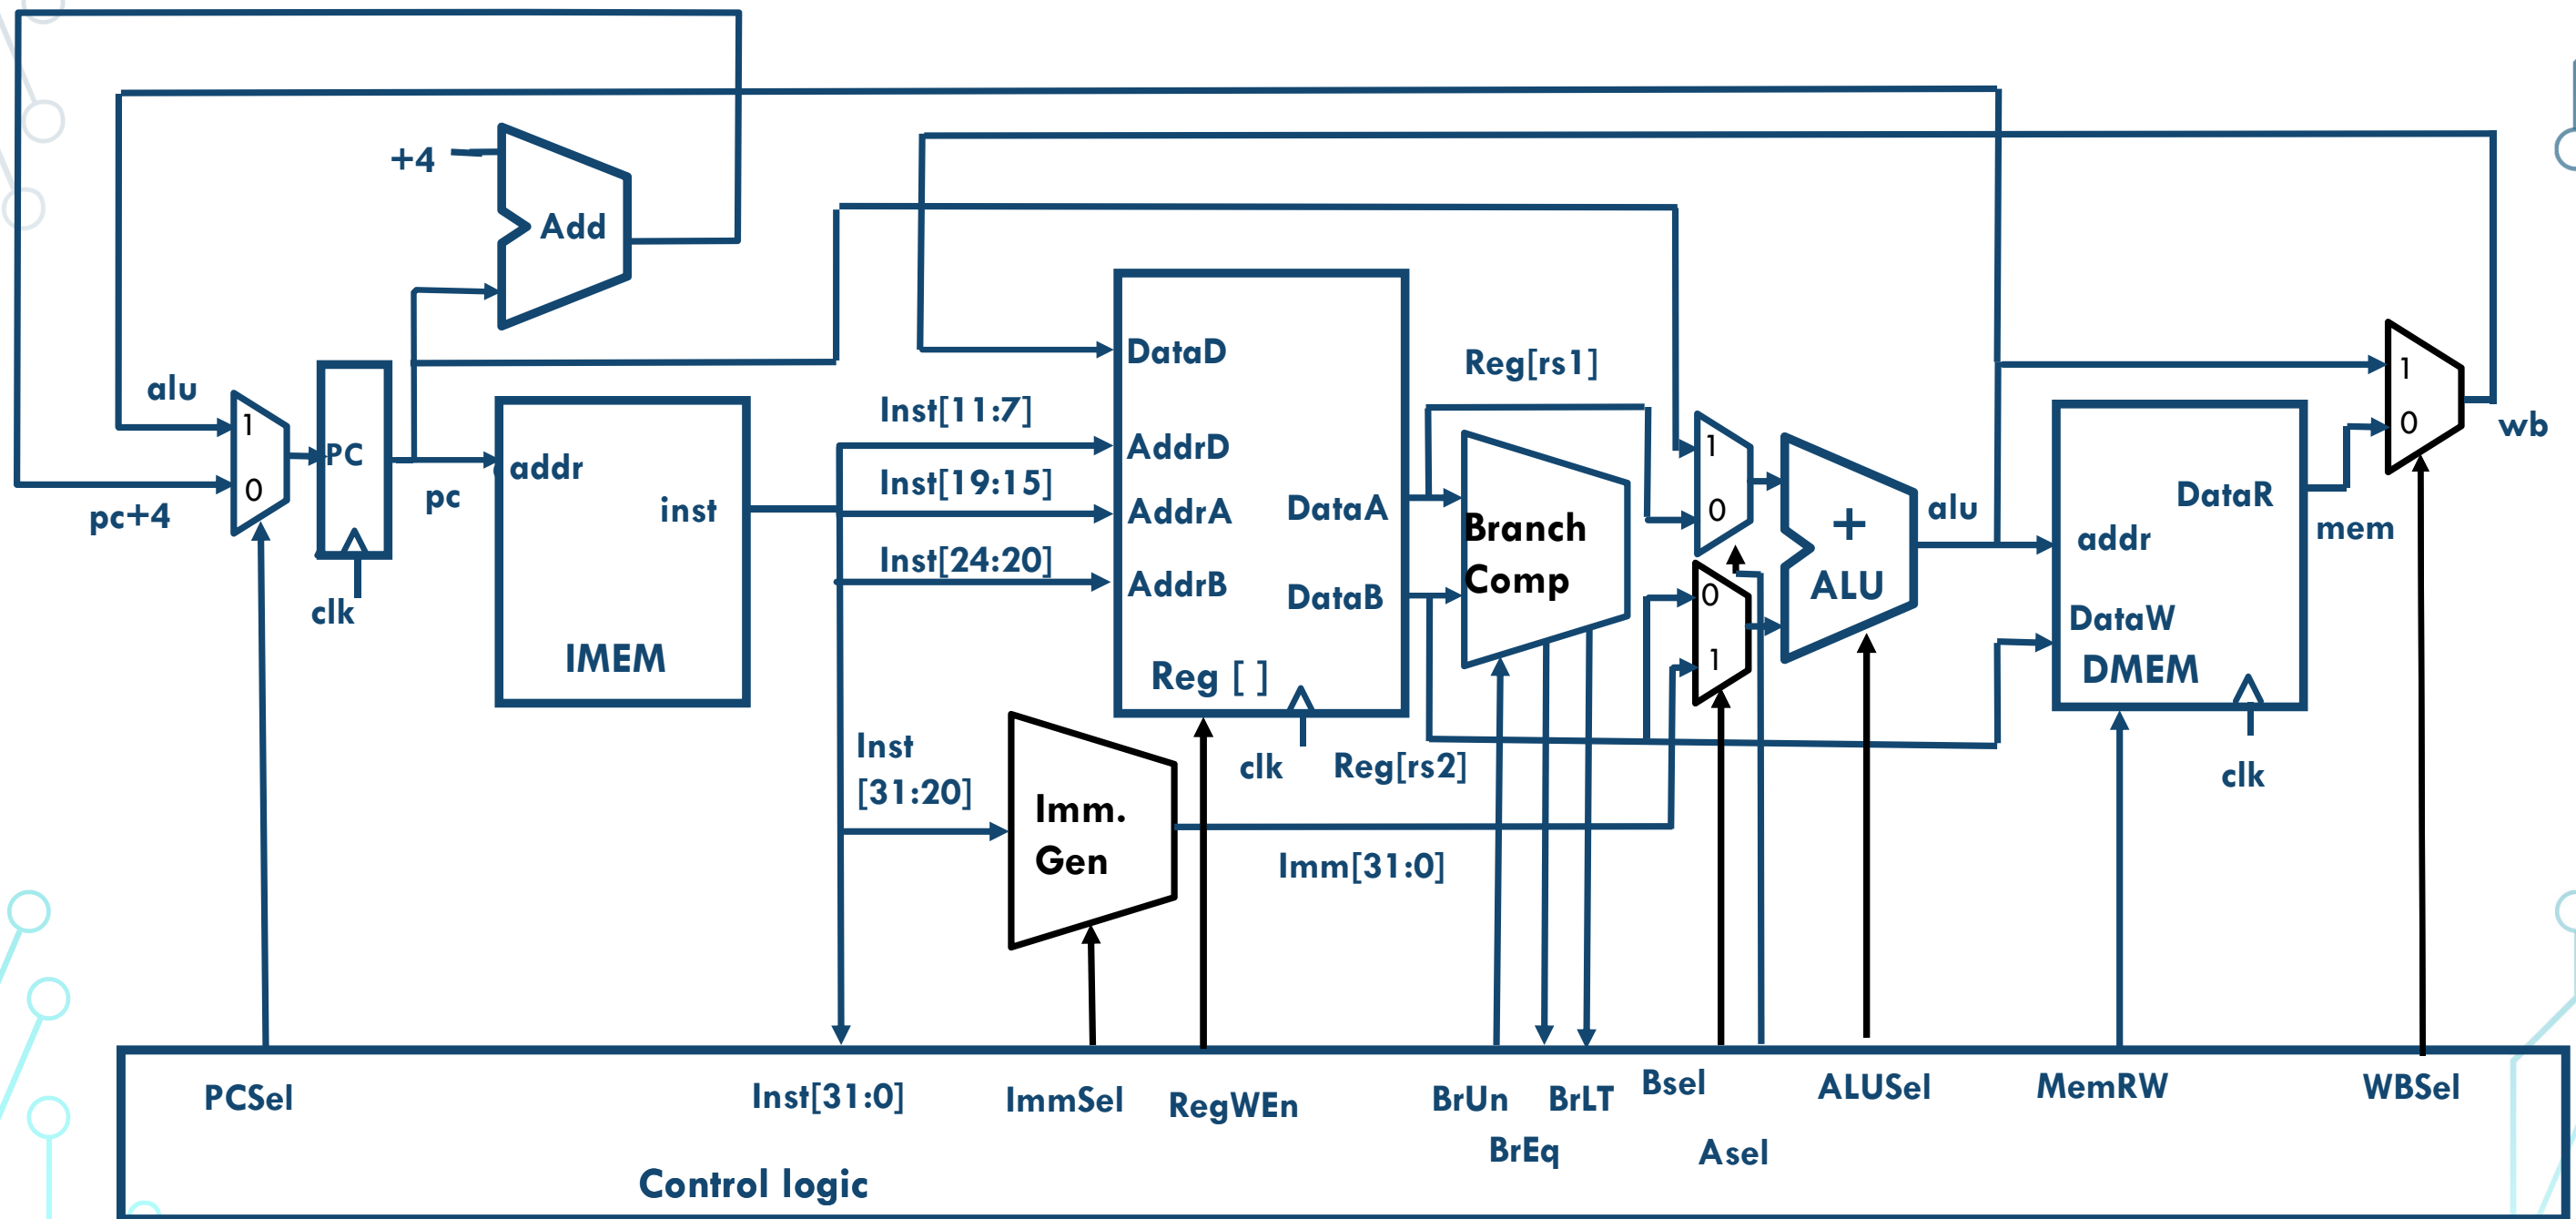
Otherwise, we would have yet another new encoding

Let's Add JALR (I-Format)

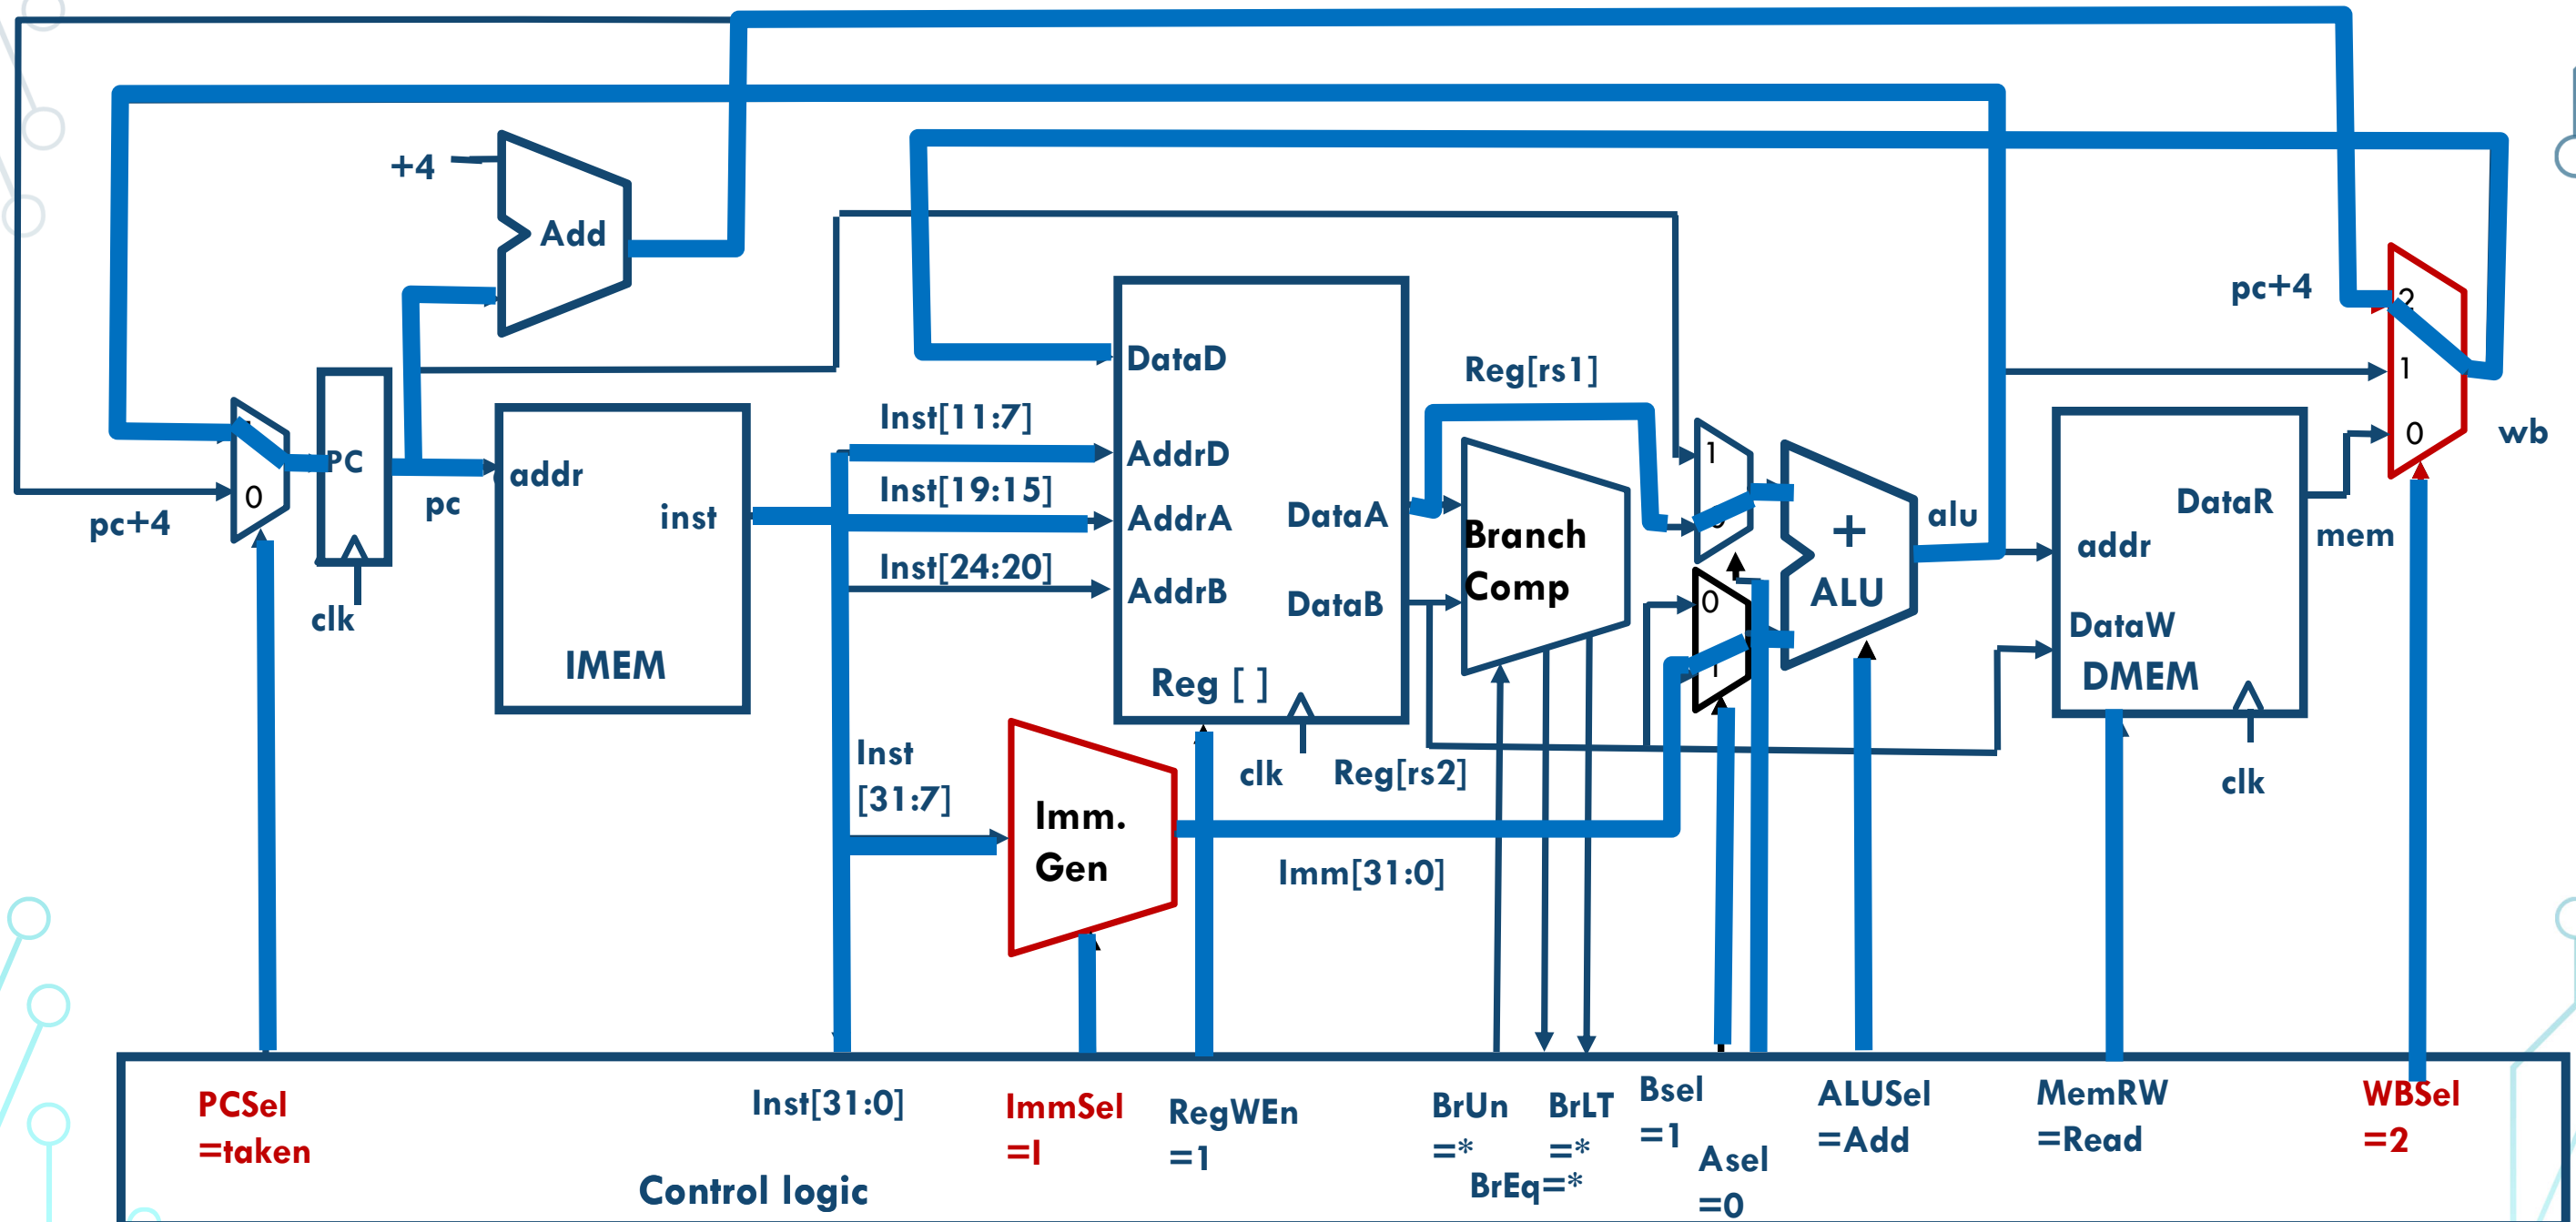


- JALR rd, rs, immediate
- Two changes to the state
 - Writes PC+4 to rd (return address)
 - Sets PC = rs + immediate
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes
 - LSB is ignored

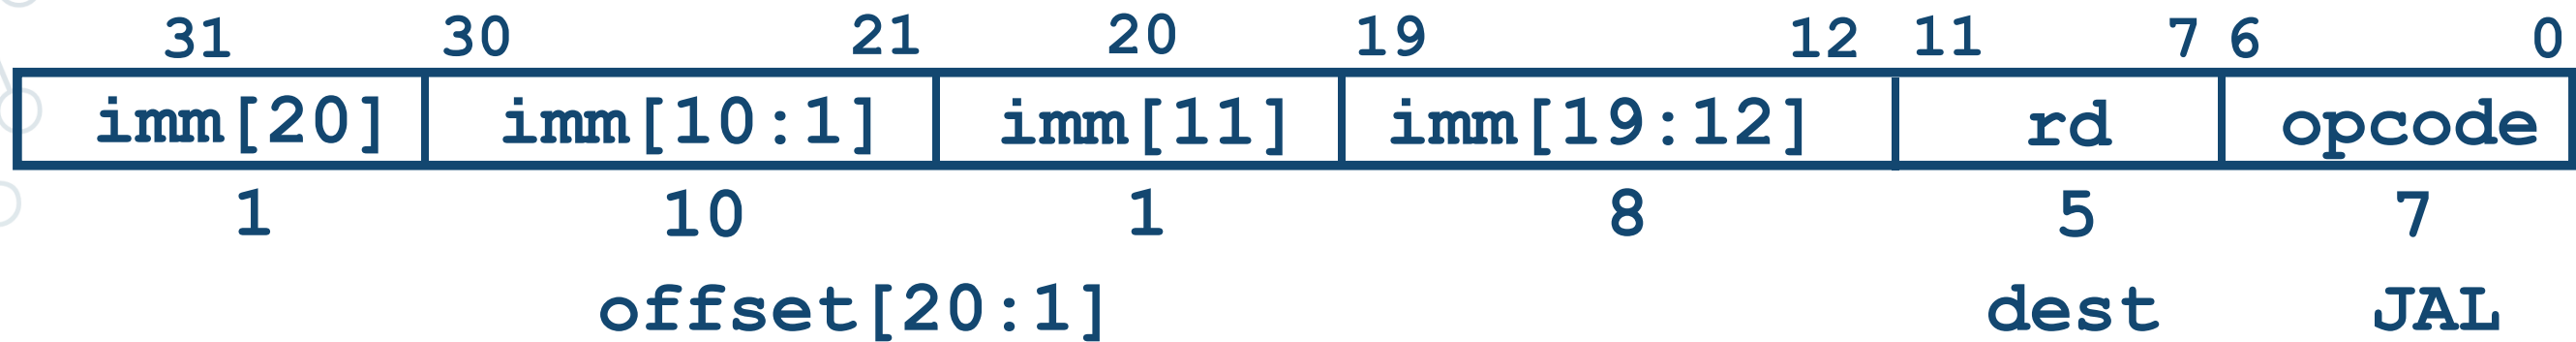
Datapath So Far, with Branches



Adding JALR

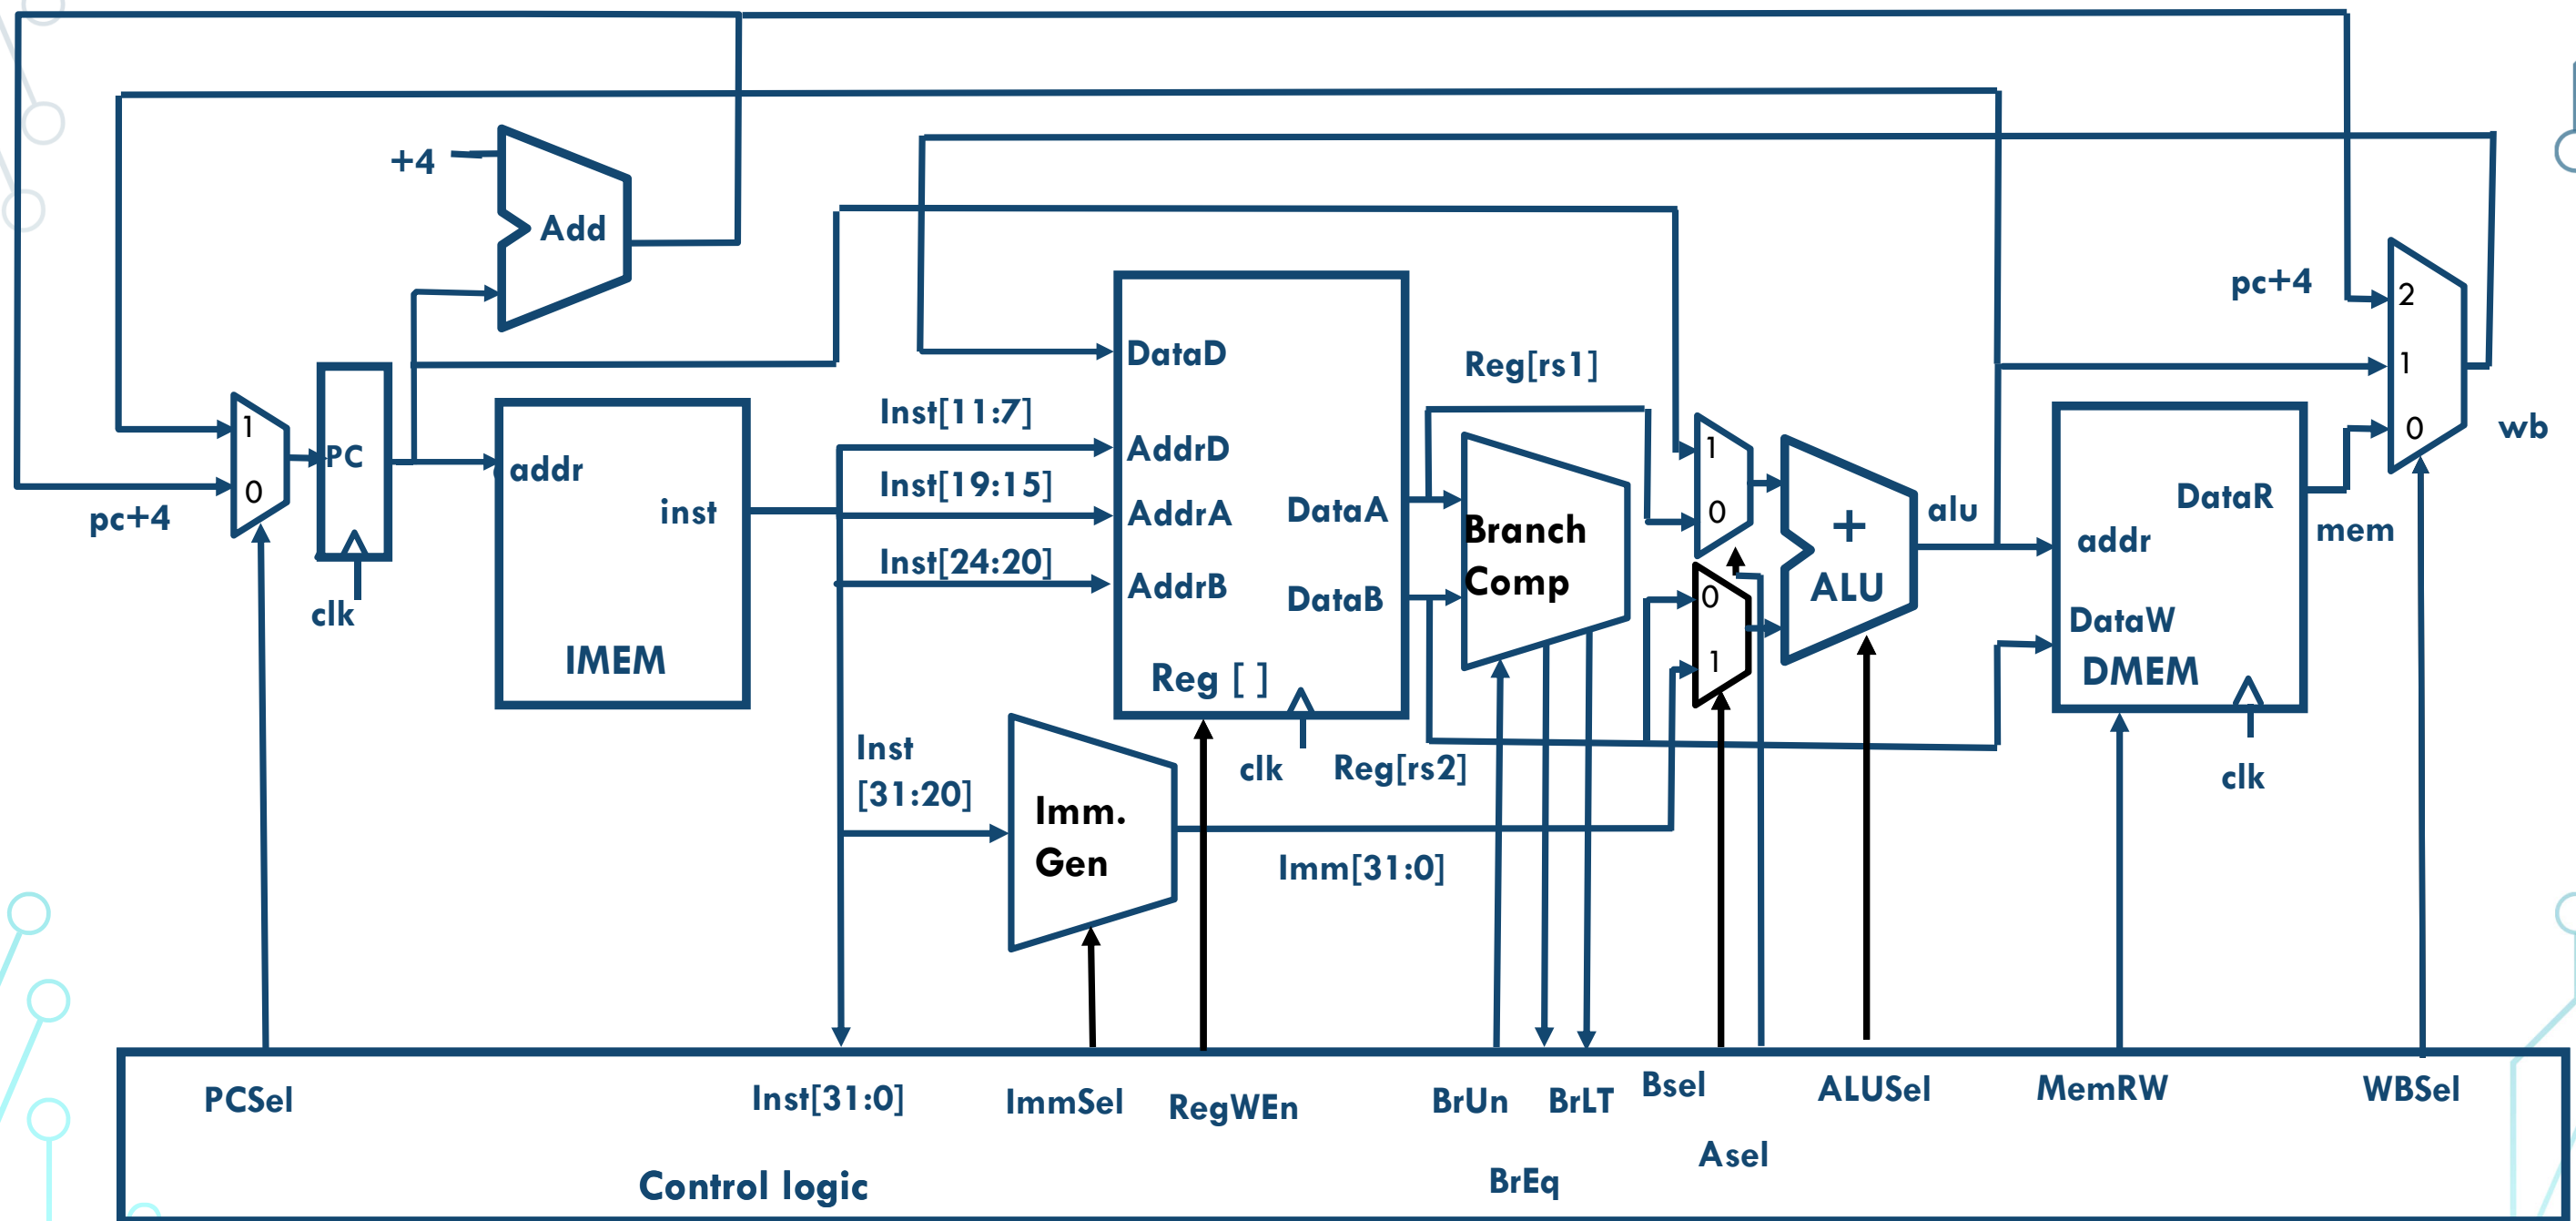


Adding JAL

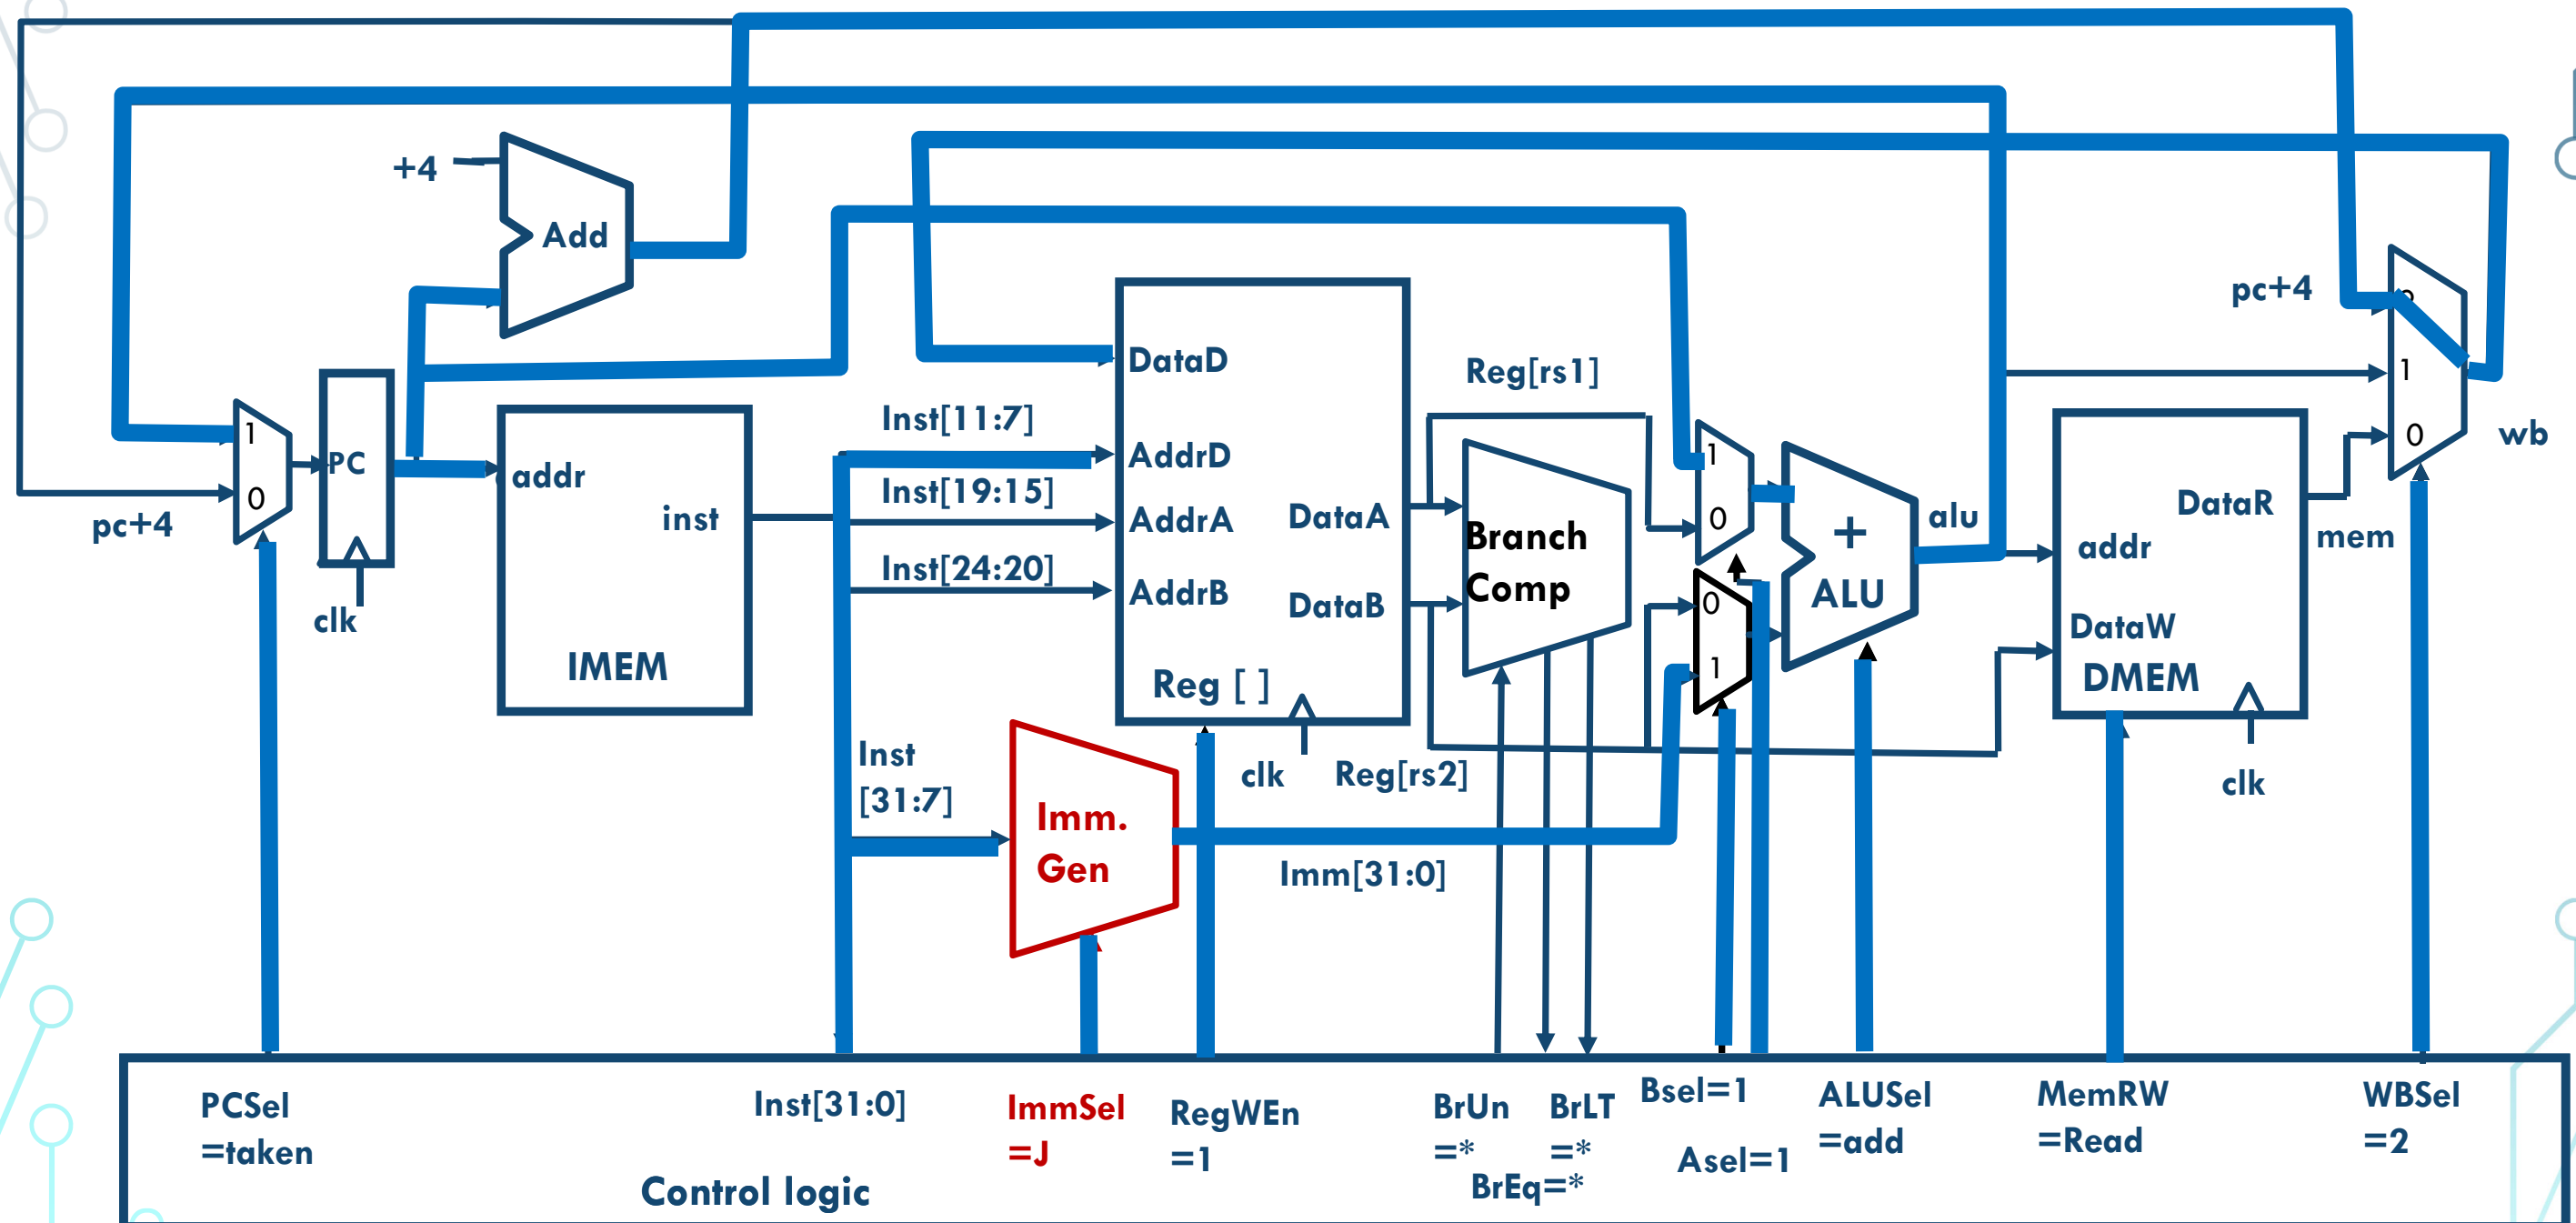


- JAL saves PC+4 in register rd (the return address)
- Set $PC = PC + \text{offset}$ (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

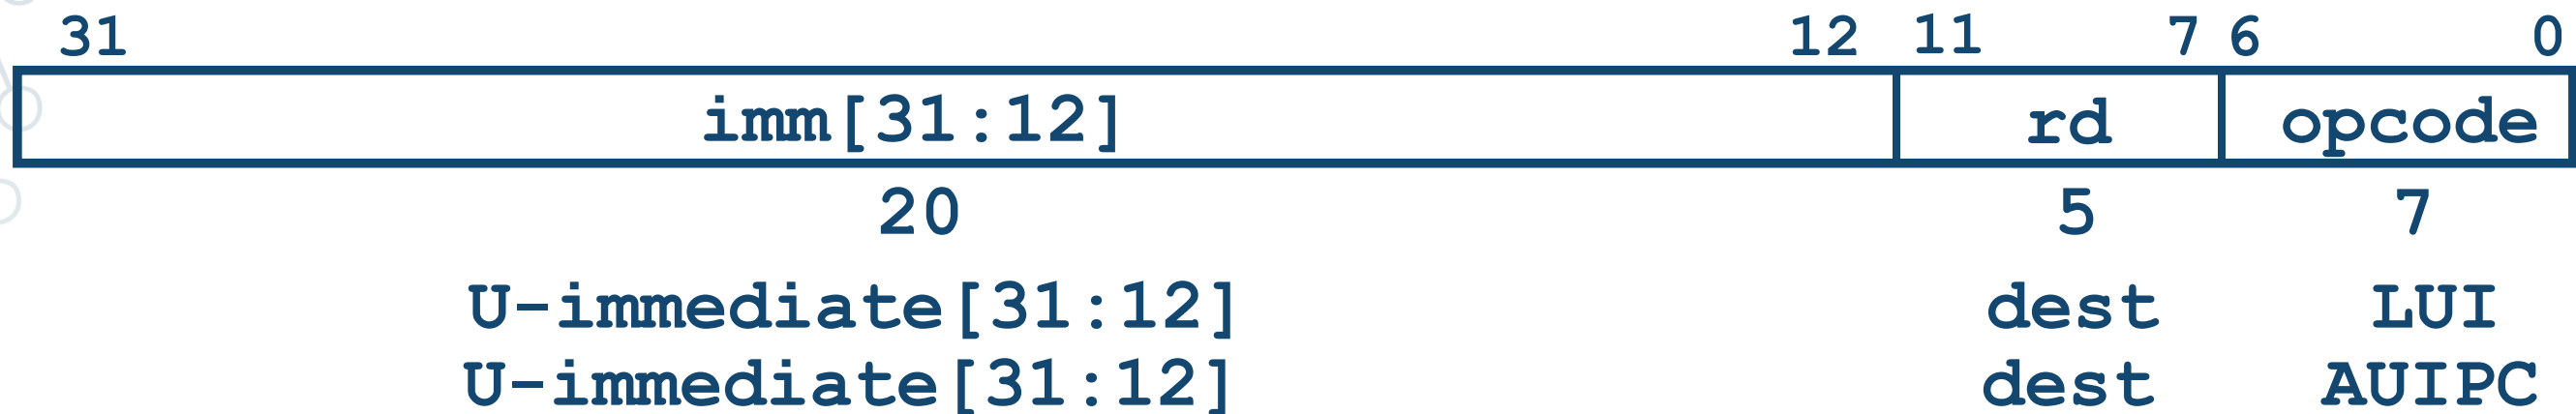
Datapath with JALR



Adding JAL

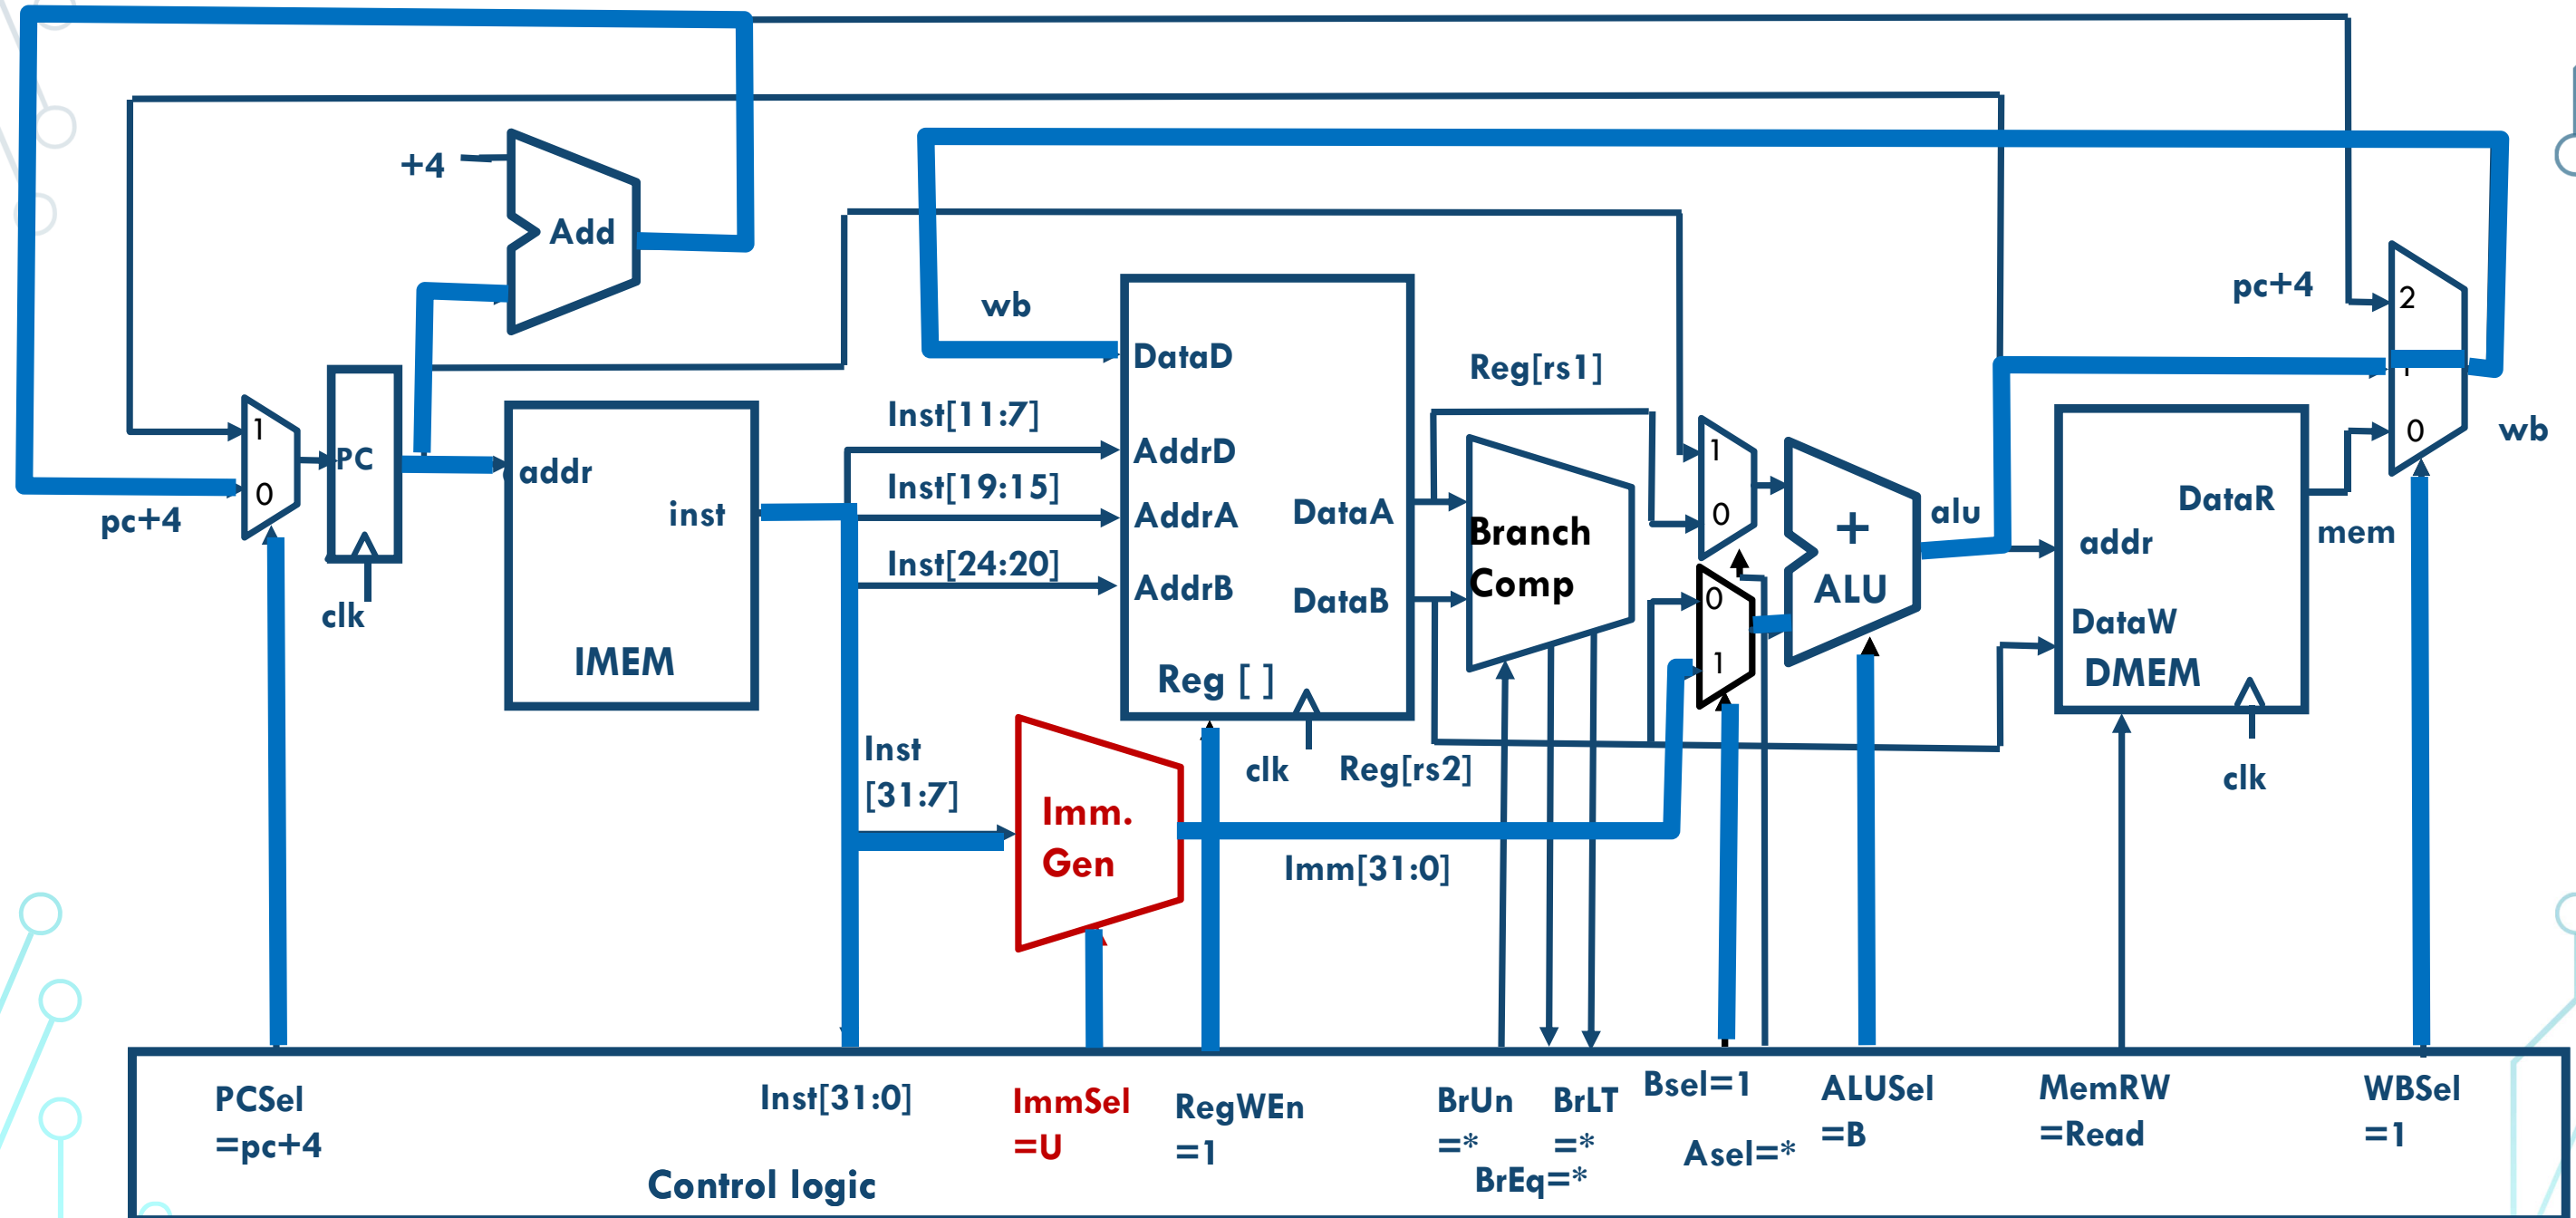


U-Format for “Upper Immediate” Instructions

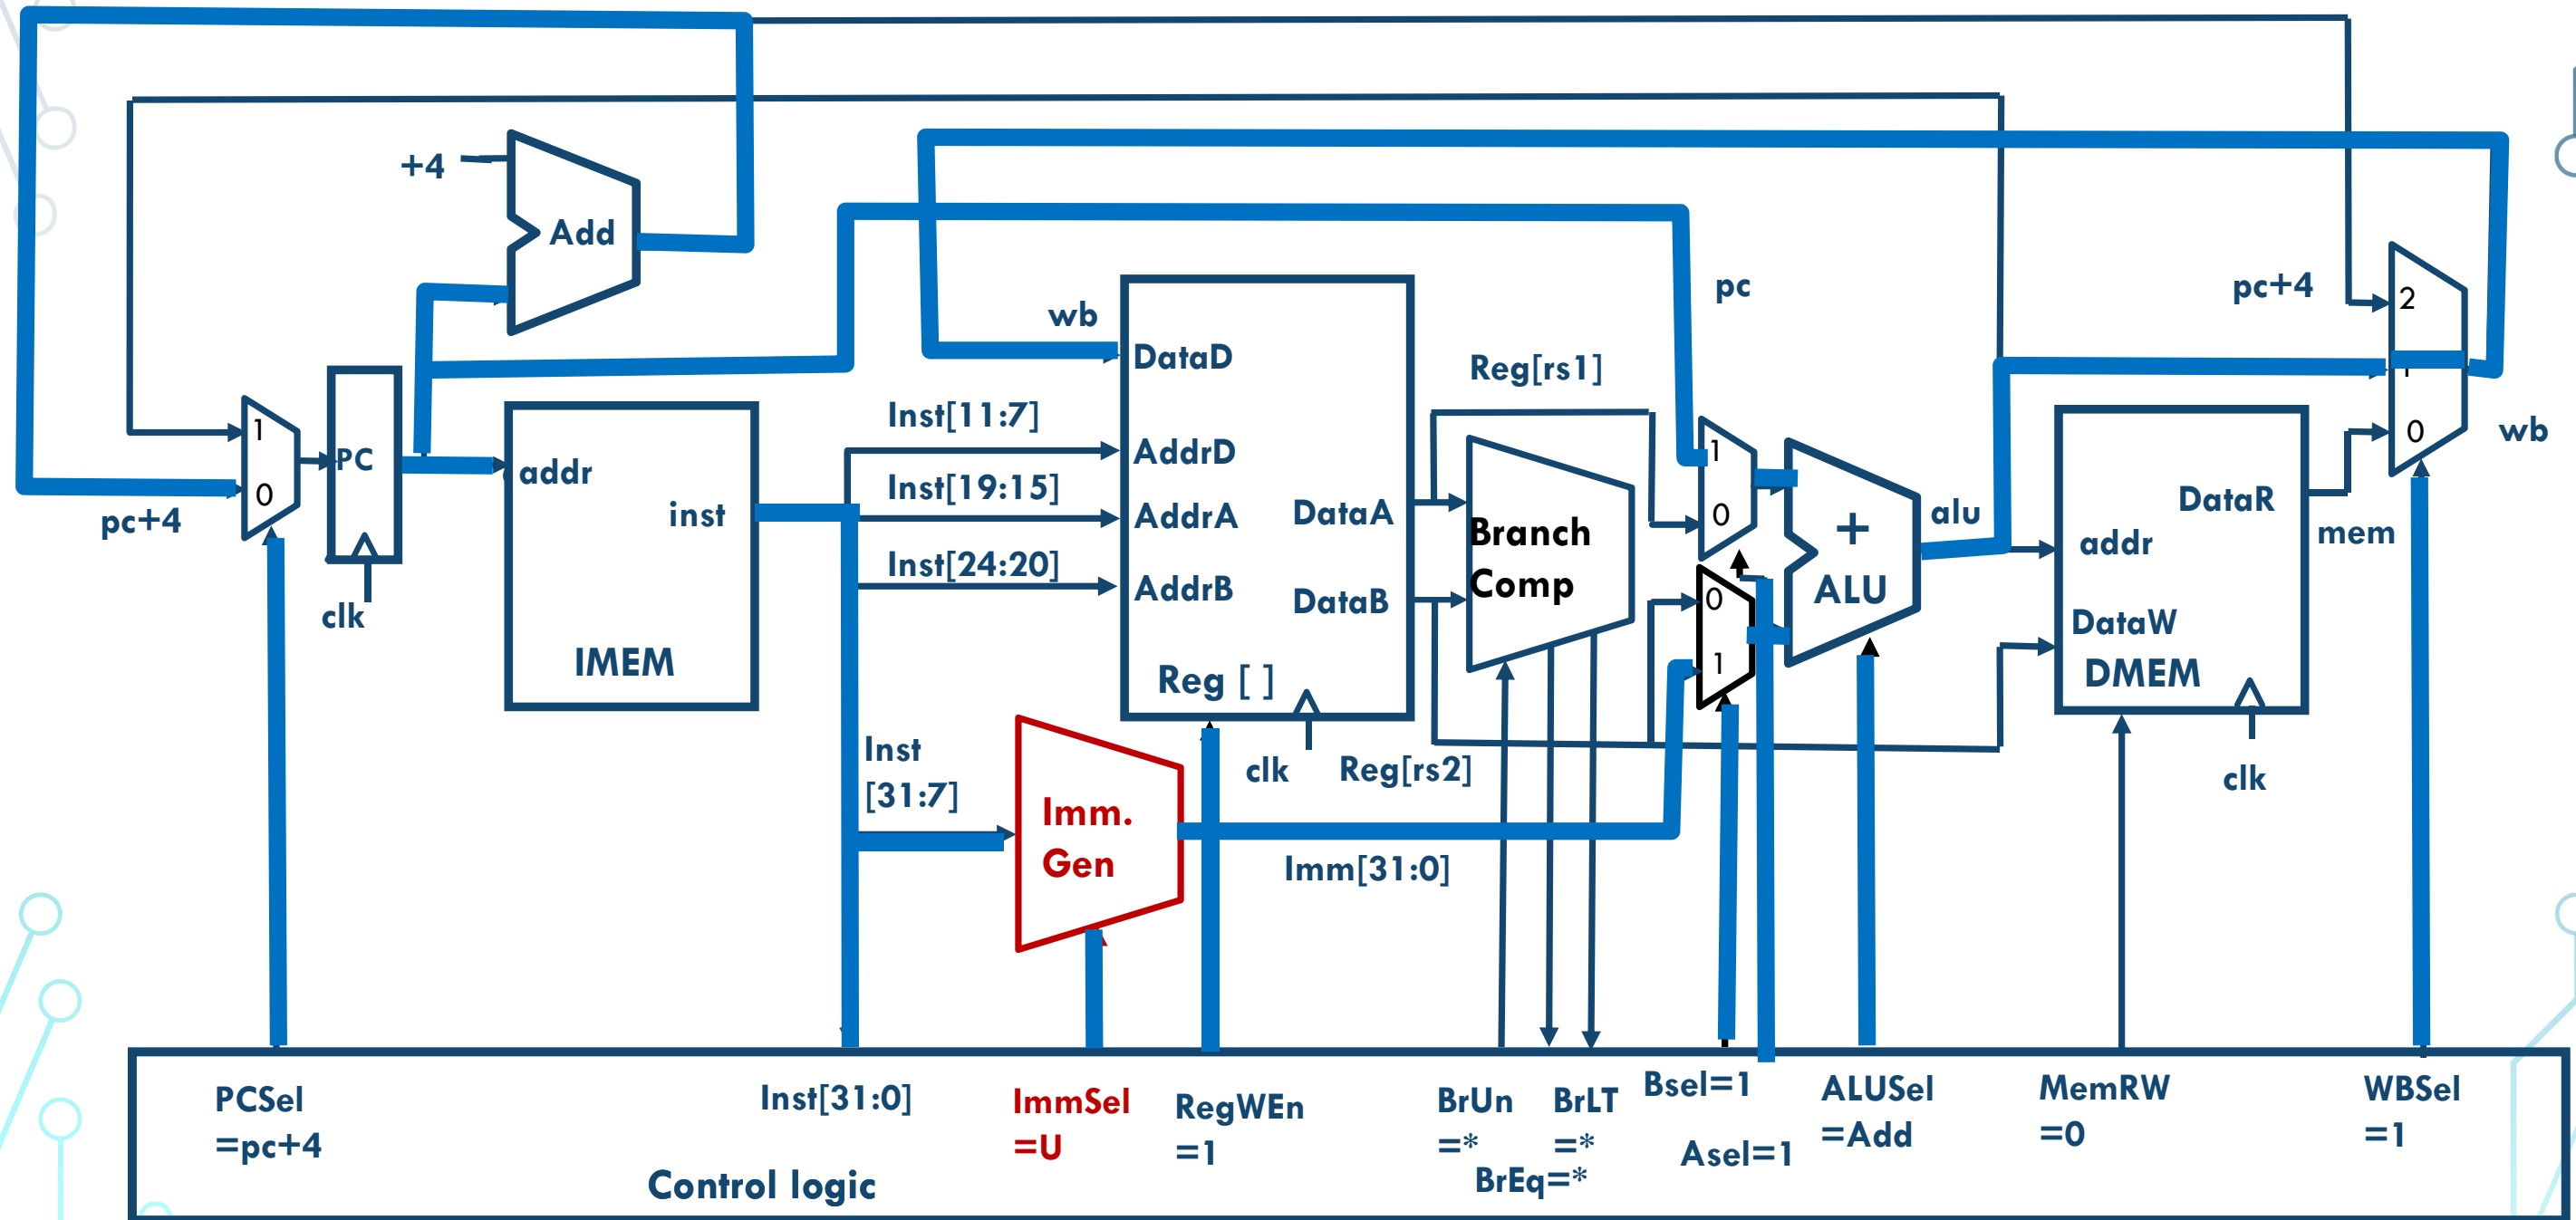


- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, rd
- Used for two instructions
 - LUI – Load Upper Immediate
 - AUIPC – Add Upper Immediate to PC

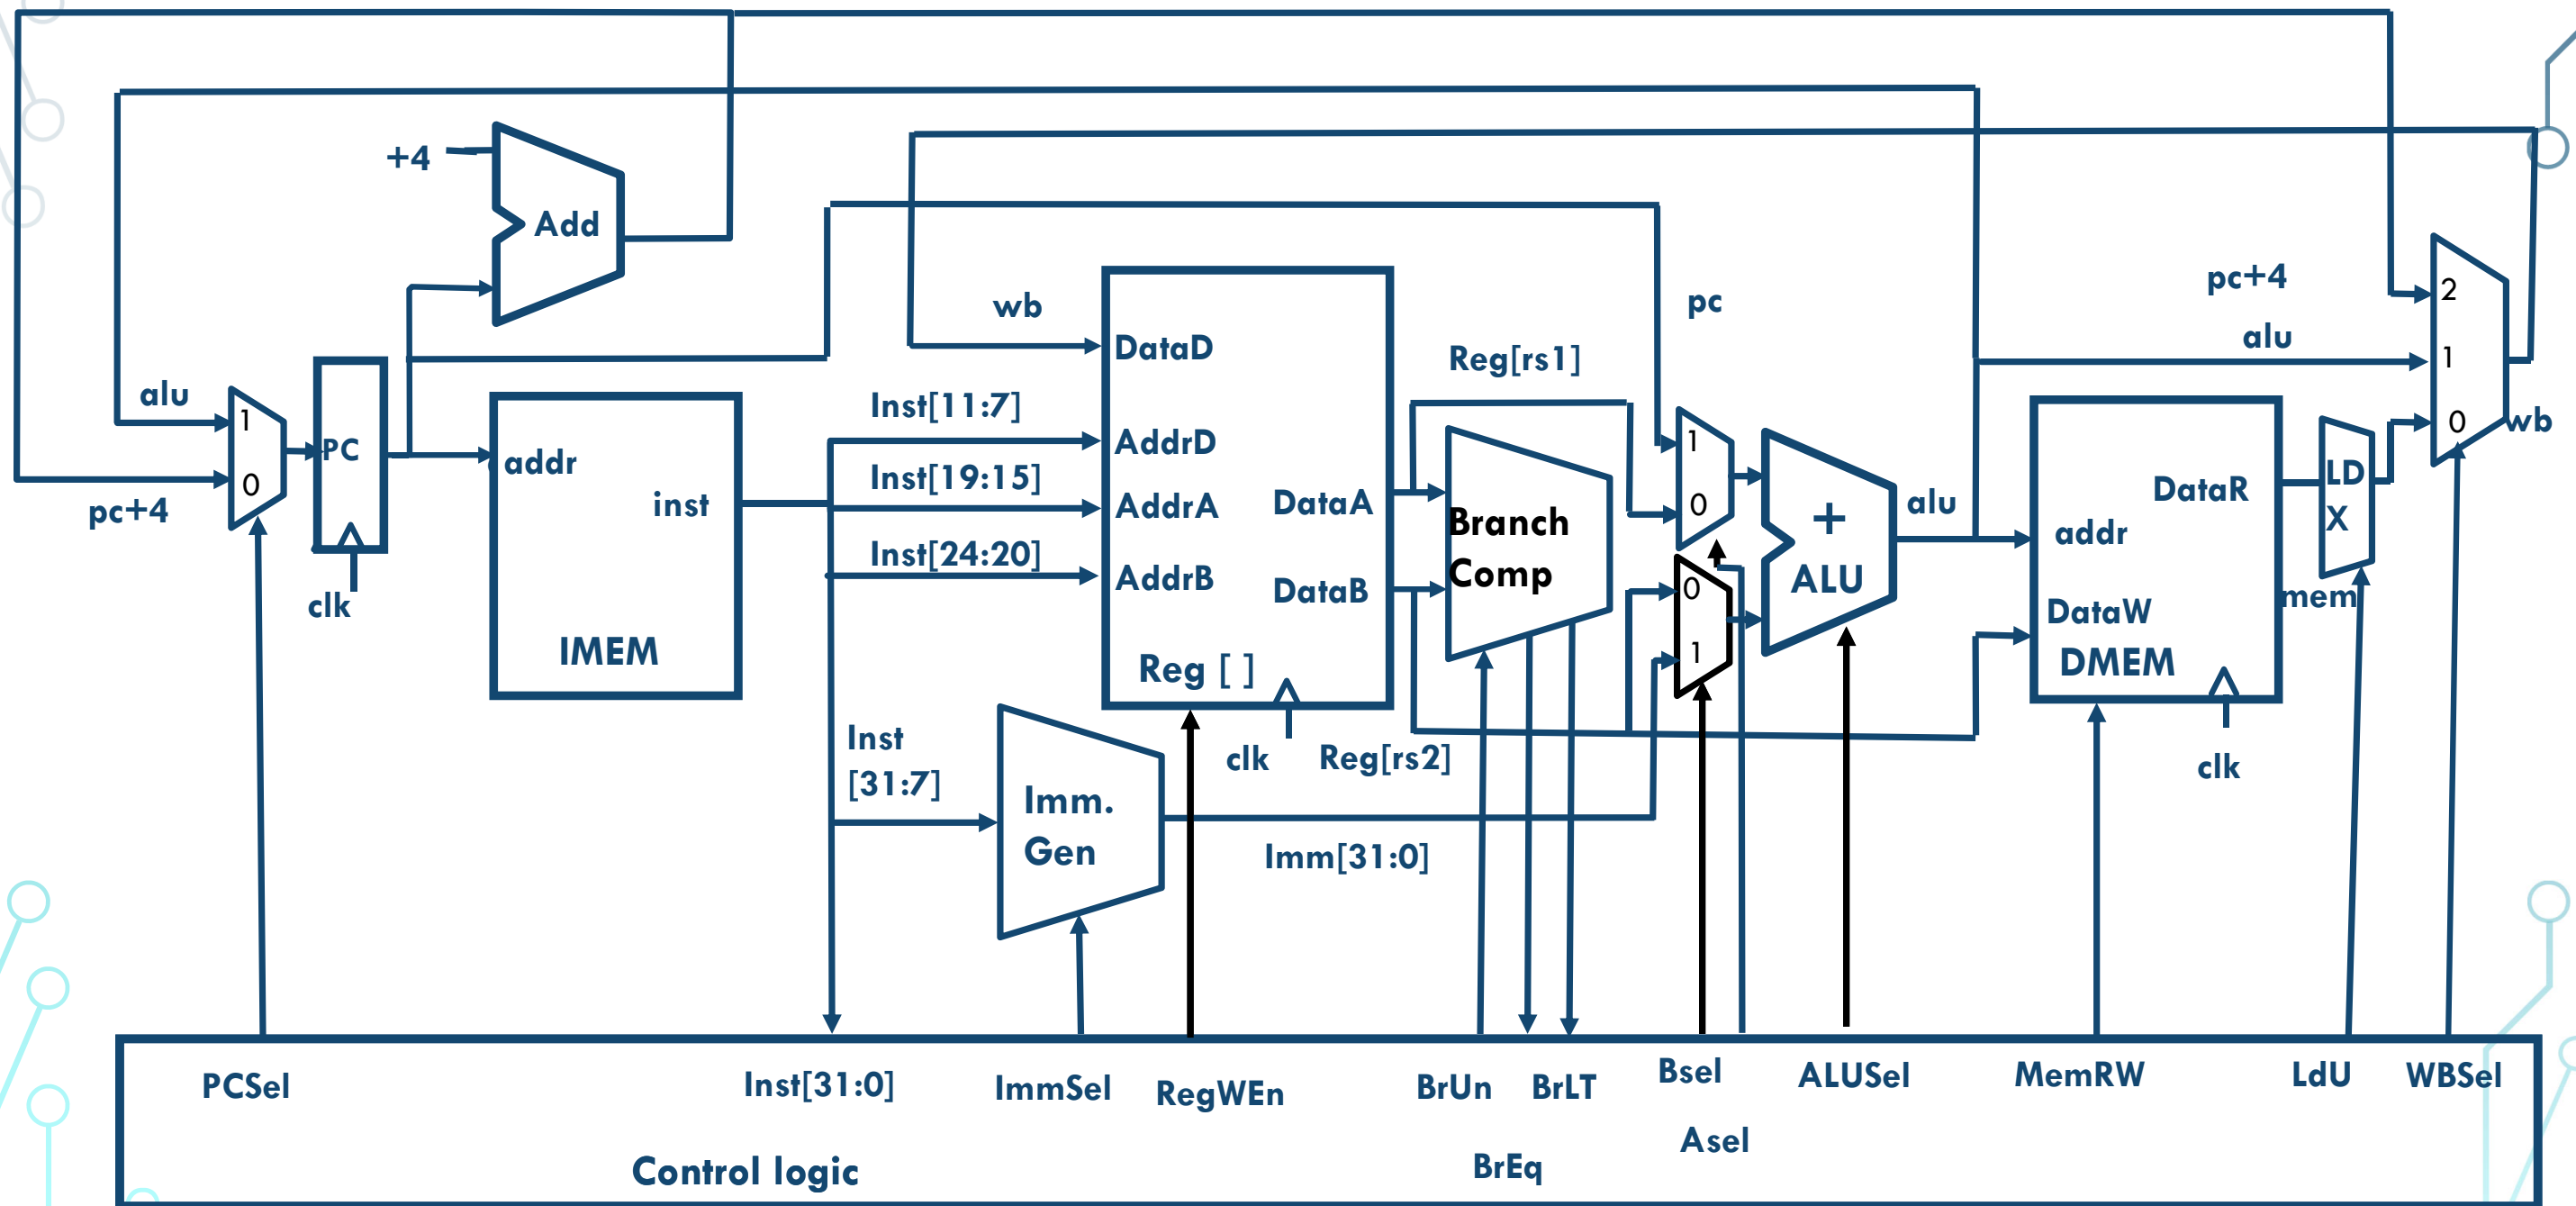
Implementing LUI



Implementing AUIPC



Complete RV32I Datapath!



Recap: Complete RV32I ISA

imm[31:12]				rd	0110111	LUI
imm[31:12]				rd	0010111	AUIPC
imm[20:10:11:19:12]				rd	1101111	JAL
imm[11:0]		rs1	000	rd	1100111	JALR
imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	1100011	BEQ
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	1100011	BNE
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	1100011	BLT
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	1100011	BGE
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	1100011	BLTU
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	1100011	BGEU
imm[11:0]		rs1	000	rd	0000011	LB
imm[11:0]		rs1	001	rd	0000011	LH
imm[11:0]		rs1	010	rd	0000011	LW
imm[11:0]		rs1	100	rd	0000011	LBU
imm[11:0]		rs1	101	rd	0000011	LHU
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	SB
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	SW
imm[11:0]		rs1	000	rd	0010011	ADDI
imm[11:0]		rs1	010	rd	0010011	SLTI
imm[11:0]		rs1	011	rd	0010011	SLTIU
imm[11:0]		rs1	100	rd	0010011	XORI
imm[11:0]		rs1	110	rd	0010011	ORI
imm[11:0]		rs1	111	rd	0010011	ANDI

0000000	shamt	rs1	001	rd	0010011	SLLI	
0000000	shamt	rs1	101	rd	0010011	SRLI	
0100000	shamt	rs1	101	rd	0010011	SRAI	
0000000	rs2	rs1	000	rd	0110011	ADD	
0100000	rs2	rs1	000	rd	0110011	SUB	
0000000	rs2	rs1	001	rd	0110011	SLL	
0000000	rs2	rs1	010	rd	0110011	SLT	
0000000	rs2	rs1	011	rd	0110011	SLTU	
0000000	rs2	rs1	100	rd	0110011	XOR	
0000000	rs2	rs1	101	rd	0110011	SRL	
0100000	rs2	rs1	101	rd	0110011	SRA	
0000000	rs2	rs1	110	rd	0110011	OR	
0000000	rs2	rs1	111	rd	0110011	AND	
0000	pred	succ	00000	000	00000	0001111	FENCE
0000	0000	0000	00000	001	00000	0001111	FENCE
000000000000	0000	0000	00000	000	00000	0110011	ECALL
0000000000001			00000	000	00000	1110011	EBREAK
csr		rs1	001	rd		1110011	CSRWRITE
csr		rs1	010	rd		1110011	CSRREAD
csr		rs1	011	rd		1110011	CSRROBUSTREAD
csr		zimm	101	rd		1110011	CSRWRITE
csr		zimm	110	rd		1110011	CSRREAD
csr		zimm	111	rd		1110011	CSRROBUSTREAD

Not in the ratified base ISA anymore, but needed for practical hardware

Not in the ratified base ISA anymore, but needed for practical hardware

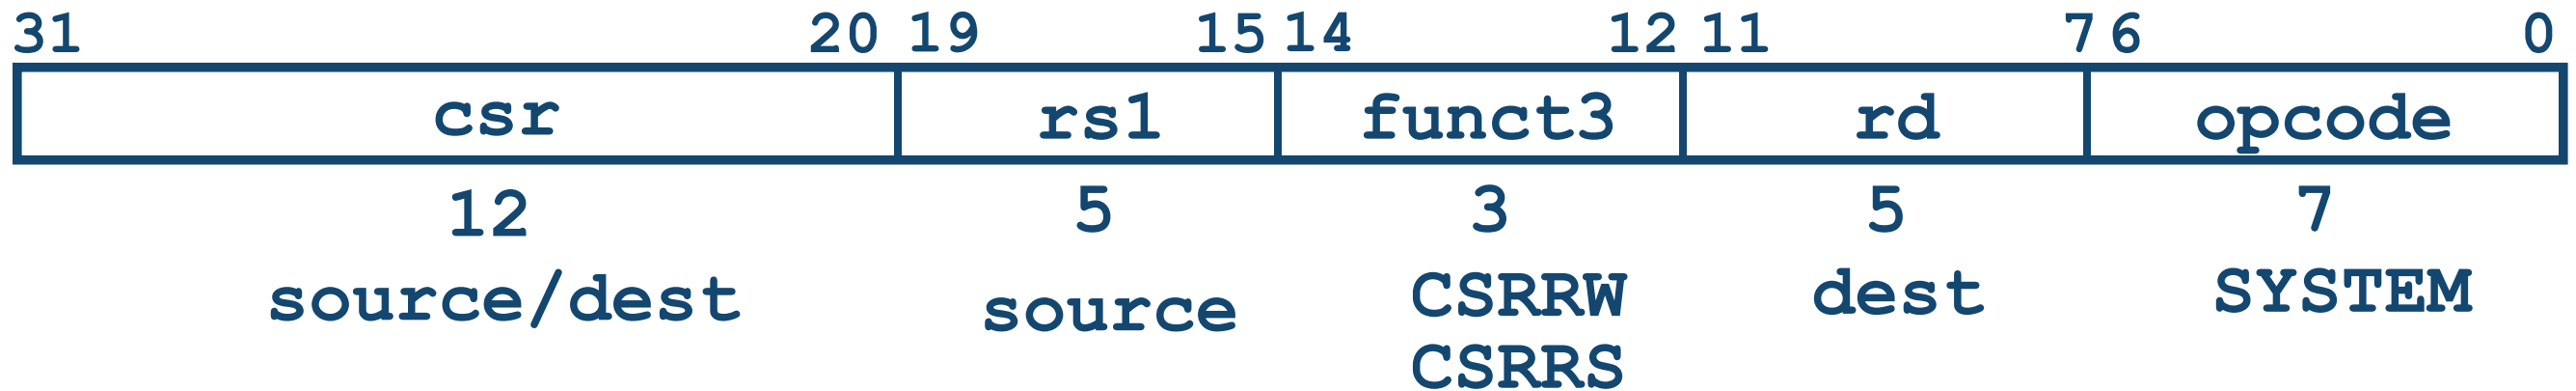
- RV32I has 47 instructions
- 37 instructions are enough to run any C program

Summary of RISC-V Instruction Formats

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0				
funct7					rs2			rs1			funct3			rd		opcode		R-type
imm[11:0]							rs1			funct3			rd		opcode		I-type	
imm[11:5]				rs2			rs1			funct3			imm[4:0]		opcode		S-type	
imm[12 10:5]				rs2			rs1			funct3			imm[4:1 11]		opcode		B-type	
imm[31:12]										rd			opcode		U-type			
imm[20 10:1 11]]						imm[19:12]				rd			opcode		J-type			

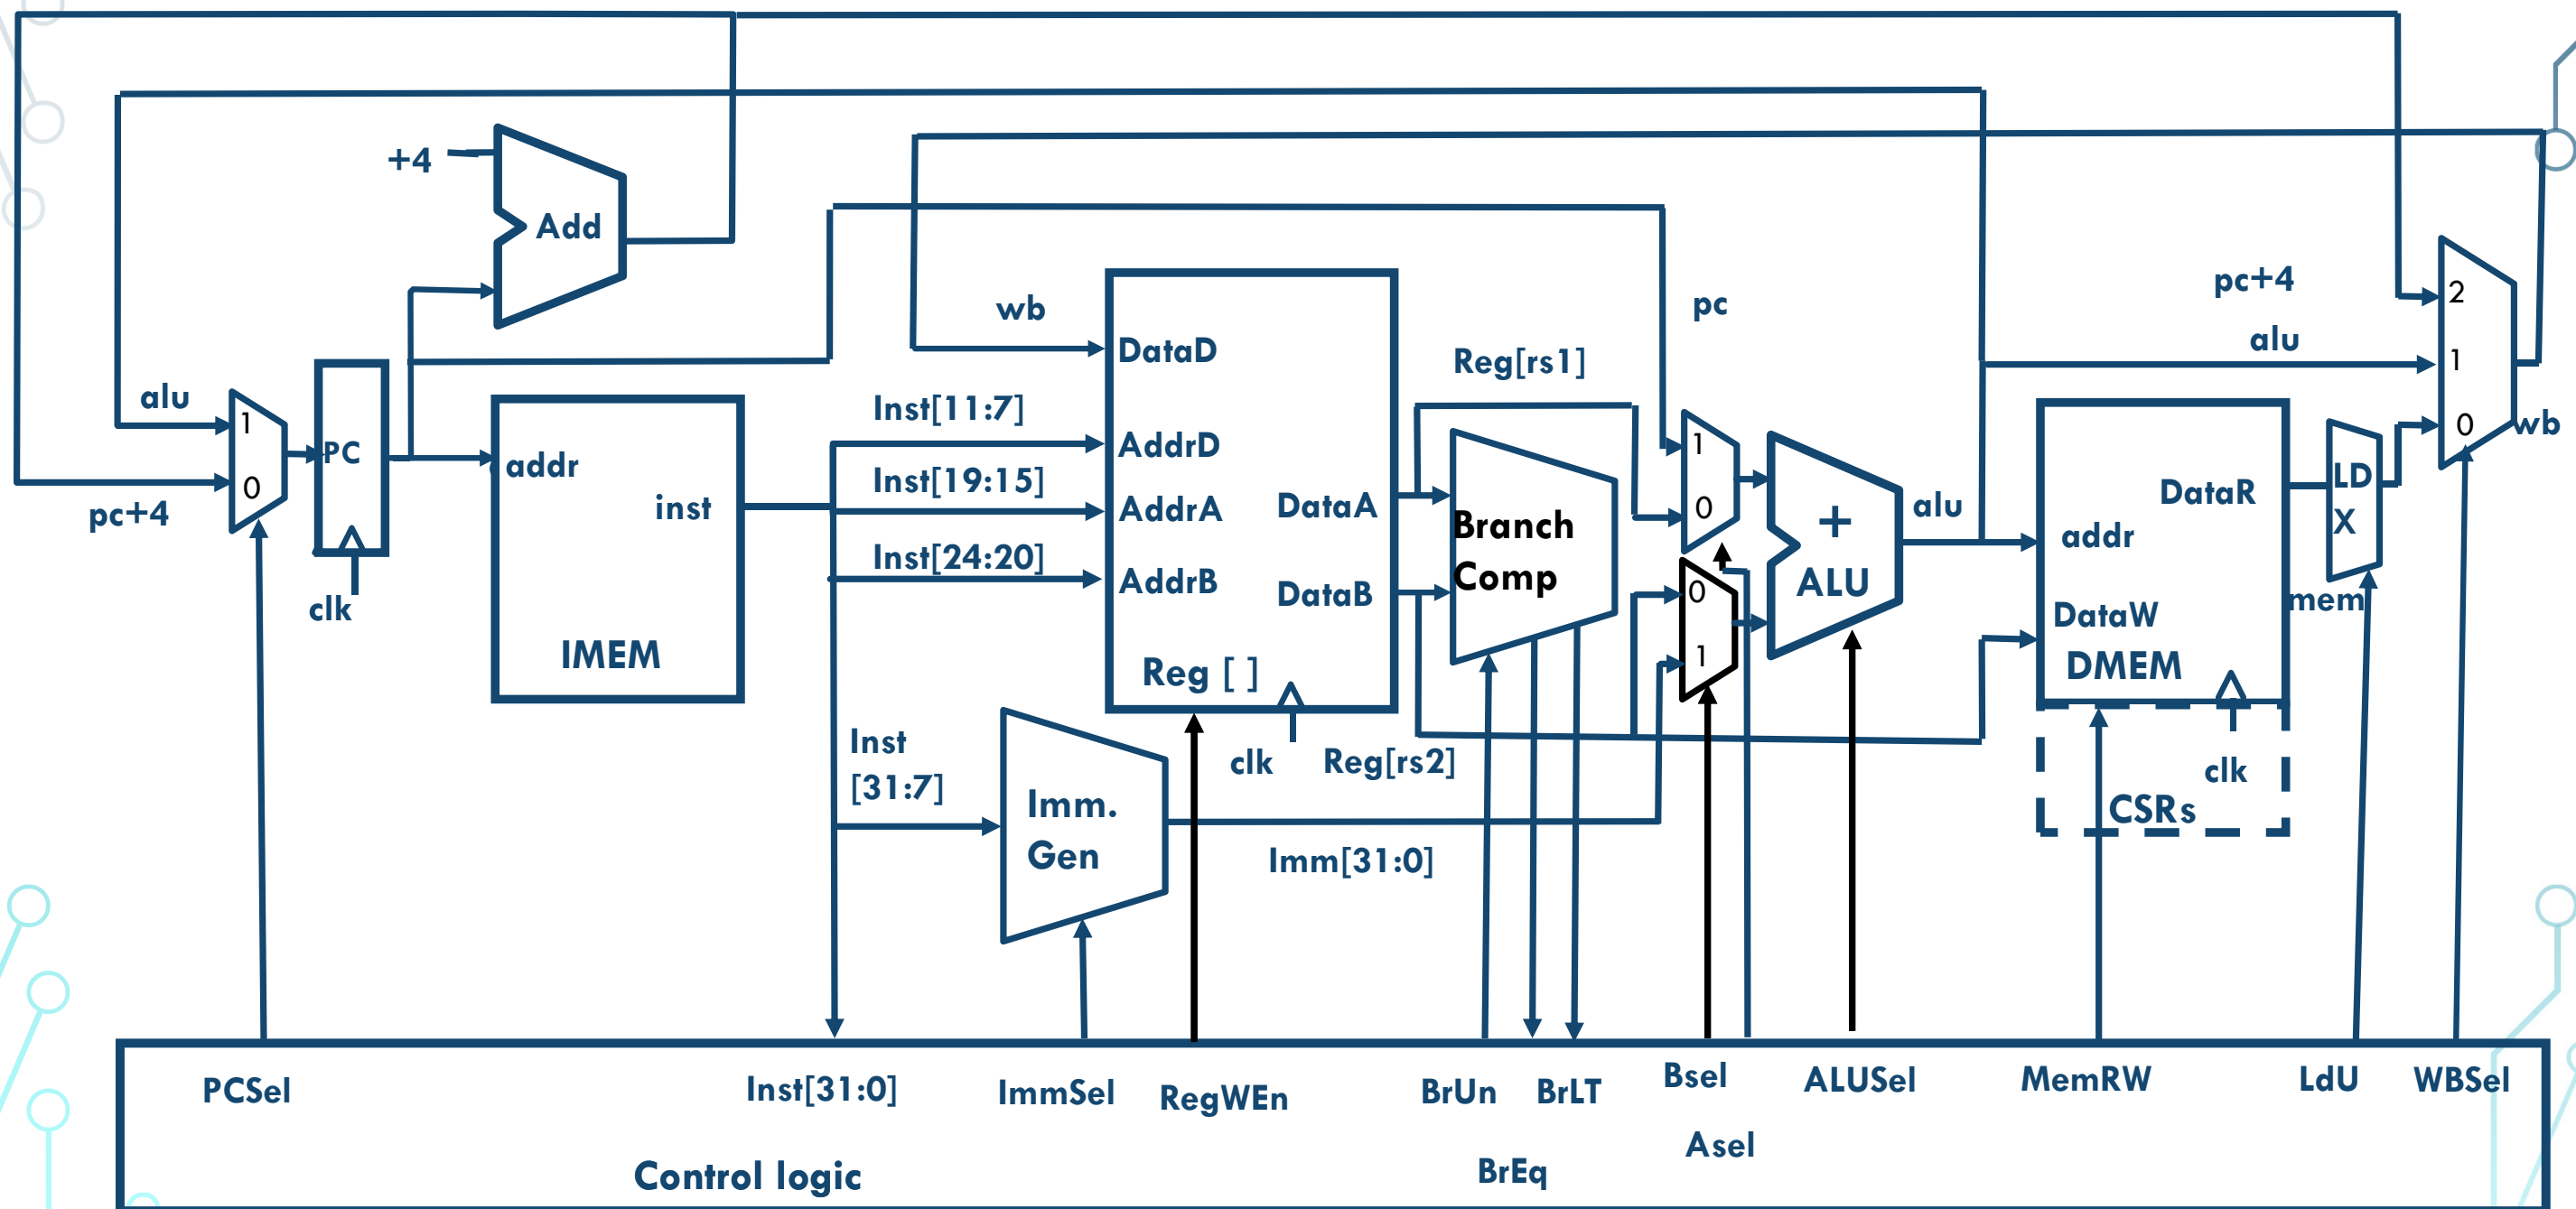
Control and Status Registers (CSRs)

- 4096 CSRs in a separate address space



- **csrrw** reads the old value of CSR, zero-extends and writes to **rd**
- Initial value of **rs1** is written to CSR
- Pseudo-instructions: **csrr**, **csrw** (**csrrw x0, csr, rs1**)

Add the `csrrw`!



Summary

- We have covered the implementation of the base ISA for RV32I
- Implementation suggested is straightforward, yet there are modalities in how to implement it well at gate level.
- Still need control
- Single-cycle datapath is slow – need to pipeline it